

DESIGN, IMPLEMENTATION, AND APPLICATIONS OF
PEER-TO-PEER VIRTUAL PRIVATE NETWORKS
FROM GRIDS TO SOCIAL NETWORKS

By

DAVID ISAAC WOLINSKY

A DISSERTATION PRESENTED TO THE GRADUATE SCHOOL
OF THE UNIVERSITY OF FLORIDA IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS FOR THE DEGREE OF
DOCTOR OF PHILOSOPHY

UNIVERSITY OF FLORIDA

2011

© 2011 David Isaac Wolinsky

I dedicate this to family and those whose have supported me.

ACKNOWLEDGMENTS

Over the past 5 years, there has been many constants and many changes, soon there will be only changes. Though through it all, I have been surrounded by wonderful people who have helped and encouraged me to succeed and can only hope that the relationship was beneficial for them as well. To begin naming names, I will begin with my advisor, Professor Renato Figueiredo. I greatly appreciate the time he has invested into me and his wisdom shared with me. I am greatly blessed to have worked so closely with a professor whom I work so well with. That leads me into Professor P. Oscar Boykin, the other head of the ACIS P2P Group, who along with Professor Figueiredo, molded me into the bold Ph.D. I am today. Professor Boykin also helped enrich my design and development skills, for which, I am extremely grateful. Rounding out the ACIS professors, leads to Professor Jose Fortes, who has always been a good source of wisdom and encouragement. As leader of the ACIS lab, Professor Fortes has always been very generous in providing both his time, which is why I am very appreciative to have him as a member on my Ph.D. committee. I would also like thank Professors Shigang Chen and Y. Peter Sheng for their time investments in my research and whose comments have been invaluable in shaping my dissertation.

My peers and family have also been critical sources of support, encouragement, and wisdom. I am thankful to the members of the P2P group, both past and present, namely, Tae Woong Choi, Heung Sik Eom, Arijit Ganguly, Kyungyong Lee, Yonggang Liu, Pierre St. Juste, and Jiangyan Xu, whose comments and contributions have paved the way for my research. I am thankful to members of my sports groups, both the Larsen-Benton Basketball Association and the Badminton Group for their friendships, as they provided a means to redirect frustrations developed along the way. I appreciate the hard work and dedication of my fellow Archer colleague, Girish Venkatasubramanian. My gratitude goes to the kind office ladies who assisted me so much, Catherine Reeves, Janet Sloan, and Dina Stoeber. I am thankful for the time put forth by my Grid Appliance

colleagues, Panoat Chuchaisri and Arjun Prakash. I would like to thank Priya Bhatt, Bingyi Cao, Xin Fu, Selvi Kadriavel, and Prapaporn Rattanatamrong for their kind hearts and encouragement and, in some cases, their spicy food. I would like to thank Donna for her support and encouragement throughout the years, likewise, I have been blessed to have parents that have encouraged me to press forward and achieve my goals in life.

Research is a collaborative effort that, for me at least, involves both professional and home life. My success has largely been the result of the quality individuals that I have been fortunate enough to have in my life. It is for those already mentioned and those remembered that this dissertation is owed. Thank you so much.

TABLE OF CONTENTS

	<u>page</u>
ACKNOWLEDGMENTS	4
LIST OF TABLES	10
LIST OF FIGURES	11
ABSTRACT	13
 CHAPTER	
1 INTRODUCTION	17
1.1 Virtual Private Network Basics	20
1.2 Computer Network Architectures	23
1.3 Structured Overlays	25
1.4 Network Asymmetries	28
1.5 Contributions	30
2 VIRTUAL NETWORK CONFIGURATION AND ORGANIZATION	34
2.1 Network Configuration	35
2.1.1 Centralized VPN Systems	35
2.1.2 Centralized P2P VPN Systems	36
2.1.3 Decentralized VPN Systems	37
2.1.4 Unstructured P2P VPN Systems	37
2.1.5 Structured P2P VPN Systems	37
2.2 Local Configuration	39
2.2.1 Local VPN Architecture	41
2.2.2 Address Resolution	43
2.2.3 Address Allocation	46
2.2.4 Domain Name Servers and Services	48
2.3 Supporting Migration	49
2.4 Evaluation of VPN Network Configuration	53
2.5 Evaluation of VPN Local Configuration	55
2.5.1 On the Grid	56
2.5.2 In the Clouds	60
3 BOOTSTRAPPING PRIVATE OVERLAYS	65
3.1 Current Bootstrap Solutions	69
3.2 Core Requirements	71
3.2.1 Reflection	71
3.2.2 Relaying	73
3.2.3 Rendezvous	74
3.3 Implementations	75

3.3.1	Using Brunet	76
3.3.2	Using XMPP	78
3.4	Evaluating Overlay Bootstrapping	80
3.4.1	Deployment Experiments	80
3.4.2	Deployment Experiences	82
4	FROM OVERLAYS TO SECURE VIRTUAL PRIVATE NETWORKS	86
4.1	Experimental Environment	87
4.2	Towards Private Overlays	88
4.2.1	Time to Bootstrap a Private Overlay	89
4.2.2	Overhead of Pathing	92
4.3	Security for the Overlay and the VPN	92
4.3.1	Implementing Overlay Security	93
4.3.2	Overheads of Overlay Security	96
4.3.2.1	Adding a Single Node	97
4.3.2.2	Bootstrapping an Overlay	98
4.3.3	Discussion	98
4.4	Handling User Revocation	99
4.4.1	DHT Revocation	99
4.4.2	Broadcast Revocation	100
4.4.3	Evaluation of Broadcast	100
4.4.4	Discussion	100
4.5	Managing and Configuring the VPN	102
4.6	Leveraging Trust from Online Social Networks	104
4.6.1	Architecture	104
4.6.2	Leveraging Trust From Facebook	105
4.6.3	Leveraging Trust from XMPP	106
4.6.4	Address Allocations and Discovery	107
4.7	Related Work	107
4.7.1	VPNs	107
4.7.2	P2P Systems	109
5	EXTENSIONS TO P2P OVERLAYS AND VIRTUAL NETWORKS	110
5.1	Built-in Self-Simulation	110
5.1.1	Time-Based Events	111
5.1.2	Network Communication	112
5.1.3	User Actions	114
5.1.4	The Rest of the System	115
5.1.5	Optimizations	115
5.2	Efficient Relays	117
5.2.1	Motivation for Relays in the Overlay	118
5.2.2	Comparing Relay Selection	119
5.3	Policies for Establishing Direct Connections	121
5.3.1	Limitations	121

5.3.2	On-Demand Connections	123
5.4	Broadcasting IP Broadcast and Multicast Packets Via the Overlay	125
5.5	Full Tunnel VPN Operations	126
5.5.1	The Gateway	128
5.5.2	The Client	128
5.5.3	Full Tunnel Overhead	130
6	AD-HOC, DECENTRALIZED GRIDS	132
6.1	WOWs	136
6.1.1	P2P Overlays	136
6.1.2	Virtual Private Networks	137
6.1.3	Virtual Machines in Grid Computing	138
6.2	Architectural Overview	139
6.2.1	Web Interface and the Community	141
6.2.2	The Organization of the Grid	143
6.2.2.1	Selecting a Middleware	144
6.2.2.2	Self-Organizing Condor	144
6.2.2.3	Putting It All Together	145
6.2.3	Sandboxing Resources	145
6.2.3.1	Securing the Resources	146
6.2.3.2	Respecting the Host	147
6.2.3.3	Decentralized Submission of Jobs	147
6.3	Deploying a Campus Grid	148
6.3.1	Background	148
6.3.2	Traditional Configuration of a Campus Grid	149
6.3.3	Grid Appliance in a Campus Grid	151
6.3.4	Comparing the User Experience	152
6.3.5	Quantifying the Experience	153
6.4	Lessons Learned	155
6.4.1	Deployments	155
6.4.2	Towards Unvirtualized Environments	156
6.4.3	Advantages and Challenges of the Cloud	156
6.4.4	Stacked File Systems	157
6.4.5	Priority in Owned Resources	159
6.4.6	Timing in Virtual Machines	160
6.4.7	Selecting a VPN IP Address Range	160
6.4.8	Administrator Backdoor	161
6.5	Related Work	162
7	SOCIAL PROFILE OVERLAYS	164
7.1	Related Works	165
7.2	Social Overlays	167
7.2.1	Finding Friends	168
7.2.2	Making Friends	169

7.2.3	The Profile Overlay	171
7.2.4	Event Based Message Notification	172
7.2.5	Active Peers	173
7.2.6	Groups	173
7.3	User Interaction	174
7.4	Challenges	176
8	CONCLUSIONS	179
	APPENDIX: STRUCTURED OVERLAY BROADCAST	182
	REFERENCES	184
	BIOGRAPHICAL SKETCH	193

LIST OF TABLES

<u>Table</u>	<u>page</u>
2-1 VPN Network Classifications	35
2-2 Qualitative comparison of the three deployment models	40
2-3 WAN Results for inter-cloud networking	61
2-4 LAN results performed at GoGrid	62
2-5 Virtual Network Comparison	62
3-1 Time in seconds for various private overlay operations	81
3-2 Public and research overlays	83
4-1 Pathing overheads	92
5-1 Relay comparison	120
5-2 Full tunnel evaluation	131
6-1 Grid Middleware Comparison	140

LIST OF FIGURES

<u>Figure</u>	<u>page</u>
1-1 A typical VPN client	21
1-2 1-D ring structured overlay	26
1-3 A typical NAT interaction	29
2-1 Three VN Approaches	40
2-2 The state diagram of a self-configuring VN.	42
2-3 VN Interface	43
2-4 VN Router	44
2-5 VN Hybrid	45
2-6 ARP request/reply interaction	45
2-7 DHCP client/server interaction	47
2-8 VN Router migration	52
2-9 VN Router migration evaluation	53
2-10 System transaction rate for various VPN approaches.	54
2-11 System bandwidth for various VPN approaches.	55
2-12 Grid evaluation setup	57
2-13 Grid Netperf bandwidth evaluation	58
2-14 Grid Netperf latency evaluation	58
2-15 Grid SPECjbb evaluation, Stream load	59
2-16 Grid SPECjbb evaluation, RR load	59
3-1 Bootstrapping a P2P system using an existing (generic) overlay.	66
3-2 Bootstrapping a P2P system using Brunet.	78
4-1 CDF of private overlay bootstrap time	91
4-2 Security filter	94
4-3 DTLS handshake	96
4-4 Joining a secure overlay	97

4-5	Bootstrapping a secure overlay	98
4-6	Overlay broadcast time	101
4-7	Bootstrapping a new GroupVPN	103
5-1	Motivation for relays	118
5-2	Motivation for relays	119
5-3	Latency in PlanetLab deployment	122
5-4	Drop rate in PlanetLab deployment	123
5-5	Time to form a direct connection	125
5-6	IP Broadcast using the DHT compared to overlay broadcast	126
5-7	An example of both full and split tunnel VPN modes	127
5-8	The contents of a full tunnel Ethernet packet	129
6-1	Grid Appliance middleware	134
6-2	Grid Appliance deployment scenario	142
6-3	A collection of various computing resources at a typical university.	149
6-4	Time to construct a grid	154
6-5	Time to run a job on a grid	154
6-6	Grid Appliance stackable file system	158
7-1	An example OverSoc social overlay network	166
7-2	Alice requests and receives a friendship from Bob.	168
7-3	Alice, already a friend of Bob, connects to his social overlay.	170
A-1	Tree-based overlay broadcast	182

Abstract of Dissertation Presented to the Graduate School
of the University of Florida in Partial Fulfillment of the
Requirements for the Degree of Doctor of Philosophy

DESIGN, IMPLEMENTATION, AND APPLICATIONS OF
PEER-TO-PEER VIRTUAL PRIVATE NETWORKS
FROM GRIDS TO SOCIAL NETWORKS

By

David Isaac Wolinsky

May 2011

Chair: Renato Figueiredo

Major: Electrical and Computer Engineering

Virtual private networks (VPNs) enable existing network applications to run unmodified in insecure and constrained environments by creating an isolated and secure virtual environment providing all-to-all connectivity for VPN members. While there exist both centralized and distributed VPN implementations, current approaches lack self-configuration and organization capabilities that would reduce management overheads and minimize effort by non-experts. Recent use of peer-to-peer (P2P) techniques have focused on alleviating pressure placed upon infrastructure nodes by allowing peers to form direct connections for communication purposes, while infrastructure nodes are used for handling session management and supporting indirect communication by relaying traffic when NAT (Network Address Translation) or firewall traversal fails. In terms of decentralized, P2P-based VPN solutions, the mechanisms explored thus far in related works employ unstructured P2P systems, which can have significant scalability limitations. This thesis constructs a novel decentralized P2P VPN that addresses the following core aspects that are integral to user-friendliness: bootstrapping, discovery, security, and endpoint configuration.

A resource joining a distributed system goes through a bootstrapping process. The target environment for VPNs include small systems with many if not all users behind NATs and firewalls making the bootstrapping process challenging. Centralized

systems address the bootstrapping problem by using a common resource for peer registration, discovery, and connection establishment. Centralized systems, however, come with additional costs in deploying and managing a dedicated resource with a public Internet address and the capability to handle demands placed upon it by clients. I have investigated, implemented, and evaluated decentralized means to bootstrap private P2P overlays for connectivity-constrained resources, with an approach that supports a recursive overlay organization or the use of third-party free-to-join public overlay infrastructures using technologies such as XMPP.

Bootstrapping helps establish connectivity into an overlay; however, many systems including P2P VPNs require a means for discovery specific peers. Existing VPNs either rely on large tables hosted on infrastructure nodes or overlay broadcast techniques to find a resource. As a system grows in capacity, these approaches have their limitations, especially in VPNs where all IP addresses are independent of their location inside the VPN. I have employed distributed hash tables to efficiently establish decentralized IP address allocation and discovery seamlessly providing scalability and resilience.

In a VPN, other peers are typically either trusted directly by the peer, or indirectly through a trusted third-party. While users may trust a third-party to assist them in creating network links to other peers, they do not desire to have intermediaries that are able to read or modify their IP packets. Unfortunately, most VPNs only encrypt messages on a point-to-point (PtP) basis allowing these intermediaries privileged access to their identity and their messages. In these cases, end-to-end (EtE) security relies on out-of-bound exchanges and applications. To transparently handle security at both PtP and EtE layers across a wide spectrum of communication transports, I have developed a novel security filter, which has been demonstrated to support existing Public Key Infrastructure based security systems (such as DTLS) for both PtP and EtE traffic inside connectivity-constrained environments.

While security primitives enable private and authenticated communication, the configuration and management overheads involved in establishing trust and maintaining secure connections in VPNs are a significant hindrance to usability and adoption. In my approach, all security links are established from exchanged certificates, so each peer is uniquely identifiable. My approach uniquely handles administrative and user aspects of certificates automatically through the use of online social networking features such as peer relationships and groups.

The above self-organizing mechanisms to create VPN links need to be complemented with approaches that support effective bindings to endpoints from which messages are captured/injected from/to the VPN. In a typical approach, called the interface model, each resource in the VPN has a local binding to the VPN by locally installed software. Unfortunately, this introduces significant overheads when two or more such systems are running inside the same trusted LAN. Alternatively, if all resources in a LAN connect to a common VPN, such as in a grid or for cloud computing environments, the resources can share a common entry point to the VPN through a router model. Unfortunately, existing approaches do not transparently configure the router and connected resources. Additionally, the router model does not work well on shared networks, where there are either untrusted users or some resources should not be accessible through the VPN. I have shown herein how all of these considerations can be handled without the introduction of new protocols by utilizing existing services commonly provided by network stacks, primarily DHCP and ARP, which enables a new type of VPN model that balances the benefits of the interface and router models.

The premise for this work is to enhance the usability of VPN systems enabling wider adoption by non-expert users in home, small/medium business, and education environments. The concepts for this work have been carefully designed, implemented, and evaluated and then demonstrated through the implementation of novel systems (SocialVPN, GroupVPN, and Grid Appliance) accessed by real users. The SocialVPN

creates user-centric VPNs so that peers only have VPN links with their social network friends, whereas the GroupVPN employs a group infrastructure to manage VPN members and distribute VPN configuration. A free GroupVPN bootstrapping environment relying on PlanetLab hosted resources has been available for over three years and has been accessed by over hundreds of users including several universities and commercial entities, whereas the SocialVPN has over 80 active members online at any given time. The Grid Appliance uses the GroupVPN to form ad-hoc and distributed computing pools, facilitating computer architecture research in the Archer project. The Archer project has been accessed by student at several universities and has accumulated over 500,000 CPU hours in a little less than three years. Furthermore, the Grid Appliance has been used as both a teaching tool in distributed computing classrooms as well as by external users to create their own grids. The challenges faced in these deployments have opened the door for other avenues of research into built-in self-simulation, P2P connection establishment, efficient IP broadcasting and multicasting, and decentralized establishment of Internet gateways.

CHAPTER 1 INTRODUCTION

A Virtual Private Network (VPN) provides the illusion of a local area network (LAN) spanning a wide area network (WAN) infrastructure by creating encrypted and authenticated, secure¹ communication links amongst participants. Common uses of VPNs include secure access to enterprise network resources from remote/insecure locations, connecting distributed resources from multiple sites, and establishing virtual LANs for multiplayer video games over the Internet. VPNs in this context differ from others that provide “‘emulation of a private Wide Area Network (WAN) facility using IP facilities’ (including the public Internet or private IP backbones). ” [47]. This style of VPN is used to connect large sets of machines behind routers to a virtual private WAN, whereas this dissertation focuses on the approach of connecting individual resources into a private LAN.

As a tool enabling collaborative environments, VPNs can be useful for various applications. If friends and family require computer assistance and their computer guru no longer lives nearby, the guru can remotely log into the machine using a VPN running over the Internet despite networking constraints between the two parties. When traveling abroad, a user may wish that their Internet traffic be kept private from the local network. A VPN can be used to route all Internet packets securely through the user’s home or office network, ensuring the user’s privacy. Many computer and video games have multiplayer networking components that require direct connectivity. Most of these games rely on centralized servers for bootstrapping limiting their lifespan. Players of these games can continue playing through VPNs. Small and medium businesses may find VPNs useful for connecting desktops and servers across distributed sites securing traffic

¹ For the remainder of this document, unless explicitly stated otherwise, security implies encryption and mutual authentication between peers.

to enterprise networked resources. Independent organizations that each have limited resources can combine together their resources through a VPN to create a powerful computing grid.

The utility of the VPN described herein is illustrated by a collaborative grid computing project, Archer [37]. Archer consists of over 700 core resources as well as voluntary resources from the community to provide a dynamic and decentralized grid environment for computer architecture researchers to share and access compute cycles. Use of centralized systems would limit the scope of Archer and require dedicated administration, whereas existing decentralized solutions require manual configuration of links between peers, which is beyond the scope of Archer's target users. Current P2P virtual network (VN) approaches either lack scalability or proper security components to be considered VPNs; whereas my approach applies naturally to such systems.

There are various VPN architectures that attempt to deal with the challenges presented in these use cases. In some, certain VPN approaches may work, where others are not applicable, and in others scenarios, no current VPN approach is applicable. In general, successful deployment and use of VPNs face the following challenges:

- **OVERLAY CONFIGURATION.** Peers must be able to find each other in order to bootstrap into the overlay and to establish links with specific users inside the overlay.
- **CONNECTIVITY.** Network asymmetries created by firewalls, NATs, and Internet router outages motivate the use for VPNs. One approach is to route all traffic through a third party, but this incurs overheads. There also exist approaches to allow two peers behind NATs to communicate directly and falling back to a relay if the two peers are behind too restrictive networks.
- **PEER MANAGEMENT.** To ensure reliability and trust, a distributed system should employ security. Peer management involves providing and obtaining security credentials as well as preventing misbehaving peers from communicating with a user's resource or excluding them entirely from the system.
- **PRIVACY.** The original intention for VPNs was network security, thus all communication between peers is private. Many VPNs only secure traffic between hops and

are thus susceptible to man-in-the-middle attacks. Unfortunately, establishing end-to-end privacy can be challenging as it requires additional out-of-band exchanges.

- **ENDPOINT CONFIGURATION.** Applications transfer packets through a network interface. Endpoint configuration is necessary to enable sending and receiving packets from userspace through the VPN.

Collaborative environments can strongly benefit from a VPN that is both user-friendly as well as scalable. A system that is only user-friendly will initially attract interest but frustrate users in the long run, while systems lacking user-friendliness may have limited user adoption. By applying these requirements to the aforementioned challenges leads to the following goals: (1) a collaborative VPN should be easy to configure, such that users should be able to deploy and use them without being experts in operating systems (OSs) or networks; (2) the system should not require additional resources to support more users; (3) adding new users and resources should be straight-forward using approaches familiar to common users; (4) peers should be able to connect to each other directly if and when possible; and (5) all communication, not just hop-by-hop, should be secure.

While existing VPNs are able to meet some of these requirements, they are unable to meet them all. Centralized approaches (e.g. OpenVPN [120]) by their very nature require dedicated infrastructures and do not allow direct communication between peers, though when configured to do so are able to guarantee all-to-all communication regardless of NAT and firewall conditions. Peer-to-peer (P2P) based approaches (e.g. Hamachi [67], Wippien [79], Gbridge [66], PVC [85]) solve the issue of direct communication, though they are vulnerable to man-in-the-middle attacks when session management is handled by an external provider, rely on a central resource for the creation of VPN links, and require managed relays if direct peer communication across NATs and firewalls fails. Distributed approaches (e.g., ViNe [108], Violin [53], VNET [106], tinc [100]) require manual configuration of links between members of

the virtual network. Existing P2P overlay approaches lack scalability (N2N [29] and P2PVPN [46]) or are difficult to configure and lack privacy (I3 [103]).

My work culminates in a novel design and implementation of a VPN from endpoint configuration to overlay construction and organization resulting in an autonomic VPN that bridges the gap between user-managed VPNs and those hosted by third-party services. The VPN builds on top of a P2P system used to transparently handle network asymmetries and support address allocation and resolution. The P2P system organizes into a structured overlay, which supports scalable, distributed data sharing via a distributed hash table (DHT). Peers search for each other by querying the DHT and then use constructs provided by the P2P layer to form direct links or relays with remote peers. Peer management is handled through common social networking interfaces such as dedicated group infrastructures or relationships based upon XMMP or Facebook. Both the VPN and the overlay are secured by a common security filter framework, which can be decentrally located and bootstrapped through existing overlays. Finally, through various VPN models, users and system administrators can take the same VPN software and install it in various environments with minimal configuration overhead.

1.1 Virtual Private Network Basics

VPNs consist of two components: clients and servers ². Clients discover other clients by means of servers or overlays. Depending on the VPN style, clients will then either communicate with each other through servers or use them to establish direct links with each other. While setup may be different amongst the various VPNs, during run time, the environment provided by a VPN client is the same regardless of how the server or overlay is implemented.

² The definition of a server is VPN dependent, the general concept is a resource or set of resources that maintain state in the system. It might be a centralized resource, an overlay, or even a client in the case of P2P systems.




Figure 1-1. A typical VPN client. A VN device makes application interaction with the VPN transparent. Packets going to the VPN destination are sent by routing rules to the VN device interfaced by the VPN client. The VPN client sends and receives packets from other VPN participants via the hosts physical network device.

Figure 1-1 abstracts the common features of all VPNs clients, a service and a virtual network (VN) device providing communication with the VPN system and host integration, respectively. During initialization, the VPN service authenticates with the overlay by means of a centralized or distributed service, independently with each peer, or some other means; then, optionally, querying for information about the network, such as network address space, address allocations, and domain name service (DNS) servers. At which point, the VPN enables secure communication amongst participants.

Clients can authenticate with the overlay using a variety of methods. A system can be setup quickly by using null (no) authentication or a shared secret such as a key or a password. Using accounts and passwords with or without a shared secret provides individualized authentication, allowing an administrator to block all users if the shared secret is compromised or individual users who act maliciously. Using unique private keys with corresponding signed certificates provide a more secure approach, because

it eliminates the feasibility of brute force attacks. The trade-offs in the approaches come in terms of security, usability, and management. While the use of signed certificates provides better security than shared secrets, certificates require more configuration and maintenance. In a system comprising of non-experts, like many university VPNs, the usual setup uses a shared secret and individual user accounts. Secrets can be packaged with the VPN application, so long as it is distributed through secure channels such as authenticated HTTPS.

A VN device allows networking applications to communicate transparently over the VPN. The VN device provides mechanisms for injecting incoming packets into and retrieving outgoing packets from the networking stack, enabling the use of common network APIs such as Berkeley Sockets, allowing existing applications to work over the VPN without modification. While there are many different types of VN devices, TAP [63] stands out from the rest due to its open source and pervasive nature. TAP allows the creation of one or more Virtual Ethernet and / or IP devices and is available for almost all modern operating systems including Windows, Linux, Mac OS/X, BSD, and Solaris. A TAP device presents itself as a character device providing read and write operations. Incoming packets from the VPN are written to the TAP device and the networking stack in the OS delivers the packet to the appropriate socket. Outgoing packets from local sockets are read from the TAP device.

VN devices are no different than any other network device. They can be configured manually through command-line tools or OS' APIs or dynamically by the universally supported dynamic host configuration process (DHCP) [4, 31]. Upon the VN device obtaining an IP address, the system adds a new rule to the routing table that directs all packets sent to the VPN address space to be directed to the VN device. Packets read from the TAP device are encrypted and sent to the overlay via the VPN client. The overlay delivers the packet to another client or a server with a VN stack enabled. Received packets are decrypted, verified for authenticity, and then written to the VN

device. In most cases, the IP layer header remains unchanged, while VPN configuration determines how the Ethernet header is handled.

1.2 Computer Network Architectures

All models for computer communication in distributed systems fall under two categories: centralized and decentralized. Sub-classes of these categories include hybrid systems with centralized session management and decentralized communication and self-configuring, dynamic P2P systems. The architectures commonly used for implementing VPN systems are centralized organization and communication, centralized organization and decentralized communication, decentralized communication with manual organization, and decentralized communication with automatic organization.

Systems with Centralized organization and communication consist of clients and servers where all distributed peers are clients discovering and connecting, or organizing, through a dedicated centralized resource. Clients never communicate with each other directly, but rather every message between two clients must traverse the server. For instance, most online social networks (OSNs) are representative of these type of systems. Users of OSNs like Facebook [36] and MySpace [74] communicate through centralized environments, never directly to each other's computers. OpenVPN [120] represents this VPN approach. These systems rely on dedicated resources. In the situation that a server goes offline or becomes overwhelmed by the clients, the system is rendered useless.

Centralized organization and decentralized communication systems include the first set of popular P2P systems, such as the original Napster, Kazaa, and VPNs like Hamachi [67]. Similar to the client-server model, clients connect to a server to find other clients, though instead of communicating through the server, the clients form direct connections with each other. These approaches are limited by network address translation (NAT) and firewalls that may prevent peers from communicating with each other. In these cases, the central server may act as a relay allowing the two clients

to communicate through it. Unlike systems using centralized communication, these systems are less susceptible to being overwhelmed by client traffic and even if the server goes offline existing client links remain active, though new connections cannot be established.

Systems employing decentralized communication with manual organization address the issues of a central system going offline, because clients are configured to connect to any number of distributed servers forming an overlay. In these systems, servers are explicitly configured to communicate with other servers. Though this approach improves upon the performance and availability issues inherent to completely centralized architectures, if a server goes offline any systems communicating through it will no longer be connected to the rest of the system until the administrator creates additional links or the server becomes active again. Clients in these systems do not typically form direct links with each other; rather, they route packets through the overlay. This approach has been used to create scalable VPNs, like ViNe [108], VNET [106], Violin [53], and Layer 2 Tunneling Protocol based VPNs [107].

In automatic organization-based decentralized communication systems, there is no distinction amongst peers as they act as both client and servers, i.e., a P2P system or overlay. P2P systems are usually distributed with a list of common peers. A peer attempting to bootstrap into the P2P overlay randomly selects peers on this list until it is able to connect with one. This connection is then used to form connections with other peers currently in the overlay. The overlay can be organized in two different forms: randomly or deterministically creating unstructured or structured overlays, respectively. In an unstructured overlay, links are formed arbitrarily, thus a peer searches for another peer by broadcasting the message or using stochastic techniques. In structured overlays, peers organize into topologies by deterministically forming connections with peers nearby in the overlay address space creating structures such as ring and hypercubes. Peers can be found deterministically using greedy routing approaches

in usually $\log(N)$ time. Gnutella [87] file sharing system and Skype [99] are popular examples of unstructured systems, while P2PSIP [14] and distributed hash tables (DHTs) [104] are popular in structured systems. A challenge in unstructured systems is finding data objects in reasonable amount of time, while structured systems suffer when large amount of peers join or leave the system, known as churn [86]. In general, both approaches are difficult to secure depending on the nature of the application and deployment. When used in private environments though, they have been shown to be very useful, exemplified by Dynamo [28] or BigTable [22].

This dissertation uses structured overlays as the foundation in building scalable, decentralized VPNs, the following section reviews structured overlays.

1.3 Structured Overlays

Structured P2P overlays provide distributed querying systems with guaranteed search time. Unlike unstructured systems [19], which rely on global knowledge/broadcast or stochastic techniques such as random walks that take $O(N)$ time to guarantee finding data in the overlay, structured overlays organize into well-defined geometries with support to resolve queries within $O(\log(N))$. There exists a plethora of structured systems found both in research and in available applications [13, 71, 72, 82, 94, 104]. In order to obtain guaranteed search time, structured systems self-organizing into well defined topologies, such as a ring (pictured in Figure 1-2) or a hypercube. Peers joining an overlay typically follow these abstracted steps:

1. generate or obtain a unique identification number (node ID) within the overlay's address space, usually on the order of 128-bits to 256-bits;
2. attempt to connect to one or more random addresses from a pre-shared list of well-known endpoints, dedicates resources from a service provider or users with high uptime;
3. become connected to at least one peer in this list (leaf connection, bootstrap peer);
4. find the set of peers in the address space closest to the node's ID;




Figure 1-2. 1-D ring structured overlay

5. establish connections or exchange connection information with those peers (neighbor or near connections);
6. and finally connect to other nodes in the overlay outside the set of near connections to enable quickly traversing the address space (shortcut or far connections).

All nodes are required to have a unique node ID. Address collisions can cause inconsistencies in the overlay, where one or both of the nodes will not be able to properly connect to the overlay. Furthermore, having uniformly distributed node IDs enhances the utility of the shortcut connections. To obtain a good distribution of node IDs, either a central server can provide the ID or each node independently of others can use a

cryptographically strong random number generator. The former approach can be used to create a trusted overlay by having the third-party sign each node IDs [20].

In a ring, each node must be connected to closest neighbors in the node ID address space, that is the node immediately before and after it. Optimizations for fault tolerance suggest that for ring topologies the amount should be at least 2 and up to $\log(N)$ on both sides. Consider the case when there is overlay disconnectivity potentially due to churn; a peer receives a packet but cannot route it closer to the destination than itself because it does not have a connection with that peer. The message may either be locally consumed or thrown away never arriving at its intended destination. Increasing the number of near neighbor peers reduces the likelihood chances of packets being lost due to churn, especially if peers leave suddenly without warning.

As mentioned, shortcuts or far connections enable efficient routing in ring-based and similarly designed structured overlays. The various shortcut selection methods include: maintaining large tables without using connections and only verifying usability when routing messages [72, 94], maintaining a connection with a peer at specific locations in the P2P address space [104], or using locations drawn from a harmonic distribution in the node address space [71].

Structured overlays support decentralized query systems that can be used to build distributed data structures such as a distributed hash table (DHT) by mapping keys via a hash function to P2P IDs in an overlay. The data associated with the key is then stored at the node closest to the P2P id of the key and for fault tolerance can be stored by other nodes nearby or more keys can be generated by recursively hashing the original key. Using the DHT primitives, Past [95] and Kosha [17] projects have designed more complex distributed data stores.

The actual mechanism for querying nodes or routing in a P2P overlay can be either iterative or recursive. In iterative routing, the querying node iteratively contacts nodes closer and closer to the address until finding the closest node at which point it makes

the request directly to that node. In more detail, the querying node directly queries the node closest to the destination, that node returns back one or more network (IP) and P2P addresses of closer peers, the querying node queries these peers, and the process continues until determining there exists no closer node. Alternatively in recursive routing, a querying peer sends the message to the peer closest to the destination from its perspective, repeating the process until the message has arrived at the closest peer to the address or the destination. Compared to recursive routing, iterative can be implemented more easily though with considerable overhead as each overlay query will cause $\log(N)$ connections to form. NATs further complicate the use of iterative routing as peers attempting to connect with another peer behind a NAT will need the assistance of a third-party, whereas recursive routing maintains active connections and messages, seamlessly traversing NAT links and non-NAT links since the connections are established prior to message transmission.

1.4 Network Asymmetries

Naive P2P systems assume network symmetry, that is any peer can communicate directly with any other peer using the underlying infrastructure. Unless the software is run inside a LAN or an environment where the network topology is well controlled and defined, symmetry cannot be guaranteed. P2P used in wide area systems often relies on the Internet. Besides the potential routing outages on the Internet, significant amount of resources which are not directly accessible are connected to it. The issue is only further pressed by the current means of connecting to the Internet: Internet Protocol (IP) version 4 (IPv4) with its limited address space of only 2^{32} (approximately 4 billion). With the Earth's population at over 6.8 billion and each individual potentially having multiple Internet-capable devices, these limitations become more apparent.

Currently the two approaches addressing IPv4 limitations are: the use of NATs to enable many machines and devices to share a single IP address but preventing bidirectional connection initiation and IPv6 which supports 2^{128} addresses. The use




Figure 1-3. A typical NAT interaction. The peer behind a NAT has a private address. When the packet is sent through the NAT, the NAT translates the source information into a public mapping, keeping the original source information so that if a packet from the remote peer comes back, it can be translated and delivered to the original source.

of NATs, as shown in Figure 1-3, complicates the bootstrapping of P2P systems as it prevents peers from simply exchanging addresses with each other to form connections, as the addresses may not be public. In addition, firewalls may prevent peers from receiving incoming connections. Thus, while the eventual widespread use IPv6 may eliminate the need for address translation, it does not deal with the issue of firewalls preventing P2P applications from communicating as well as routing outages, and it is not clear that IPv6 users will not continue to rely on NAT/firewall devices to provide a well-defined boundary of isolation for their local networks.

When a machine, A , behind a typical NAT, B , sends out a packet to an Internet host, C , the NAT device translates the packet so that it appears it is coming from the NAT device making the NAT device a gateway. When the packet is sent from A to C , the source and destination are listed as $IP : port$ pairs, where the source and destination are $IP_A : Port_A$ and $IP_C : Port_C$, respectively. A forwards the packet to B who transforms the source from $IP_A : Port_A$ to $IP_B : Port_B$, where $Port_A$ may or may not be equal to $Port_B$. This creates a NAT mapping so that incoming packets from $IP_C : Port_C$ to $IP_B : Port_B$ are translated and forwarded to $IP_A : Port_A$.

There are a handful of recognized NAT devices as presented in [93, 101]. The following list focuses on the more prevalent types:

- **FULL CONE.** All requests from the same internal IP and port are mapped to a static external IP and port, thus any external host can communicate with the internal host once a mapping has been made.
- **RESTRICTED CONE.** Like a full cone, but it requires that the internal host has sent a message to the external host before the NAT will pass the packets.
- **PORT RESTRICTED CONE.** Like a restricted cone, but it requires that the internal host has sent the packet to the external hosts specific port, before the NAT will pass packets.
- **SYMMETRIC.** Each source and destination pair have no relation, thus only a machine receiving a message from an internal host can send a message back.

Of the various scenarios involving peers and NATs, so long as one peer is on any of the cone NATs and there are no firewalls, it can receive incoming connection requests. Challenges to this approach exist when firewalls are introduced or both peers are behind symmetric NATs. Firewalls may traffic that would otherwise allow NAT traversal, whereas symmetric NATs require complex mechanisms in an attempt to have incoming connection requests. These types of systems typically rely on a third-party to pass messages between the peers.

1.5 Contributions

The resulting expectations of a collaborative environment that addresses the challenges listed in the introduction are self-configuring environments enabling even non-experts to setup, deploy, and manage their own VPNs; peers should communicate with each other directly when possible or through efficient indirect paths when constrained; and the system should be reliable and ensure the privacy of its users. To address these requirements, I propose a novel GroupVPN using structured overlays consisting of the following novel contributions:

SECURE OVERLAYS. Typical overlays are secured using heuristics that limit the effects of malicious users. Challenges of using secure sessions for instituting trust or security into an overlay depends on the communication path ways. If the goal for the system is to support asymmetries on the network, then the system will have to make

significant use of datagram technologies. This work proposes a unique filter mechanism to support encrypting any form of communication between two parties and examines the overheads of deploying it in simulated and real environments.

BOOTSTRAPPING AD-HOC, DECENTRALIZED SYSTEMS. Secure overlays present a challenge when there is a one to one mapping between overlay and VPN in order to securely isolate a VPN. This stems from the fact that at any given time, peers may or may not be connected to the overlay. When used in small groups, most or all members may be behind NATs or remain online for short periods of time, creating a situation where not a single user on a publicly addressable resource will be online, limiting the use of private overlays. To address this issue, I propose the reusing of public free-to-join overlays to bootstrap into a private overlay. Peers use the public overlay to find each other and exchange connection information using secure messages. Only peers with appropriate security credentials are able to join the private overlay.

DECENTRALIZED RELAYS. In collaborative environments, most peers are behind NATs and potentially firewalls as well. While in general most NATs are traversable through existing approaches, not all are. Firewalls only complicate the matter. While these peers may be able to communicate through the overlay, as the overlay grows, this latency can become a hindrance to usability and interactivity. To improve this situation, I propose the creation of autonomic 2-hop relays between the peers.

USING SOCIAL INFRASTRUCTURES FOR MANAGEMENT AND DISTRIBUTION OF SECURITY CREDENTIALS. In order to simplify the management and access to a VPN, this component explores the use of social networks in terms of both groups and peers to facilitate trust establishment for a VPN. Beyond the contribution of uniquely using social networks to establish VPN trust, this work shows how systems can leverage trust in an existing environments for use in another.

SELF-CONFIGURING VPN ARCHITECTURES. Many existing VPN approaches require the users to setup their environment and do not provide a plug and play

system. In addition, different environments call for different types of VPNs, explicitly, individual users connect via their own VPN connections, while clusters may benefit from a shared VPN or may desire fault tolerance of having many but do not want the communication overhead when talking to VPN peers on the LAN. I address this issue with a self-configuring VPN approach that can be applied to various local environments scaling from a single computer to many.

P2P VPN ENABLED INTERNET TRAFFIC TUNNELING. When in insecure environments such as browsing private information in a coffee shop, users may desire to prevent local users and administrators from sniffing their traffic. Traditional VPNs support this behavior, but the approach is difficult to implement in P2P systems due to their dynamic nature. Currently, no decentralized VPN supports the ability to perform this behavior. I propose a method that not only works for decentralized and P2P systems but ensures a greater level of security than existing approaches by securing other non-VPN communication between the peer and gateway resources.

APPLICATIONS IN AD-HOC, DISTRIBUTED SYSTEMS. The value in a complex system like the one proposed herein can be realized when tied together for the creation of ad-hoc, distributed systems. The type of system focused on in this dissertation is a grid. While there are many grid topologies, approaches that share resources amongst users and even most that are used by a single user require a user with expertise in operating systems, networks, and middleware. This dissertation shows the applicability of P2P VPN methods and techniques that can be used to create a trusted, ad-hoc, distributed grid that requires little if any expertise in the underlying technology being utilized.

DECENTRALIZED SOCIAL NETWORKS. Traditional approaches to social networks, such as Facebook and MySpace, requires trust in a third-party entity. These third-parties mine users information for advertisements, potentially violating user's privacy. This dissertation presents a decentralized social network that addresses real

problems by taking advantage of the P2P system described herein by providing each user in a social network their own private overlay whose members constitute the friends of that individual.

IMPROVED MODELS FOR DIRECT CONNECTION ESTABLISHMENT. Originally, direct links in the P2P VN were based upon packet flow passing a threshold. Through the use of profiling real systems and published results of Internet behavior, I have concluded that this model does not scale well and have design and implemented a model that satisfactorily solves this problem.

The rest of this dissertation is organized as follows. Chapter 2 overviews existing VN and VPN approaches and discusses configuration and organization of the VPN including end-point configuration. In Chapter 3, I review the challenges to bootstrapping overlays and present my solution that reuses existing overlays to bootstrap smaller, ad-hoc overlays. This leads into Chapter 4, which discusses security issues in structured overlays and addresses the means to boot private and secure VPNs. Chapter 5 covers extensions to the VPN based upon practical demands and experiences. Chapter 6 describes the Grid Appliance, the target application for my research. Chapter 7 presents a proposed idea on how to use the technology discussed thus far to create a decentralized online social network. Finally, I conclude in Chapter 8 by discussing the value in my contributions and challenges that were revealed but not addressed in this body of work, thus motivating future work.

CHAPTER 2

VIRTUAL NETWORK CONFIGURATION AND ORGANIZATION

VPNs enable seamless communication in distributed computing particularly when combining large sets of remote resources or connecting to centralized or personnel resources. Similar use cases can be extrapolated onto other collaborative environments such as multiplayer games, merging home networks over a VPN, or accessing a work computer remotely. Each application has different requirements and in review of related research [29, 41, 46, 53, 55, 57, 66, 67, 79, 85, 100, 106–108, 120] not a single approach efficiently supports these dynamic environments. Certainly, ISP large scale VPNs such as Multiprotocol Label Switching [88] (MPLS) do not as well due to the manual configuration and expertise required. An overview of these and the one described herein are presented in Table 2-5

The organization of a VPN has a direct effect on the amount of user effort required to connect multiple sites. In this regard there are two components of a VPN, the local organization and the remote or network organization. The setup of the virtual network in order to have be a destination and recognized source for remote packets constitute the local organization, whereas the routing of the packets amongst peers is handled by the network organization. Prior research works primarily focused on the latter issue, while ignoring the former. This left users to setup their own address allocations either through manually configuring each environment or dealing with the problems caused by DHCP servers in cross domain network construction, as well as their own security distribution systems. In addition, organizing a network can be an even more complicated task than locally configuring the network, because it may require the cooperation of many administrators at the various sites. This chapter presents a novel approach to VPNs that achieves both local and network self-configuration.

Table 2-1. VPN Network Classifications

Type	Description
Centralized	Clients communicate through one or more servers which are statically configured
Centralized Servers / P2P Clients	Servers provide authentication, session management, and optionally relay traffic; peers may communicate directly with each other via P2P links if NAT traversal succeeds
Decentralized Servers and Clients	No distinction between clients and servers; each member in the system authenticates directly with each other; links between members must be explicitly defined
Unstructured P2P	No distinction between clients and servers; members either know the entire network or use broadcast to discover routes between each other
Structured P2P	No distinction between clients and servers; members are usually within $O(\log N)$ hops of each other via a greedy routing algorithm; use distributed data store for discovery

2.1 Network Configuration

The key to communicating in a VPN is creating links to the VPN and finding the peer in the VPN. The different architectures for VPN link creation are based on the methods described in Table 2-1. These approaches are described in more detail below.

2.1.1 Centralized VPN Systems

OpenVPN is an open and well-documented platform for deploying centralized VPNs. In this dissertation, it is used as the basis for understanding centralized VPNs as it represents features common to most centralized VPNs.

In centralized VPN systems, clients forward all VPN related packets to the server. Client responsibilities are limited to configuring the VN device and authenticating with the VPN server, whereas the servers are responsible for authentication and routing between clients and providing access to the servers' local resources and the Internet (full tunnel). Likewise, broadcast and multicast packets also must pass through the central server.

Centralized VPNs can support multiple servers: upon starting, the client can randomly select from a list of known servers, implementing a simple load balance. Once connected, the servers provide the client an IP address in the VPN address space. Depending on configuration, this allows a client to communicate with other clients, resources on the server's network, or Internet hosts via the VPN. Servers require additional configuration to communicate with each other.

All inter-client communication flows through a central server. By default, a client encrypts a packet and sends it to the server. Upon receiving the packet, the server decrypts it, determines where to relay it, encrypts it, and then sends the packet to its destination. This model allows a server to eavesdrop on communication. While a second layer of encryption is possible through a shared secret, it requires out-of-band communication and increases the computing overhead on communication.

2.1.2 Centralized P2P VPN Systems

Hamachi [67] is the first well-known centralized VPN that used the ambiguous moniker “P2P VPN”. In reality, these systems are better classified as centralized VPN servers with P2P links. Similar VPNs include Wippien [79], Gbridge [66], PVC [85], and P2PVPN¹ [46]. The P2P in these systems is limited to direct connectivity between clients orchestrated through a central server: in Wippien it is a chat server, while P2PVPN uses a BitTorrent tracker. If NAT traversal or firewalls prevent direct connectivity, the central server can act as a relay. Each approach uses their own security protocols with most using a server to verify the authenticity and setup secure connections between clients. In regards to the P2PVPN, long term goals involve the creation of an unstructured, which would provide a method of decentralized organization.

¹ Due to the similarities between the name P2PVPN and focus of this dissertation, “P2PVPN” refers to [46] and “P2P VPN” to the use of P2P in VPNs.

2.1.3 Decentralized VPN Systems

Some examples of systems that assist in distributing load in VPN systems are tinc [100], CloudVPN [34], ViNe [108], VNET [106], and Violin [53]. These systems are not autonomic and require explicit specification of links between resources. This means that, like OpenVPN, these systems can suffer VPN outages when nodes go offline, thus administrators must maintain the VPN connection table. Unlike OpenVPN, these approaches typically do not require all-to-all direct connectivity for all-to-all communication. Users can either setup out-of-band NAT traversal or route through relays. Links are manually configured.

2.1.4 Unstructured P2P VPN Systems

Unlike centralized and decentralized systems, P2P environments require the user to connect to the overlay, which then automatically configures links. The simplest form of overlays are unstructured, where peers form random connections with each other and use broadcast and stochastic (e.g. random walks) techniques to find information and other peers; however, due to its unstructured nature, the system cannot guarantee distance and routability between peers. The only example of an unstructured VPN is N2N [29]. In N2N, peers first connect to a super node and then, to find another peer, they broadcast discovery messages to the entire overlay. In the case that peers cannot form direct connection, peers can route to each other over the N2N overlay. In the realm of VPNs, all client VPNs are also servers performing authentication though neither approach deals with decentralized address allocation.

2.1.5 Structured P2P VPN Systems

To address the scalability concerns in unstructured systems, this work uses structured P2P overlays. As described in the first chapter, structured P2P overlays provide distributed look up services with guaranteed search time with in $O(\log(N))$ time in contrast to unstructured systems with $O(N)$ time. In general, structured systems are

able to make these guarantees by self-organizing a structured topology, such as a 1-D ring or a hypercube, deterministically by randomly generated node identifiers.

The primary feature used by structured overlays is a distributed data store known as a distribute hash table (DHT), which stores key, value pairs. In the overlay, the key is an overlay address, where the value is stored. The peer closest to the key's overlay address is responsible for maintaining the value. Cryptographic hashes like SHA and MD5 can be used to obtain the key's overlay address from a string or some other byte array.

In [44, 103], a method for address allocation is described by using the DHT. Each VPN has a unique name or namespace, when a peer requests an IP address, a mapping of $hash(namespace, IP)$ to the peers overlay address is atomically written to the DHT. A success implies that the writer was the first writer to that value and other peers reading that value will be able to identify that peer as owner of that IP address in that namespace. Likewise when a peer wants to route a packet to a remote VPN peer, they query the DHT using the mapping, which returns the overlay address. The IP packet is then sent to the overlay destination.

Unicast messages are sent between two end points on the overlay using normal overlay routing mechanisms. Direct overlay links can be used to improve performance between end points. [41] describes a method by which peers can form autonomic direct connections with each other using an unstructured overlay. As IP traffic increases over a period of time, a direct connection to bypass the overlay is initiated by the receiver of the packets. Alternatively, a VPN may wish to form all-to-all connections with VPN peers as described in [38].

To support broadcast and multicast in an overlay, all members of a subnet associate through the DHT by placing their overlay address at a specific key, i.e., *namespace : broadcast*. Then when such a packet is received, it is sent to all addresses associated

with that key. It is up to the VN at each site to filter the packet. This is sufficient to support deployments where multicast or broadcast is not relied upon extensively.

2.2 Local Configuration

At first order, there are two approaches to local VPN configuration: a single VN endpoint per a host (Interface) and a VN router endpoint for many hosts on the same LAN (Router). The components differing between the two approaches are:

- **SOFTWARE LOCATION.** Interfaces execute the software on each VPN connected resource, whereas any machine connected to the same LAN as a Router will be able to access the VPN. The Router requires a dedicated resource.
- **NETWORK CONFIGURATION.** Since the Interface software runs on each machine, it is able to directly configure networking parameters, whereas a Router must use external methods to configure the resources.
- **COMMUNICATION ON A LAN.** When two peers on a LAN using a VPN Interface to communicate, all traffic must pass through the VPN adding unnecessary overhead, though in a Router the two peers have a merged physical and virtual network between them and the traffic is able to bypass the VPN.
- **FAULT TOLERANCE.** The Router only has a single instance running, when it goes offline, all resources will lose their VPN access, whereas each individual resource has their own Interface and is responsible for their own VPN connectivity.
- **COMMUNICATION OVER THE WAN.** Performing encryption can be expensive and may limit the bandwidth available due to CPU constraints. A Router may struggle to use all the available bandwidth, whereas enough Interfaces will eventually be able to use all the bandwidth. Although each additional VPN Interface also has idle traffic, potentially reducing usable bandwidth.

This dissertation identifies methods by which a single software stack can be implemented to support self-configuration and resource migration in a way that is platform independent. This method lends itself to a new architecture known as Hybrid, allowing an instance to be run on each VPN resource but enabling direct communication amongst peers on a LAN as described in [116]. The architectures are shown communicating via an overlay in Figure 2-1 and compared in Table 2-2. The two aspects that need configuration in the local configuration beyond the VPN architecture are address allocation, obtaining and setting an IP address on a resource,

Table 2-2. Qualitative comparison of the three deployment models

	Interface	Router	Hybrid
Host LAN	No assumption	Ideally, VLAN	No assumption, though may have duplicate address allocation in the same subnet for different namespaces. ²
Host software	IPOP, tap	End node: none. Router: IPOP, tap, bridge	IPOP, tap, VETH, bridge
Host overhead	CPU, memory	End node: none. Router: CPU, memory	CPU, memory
LAN traffic	Through IPOP	Bypasses IPOP	
Migration	Handled by node	Involves source and target routers	Handled by node
Tolerance to faults	Nodes are independent	Router fault affects all LAN nodes	Nodes are independent




Figure 2-1. Illustration of the three different deployment models considered in this dissertation. In VN interface mode (1), each node has an overlay ID and communicates to all other nodes through VN tunneling. In VN router mode (2), only the router has an overlay ID and routes for a set of resources; LAN communication does not require VN tunneling. In hybrid mode (3), each host has an overlay ID; LAN communication does not require VN tunneling.

and address resolution, determining where to route a VPN packet. The keys to creating this environment involve the use of standard network protocols implemented uniformly across operating systems, including DHCP (dynamic host configuration protocol) and ARP (address resolution protocol). Many applications make use of names instead of IP addresses to resolve peers, as such a naming system, like DNS (domain name service) is almost as important as address resolution and allocation. A state machine representation of this architecture is shown in Figure 2-2.

2.2.1 Local VPN Architecture

As described in the introduction, the TAP device is the glue by which the local resources communicate with the VPN. Each approach relies on the TAP device though in different configurations. In the Interface (Figure 2-3), the TAP device is used directly by the user as any other network device. In short, packets are written to the TAP device by the O/S sockets and read by the VPN software to send to the remote location, packets received by the VPN are written to the TAP device and delivered to sockets by the O/S. The Router (Figure 2-4) bridges the TAP device to a LAN, thus packets can be routed to it and sent through the VPN. TAP device virtualizes a bridge to other physical networks.

Finally, the Hybrid (Figure 2-5) like the Router connects to the LAN but only allows configuration from the local host. In Linux this is possible through the use of a VETH pseudo device that provides a virtual Ethernet pair, so that one end can be bridged with the TAP device and LAN while the other provides another interface that can be configured on the LAN, which will be used by the VPN. The reason for this lies in the nature of the state of the interfaces connected to the bridge, which go into promiscuous mode, so that all packets sent to them are forwarded on as if they are on a wire as if there were only a single network interface. In non-promiscuous mode, the network card will drop packets that are not destined for that network card. So in that case, it is not possible to assign more than one IP address to a bridge, because it and all devices

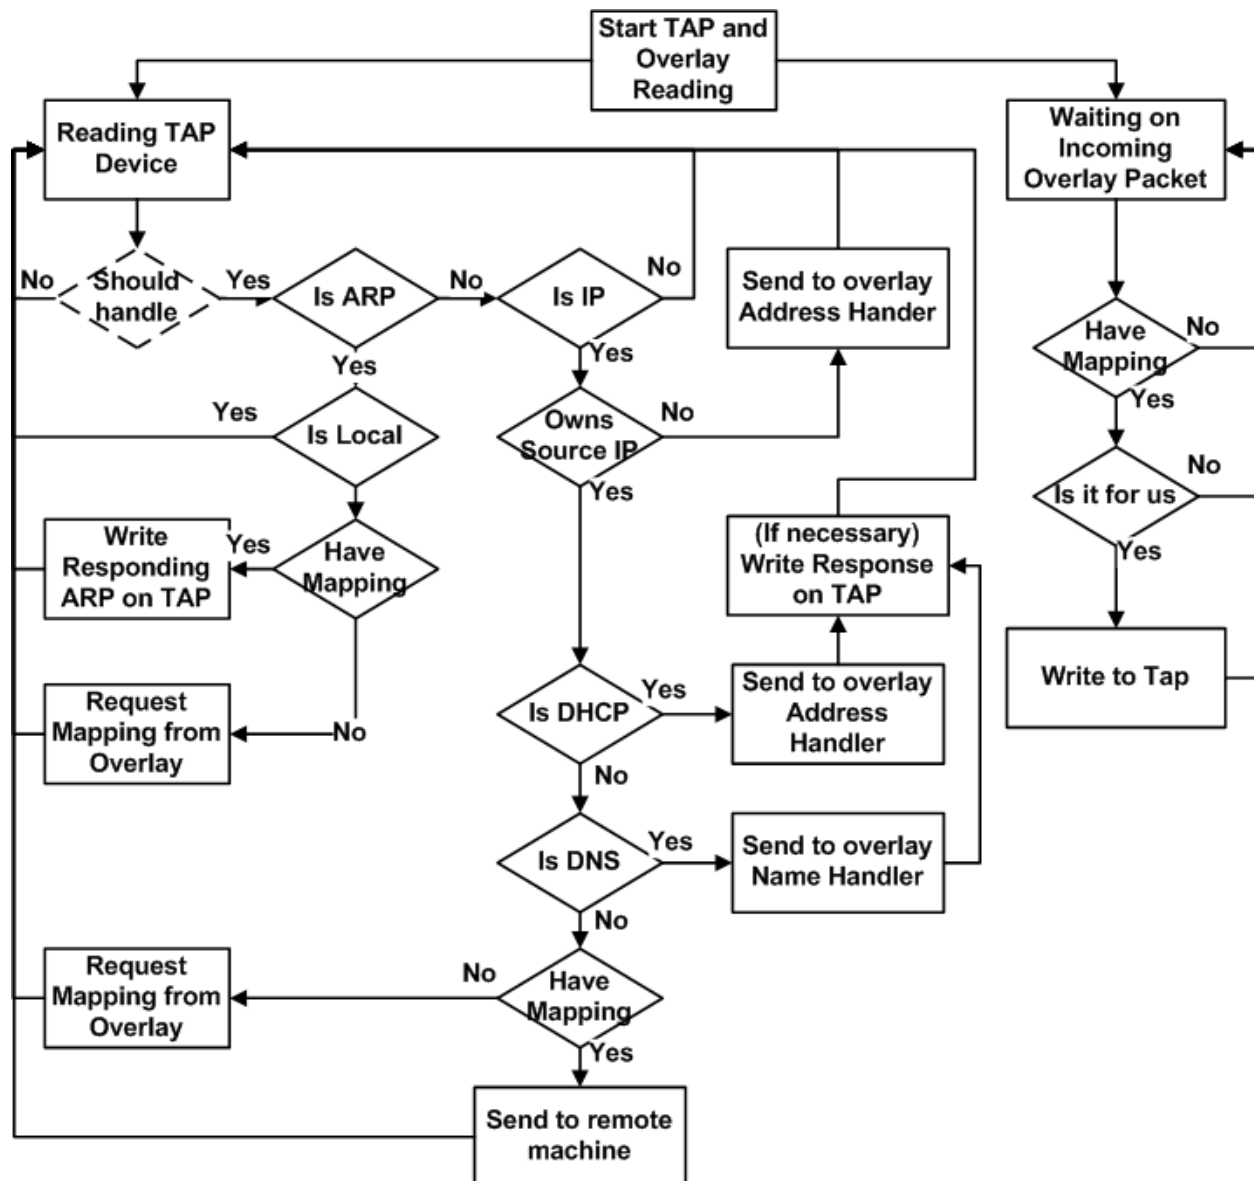


Figure 2-2. The state diagram of a self-configuring VN. In this model, a VN interface is identical to a VN router with the caveat that the TAP device is not bridged, thus isolating the VN traffic. The “Should Handle” with dashed lines is a feature that is specific to the VN hybrid; that is, a VN hybrid must be configured to communicate for a single network device.

connected to it are viewed as one big network interface. Connecting the VETH device allows an additional uniquely identifiable Ethernet addresses and thus additional IP addresses. In contrast, aliasing a Ethernet card only provide additional IP addresses




Figure 2-3. A VN deployed as an interface for single machine usage. The user of the machine is presented two interfaces on two different IP subnets. All non-VN subnet based traffic is routed normally via the default interface.

and services that rely on layer 2 networking. In this case, some services may not work, for example, DHCP does not work on aliased network cards.

2.2.2 Address Resolution

IP is a layer 3 protocol. Layer 2 devices such as switches, bridges, and hubs are not aware of IP addresses. When a system wants to send a layer 3 packet over a layer 2 network, it first uses ARP to find the layer 2 address owning the layer 3 address. This process, as shown in Figure 2-6, begins by the sending of a layer 2 broadcast message which contains an ARP request, asking all members in the LAN that the node owning the target IP address respond to the sender of the request. If a node owns the target IP address, it responds with an ARP reply, making themselves the sender and the original sender is the message recipient. The Ethernet header consists of the source address being the sender and the target being the destination. By listening to these requests, layer 2 devices such as a switch can autonomously learn the location of nodes holding




Figure 2-4. A VN deployed as a router providing virtual network access for an entire layer 2 network. Each machine in the network only has a VN-based address, though they can communicate directly with each other (and with proper routing rules and NAT setup the Internet as well). The machine hosting the VN can also have an IP address in the network by assigning one to the bridge.

Ethernet addresses are and can forward packets through appropriate ports as opposed to broadcasting or dropping them.

In a typical IP subnet, all machines talk directly with each other through switches. As such, they must learn each other's Ethernet address. The VN model used herein focus on a large, flat subnet spanning across all nodes connected to the VPN. To accomplish this, the VN provides the ability to virtualize a bridge, similar to proxy ARPs [80] used to implement a transparent subnet gateway [18]. In this scenario, the VN would need to respond to the ARP packets with a fake layer 2 address. Layer 2 devices in the system would then route all packets destined for that layer 2 address to the VN.




Figure 2-5. A VN deployed in a hybrid mode providing virtual network access for a single machine but bypassing the VN when a VN peer is local. This model is similar to having two network cards from a single machine going to one switch. The key feature is that this model allows a machine to be in multiple IP address space subnets and have layer 2 traffic as well.

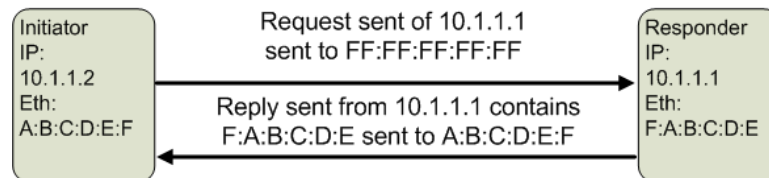


Figure 2-6. ARP request/reply interaction

As shown in the state machine (Figure 2-2), ARPs are only responded to if (a) they are inquiring about a VN IP address, (b) the VN address is not locally allocated, and (c) there is a P2P:IP mapping. If all those are true, then an ARP response is sent back to the sender. ARPs are occasionally sent out during the course of communication and thus if a machine migrates to a VN router, the VN router will no longer respond with ARPs. An ARP response sent by the VN requires a source Ethernet address, bridges and switches will see the response and will forward all traffic towards the TAP device

for that Ethernet address. A VN device can use the same Ethernet address for remote entities.

Prior to the introduction of the VN hybrid, the VNs used the Ethernet address FE:FD:00:00:00:00 to refer to remote entities. If each VN hybrid used this address, there would be layer 2 collision causing a single hybrid to have all traffic sent to it. In hybrid mode, each VN must generate a unique “remote” Ethernet address at run time. Experience and research has led to the following solution: (1) use FE:FD for the first two bytes as they tend to be unallocated and (2) assign random values to the 4 remaining bytes. Applying the birthday problem in this context, the expected probability of address collisions is small for typical LAN environments (less than 50% if the average number of VN hybrid nodes on the same L2 network is 65,000).

The key difference from the Hybrid and Router is that the Hybrid routes for only a single node, say “A”, and thus must ignore messages that do not originate from “A”. The Hybrid model does not necessarily know about the existence of all machines in a LAN, because it does not own them. So when an ARP request of some remote machine, say “B”, is sent by “A”, the Hybrid must send out a matching request with the result being sent back to the pseudo-entity of the transparent subnet gateway so that the VPN can determine if “B” exists locally. If no message is returned after a set amount of time (the reference implementation used 2 seconds), then assuming that there is a peer in the overlay with the IP address, the original ARP will be responded to with the pseudo-entity being the target.

2.2.3 Address Allocation

IP addresses are traditionally allocated in one of three ways: 1) statically, 2) dynamically through DHCP, or 3) through pseudo-random link-local addressing. This model focuses on static and dynamic addressing.

The network components configurable by DHCP as defined by [4, 31] that are interesting to a VPN are addresses, routing, and other networking related features.

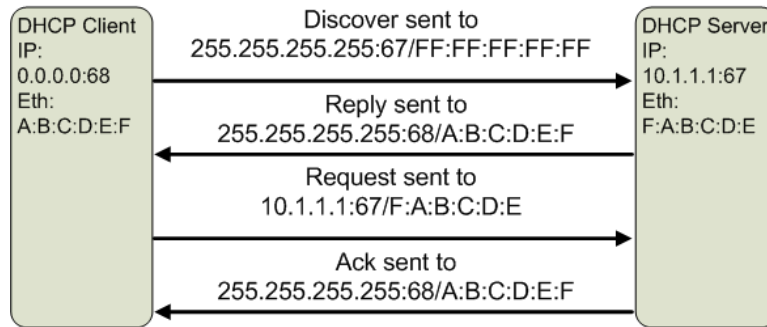


Figure 2-7. DHCP client/server interaction

While many different client and servers exist, they all tend to support the basic features of allowing the server to specify to a client an IP address, a gateway address, and domain name servers. As shown in Figure 2-7, the steps in DHCP are:

1. Client sends Discover packet requesting address.
2. Server receives the packet, allocates an address, and sends an Offer of the address and other network configuration.
3. Client receives and acknowledges the Offer by sending a Request message to accept the Offer.
4. Server receives Request message and returns an ACK message containing the same details as the Offer.

During the DHCP phase, the VPN communicates with a DHCP server for the VPN, which will allocate an address for the requester. Similarly, a VN model can review packets coming into the VPN, review the sender IP address, and request and notify the server of this allocation. Treating static addresses like DHCP enables easier configuration of the VPN, though it is difficult to handle address conflicts. In this model, this is done by the server ignoring the duplicate requests, and it is up to the user to configure for a new address. Thus DHCP provides a more reliable method in these systems.

To support scalable address allocation in decentralized systems, the DHCP server is a virtual entity, parsing DHCP packets and interacting with an overlay based DHT.

This approach does not need to be limited to structured overlay based VPNs but can be introduced as an added value component.

An important aspect of DHCP is that after a machine has received an IP address from the DHCP server, it always checks to ensure that the address has not been allocated, as such the VPN should never respond to address resolutions for local IP addresses.

If an overlay allocates an address to the VN, then the VN owns it. The other address that the VN owns is the null address, 0.0.0.0, which is sent during DHCP to indicate that the machine has no address prior to the request.

2.2.4 Domain Name Servers and Services

Name services allow machines to be addressed with names that are more meaningful to users than numeric addresses. Certain applications and services require domain name checking, such as Condor. To support DNS, this requires that the OS be programmed with the VN's DNS servers IP, typically the lowest available IP address in a subnet. In static configuration, this process requires the user to manually add this address, though through DHCP this is set automatically.

In the state representation of the VN (Figure 2-2), the VN checks the IP packet to ensure that the destination IP and port match that of the virtual DNS server and the well-known DNS port, 53. In the event of a match, the packet is passed to the VN's handler for domain names. Names are typically used for the following purposes: 1) because applications require it, and 2) to assist users in finding resources. To deal with 1), the DNS can deterministically maps IP addresses to names, such as 10.250.5.5 maps to C250005005. 2) can be solved by using the DHT and placing key:value pairs of the form hash(namespace:hostname) to IP address and hash(namespace:IP address) to hostname.

2.3 Supporting Migration

There has been a rapid increase in the deployment of Virtual Machines (VMs) for use in resource consolidation in the server industry as well as the domain of cloud computing. Providers of cloud computing services have adopted virtual machines as the unit of granularity for providing services and service level agreement to the users. Users are billed according to the number of VMs and their uptime. Major cloud-computing providers including Amazon EC2 and Go-Grid have adopted Xen as the virtualization platform for their services and sell compute resources in the form of virtual machines.

Apart from advantages like performance isolation, security, and portability, one of the significant advantages of using VMs is the capability to migrate the VM with its entire software stack from one physical host to another. This migration may be performed in a stop-restart manner, where the VM is paused, migrated to another host and restarted, or in a live mode, which attempts to minimize down time to reduce interruption of services running on the VM.

VMs including Xen [61], VMware ESX [75] and KVM [81] support migration with two critical requirements: (1) file systems (disk images) must be on a shared storage system (i.e. network file systems or storage area networks) and (2) to maintain network connectivity, the migration must occur within an IP subnet. In order to retain network connectivity after migration, the VMM must notify the LAN of the VM's new location. The new VMM host generates an unsolicited ARP reply which broadcasts to the entire network the VM's new location.

The VN Interface and Hybrid models support migration of the virtual address using techniques previously described in [41]. This is a product of the decentralized, isolated overlay approach where each overlay end point has a one-to-one mapping to VN end point, e.g., P2P to IP. When a VN Interface or Hybrid model migrates, the overlay software must reconnect to the overlay, at which point, packets will begin to be routed to the VN endpoint again, completing migration.

Unlike Interface and Hybrid models, the VN Router does not support a one-to-one mapping. In fact, a VN router tends to have one P2P address for many IP addresses. When a machine with a VN IP wants to migrate, it cannot also take its P2P address with it otherwise it would end connectivity for the rest of the members of the VN router shared overlay end point. A solution to this problems requires the ability to delete IP-to-P2P mappings in the DHT, detect new addresses on the network, and inform senders that an IP is no longer located at that overlay end point. With these capabilities, transparent migration can be achieved for the VN router model as follows.

The VMM initiates a migration on a source node. Until the migration completes, the VN router at the source continues to route virtual IP packets for the VM. Upon completion of migration, the VN router at the target learns about the presence of the migrated VM by either receipt of an unsolicited ARP or by proactively issuing periodic broadcast ICMP messages on its LAN. The VN router attempts to store (Put) the IP:P2P address mapping in the DHT, and queries for the existence of other IP:P2P mapping(s). If no previous mappings are found, the VN router assumes responsibility for the IP address. Otherwise, the VN router sends a migration request to each P2P address returned by the DHT. The VN router receiving a migration request confirms the existence of the IP address in its routing table and that if there is that there is no response to ARP requests sent to the IP address. If these conditions hold, it deletes its IP:P2P mapping from the DHT and returns true to the migration request; otherwise, it returns false. If the migration request returns true, the VN router at the target LAN starts routing for the virtual IP address; if it returns, false, the VN router does not route for the virtual IP address until the previous IP:P2P mapping expires from the DHT.

In addition to VN routers synchronizing ownership of the migrated virtual IP address, any host that is connected to that machine must be informed of the new P2P host. Over time, this will happen naturally as ARP cache entries expire and the IP:P2P mapping is looked up from the DHT. Additionally, the VN router at the source

may keep forwarding rules for the migrated IP address for a certain period of time, akin to mobile IP but not on a permanent basis. A more direct approach, as implemented in the prototype, involves the VN router notifying the connected host of a change in ownership, resulting in the host querying the DHT for the updated P2P end point. An evaluation of trade offs in the migration design, while interesting, is outside the scope of this dissertation.

A static address allocation is similar to a migration without there being an IP:P2P value in the DHT, though without querying the DHT, the situation is unclear. Systems that use DHCP only must have some method for detecting new addresses, because there is no guarantee that a DHCP will occur immediately following migration, in fact, depending on the lease time that is highly unlikely. Using an insecure DHT that supports deletes is sketchy as it would be relatively easy for machines to perform man in the middle attacks by deleting keys which they do not own. Even the use of passwords mentioned in DHT literature is not sufficient as it is not immune to collusion, or Sybil, attacks.

VN router migration was analyzed through the use of two Xen-based VMware VMs co-located on the same quad-core Xeon 5335 2 GHz machine each with 1 GB memory allocated using a minimally configured OS with a SSH server. The evaluation attempts to understand overlay overheads of the approach. The experiment, as shown in Figure 2-8, involved migrating a Xen guest VM between two Xen host VMMs running in VMware. Although they are hosted in the same infrastructure, the two domains are connected to two separate VLANs, and thus isolated. The resource information is stored in a DHT running on top of PlanetLab. Thus the migration overheads in the experiment capture the cost of wide-area messaging in a realistic environment. During the course of the experiment, over 50 different IP addresses were migrated 10 times each in an attempt to gain some insights in the cost of using the DHT with support for deletes and VN router messages as a means to implement migration. The result, presented in

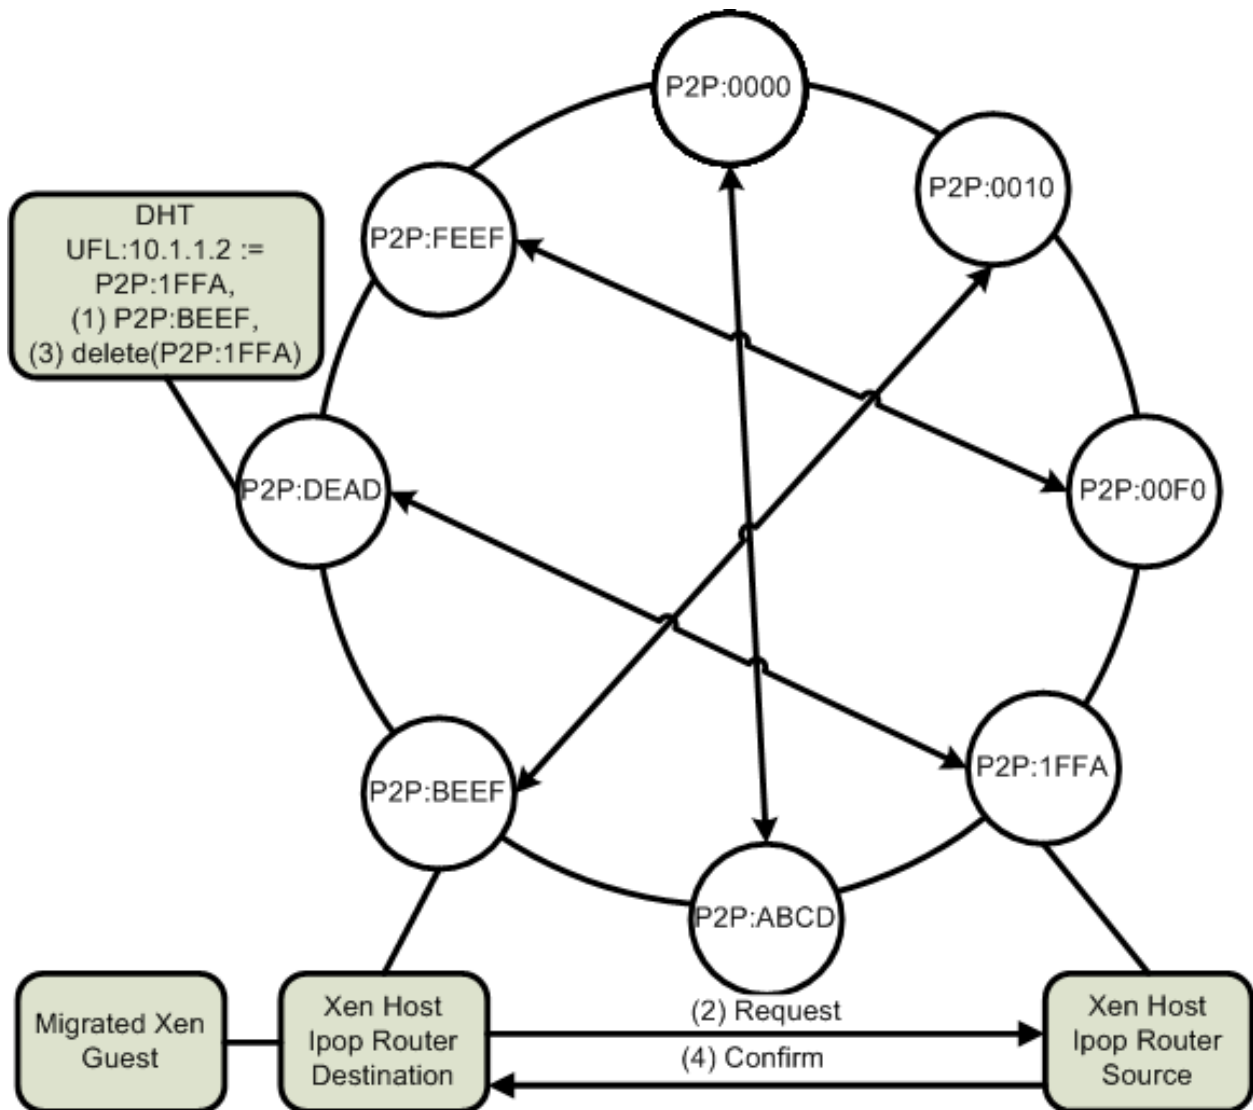


Figure 2-8. The VN operations that occur after a guest (VM) has been migrated. (1) The destination retrieves the P2P information of the source from the DHT and optimistically places its information into the DHT. (2) The destination requests that the source delete its information from the DHT. (3) The source confirms that the VM is no longer present and performs the delete. (4) The source notifies the destination that its request has finished successfully.

Figure 2-9 gathered from the experiment was how long the VN IP was offline, measured by means of ICMP ping packets. On average, the overhead of VN migration was 20 seconds. This overhead is in addition to the time taken to migrate a VM, since the VN routers begin to communicate only after migration finishes.




Figure 2-9. VN Router migration evaluation. Over 50 different IPs migrated about 10 times each. The average was 20.11 seconds with a standard deviation of 10.89. In this experiment, the majority of this time comes from VN migration, whereas VM migration requires less than a second.

2.4 Evaluation of VPN Network Configuration

This experiments explores bandwidth and latency in a distributed VPN system to motivate the usage of P2P links in a VPN. The VPNs used are include the prototype (IPOP), OpenVPN, and Hamachi. OpenVPN represents a typical centralized VPN, while Hamachi represents a well-tuned P2P-link VPN. The evaluation was performed on Amazon EC2 using small instance sized Ubuntu i386 instances to create various sized networks ranging from 1 to 32. OpenVPN uses an additional node as the central server and Hamachi has an upper bound of 16 due to limitations in the Linux version at the time of this evaluation. To perform bandwidth tests, the instances are booted and query




Figure 2-10. System transaction rate for various VPN approaches.

an NFS for the list virtual IP addresses, peers are ordered such that half the peers are act as clients and the other half the peers creating a 1 to 1 mapping between all sets. Latency and bandwidth tests are performed using netperf's request-reply and streaming tests respectively. Prior to the start of the tests, peers have no knowledge of each other, except the virtual IP addresses, thus connection startup costs are included in the test. Test are run for 10 minutes diluting the connection initiation overhead but represent an example of real usage. Results from the clients are polled at all locations and averaged together, though the OpenVPN server is measured separately. IPOP and OpenVPN use authenticated 128-bit AES, while Hamachi does not allow configuration of the security parameters and uses the default Hamachi settings.

Figure 2-10 and 2-11 present the results for latency and bandwidth respectively. Latency is measured in transactions of successful request/reply messages. In the latency test, it is obvious that having the central server increases the delay between the client and server and the results degrade more quickly as additional peers are added to the system. In small systems, OpenVPN shines probably due to optimized software, though as the system grows, the system bandwidth does not. By the time 8




Figure 2-11. System bandwidth for various VPN approaches.

peers have entered into the system, both decentralized approaches perform better than the OpenVPN solution. To summarize, decentralized VPN approaches provide better scalability, which can be immediately noticed by low latency times and, as the system grows, available bandwidth.

2.5 Evaluation of VPN Local Configuration

This section presents an evaluation of the different VN models, using prototype implementations built upon IPOP. The grid evaluation simulates a client/server environment and investigate CPU / networking overheads related with each approach. In addition a cloud deployment shows a proof of concept that connects multiple cloud and local resources as well as evaluation of overhead of the different approaches in WAN and LAN environments. In all WAN experiments, a wide-area IPOP overlay network with approximately 500 overlay nodes distributed across the world on PlanetLab resources is used to bootstrap VN connections and to support DHT-based storage and P2P messaging.

The proposed VN models place varying demands on the resources of the systems involved. The evaluation focuses on CPU as experience suggest that this is the most

significant limiting factor. As will be presented, the CPU load offered by these models depends on the bandwidth of the underlying network link, since a larger bandwidth requires more processing of packets. The tools for evaluating these VN models are Netperf and SPECjbb.

Netperf [54] is used to estimate the latency and bandwidth of the different VN models. The latency is measured by deploying Netperf in the TCP_RR mode, which measures the number of 1-byte request-receive transactions that can be completed in a second. The bandwidth is estimated by running Netperf in the TCP_STREAM mode, which is a bulk transfer mode. It should be noted that in situations where the link bandwidths were asymmetric, Netperf is deployed in both directions. Since both latency and bandwidth are dependent on the CPU comparison, evaluations that include CPU utilization tasks require creating a baseline first where only Netperf is the only active workload.

SPECjbb [102] simulates a three-tier web application with all the clients, the middle tier, and the database running on a single system in a single address space (inside a JVM). On completion, the benchmark provides the metric in terms of business of operations per second (bops). The bops score of the system under test depends on both the CPU and the memory in the system, as the entire database for the benchmark is held in memory. This benchmark generates negligible disk activity and no network activity.

2.5.1 On the Grid

The initial evaluation involves testing a client-server environment. The baseline hardware consisted of quad-core 2.3GHz 5140 Xeon with 5 GB memory and Gigabit network connectivity. Each VM was allocated 512 MB of RAM and ran Debian 4.0 using a Linux 2.6.25 kernel. The client side consisted of 4 VMs on 5 machines. The server side consisted of 5 VMs on one machine with 4 acting as servers and 1 acting as a gateway, which was necessary to control bandwidth into the system, done through




Figure 2-12. The system setup for the grid experiments. The VM “Servers” ran SPECjbb and were also the site for the collection of the netperf benchmarks. All the VM “Servers” were connected through the TC Gateway through host-only networking to the VM “Clients”. All traffic for the VM “Servers” passes through the TC Gateway, which also doubled as the Router in the Router experiments.

the Linux utility **tc** [52], traffic control. In this environment, each server had 5 clients communicating with it. The setup is shown in Figure 2-12.

The maximum bandwidth of 600 Mbps is achieved when neither virtual network nor traffic shaping are enabled (“no spec.phys” at 1000 Mbps limit in Figure 2-13), which is only 60% of the theoretical maximum. This limit is most likely the cost of VMs, specifically the time required for a packet to traverse both VMs networking stack as well as the hosts networking stack. Another observation was that transactions per second (Figure 2-16) do not improve significantly for **tc** bandwidth limit above 25 Mbps in all cases; thus focus is on only the relevant data up to this limit.

Distinguishing features of the different VN models include the following. Figure 2-13 shows that bandwidth in all VN models is comparable with traffic control limit up to 75 Mbps. Beyond this point, the interface model achieves better bandwidth than the router model (VN processing is distributed across multiple processes); the spec/no spec ratio in the router model is smaller than in the interface model because there is less resource contention caused by VN processing on end nodes. For the same reason, the router tends to achieve better SPEC results (Figure 2-15) than the interface. Figure 2-14




Figure 2-13. Netperf TCP Stream bandwidth measurements with and without SPECjbb load. Lines are of the form (no spec, spec).(phys, interface, router). Where “spec” indicates SPECjbb benchmark is active, while “no spec” indicates that SPECJbb is inactive. “phys” implies the absence of IPOP with benchmarks occurring directly over the “physical” network card. “interface” and “router” present the results for VN interface and router models respectively.




Figure 2-14. Netperf TCP RR latency measurements with and without SPECjbb load. Lines are of the form (no spec, spec).(phys, interface, router). Where “spec” indicates SPECjbb benchmark is active, while “no spec” indicates that SPECJbb is inactive. “phys” implies the absence of IPOP with benchmarks occurring directly over the “physical” network card. “interface” and “router” present the results for VN interface and router models respectively.




Figure 2-15. SPECjbb scores with and without Netperf Stream load. Lines are of the form spec.(control, phys, interface, router). “spec” implies that SPECJbb executes in all tests. In “control” Netperf is inactive, that is, it is the maximum attainable value for SPECJbb. “phys” implies the absence of IPOP with benchmarks occurring directly over the “physical” network card. “interface” and “router” present the results for VN interface and router models respectively.

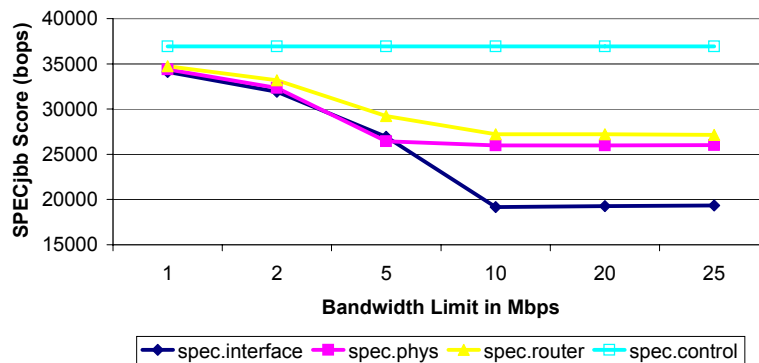


Figure 2-16. SPECjbb scores with and without Netperf RR load. Lines are of the form spec.(control, phys, interface, router). “spec” implies that SPECJbb executes in all tests. In “control” Netperf is inactive, that is, it is the maximum attainable value for SPECJbb. “phys” implies the absence of IPOP with benchmarks occurring directly over the “physical” network card. “interface” and “router” present the results for VN interface and router models respectively.

shows that the router performs poorly compared to the interface model in terms of transactions/second, though it achieves a better ratio of SPECjbb score (Figure 2-16) to transactions than the interface at constrained bandwidths (less than 5 Mbps).

The hybrid method was tested, and results were nearly identical to those of the interface, from the point of view of the WAN part of the VN, it is the same architecture.

These results are not reported in the plots as they add little value and further obfuscate the results.

The bandwidth cap observed in the router approach reflects the performance achieved by the current prototype of the router, subject to VM overheads. The use of VM is an assumption that is valid in the domain of cloud computing where all resources run in a VM. This experiment focused on the interplay between resource consumption by overlay routers and application performance. Optimized user-level overlay routers running on dedicated physical machines have been reported to achieve performance near Gbit/s in related work [109].

One thing that left unevaluated that may provide more interesting data would be providing the VN router dedicated hardware. In the test environments, this was infeasible, because all but one of the machines in the lab run VMware Server 1, which has a bug with setting the virtual network card in promiscuous mode. This effectively makes it impossible for a VM to be a VN router as no packets will ever make their way into the VM, as the VMM will reject all packets. As such, the machines hosting the servers had VMware Server 2, which does allow setting a network interface into promiscuous mode.

2.5.2 In the Clouds

The goal of this experiment is to demonstrate the feasibility of connecting multiple cloud providers as well as local resources together through virtual networking. The sites chosen for evaluation were local resources at University of Florida and cloud resources provided by Amazon EC2 and GoGrid. A qualitative observation here was that the differences in the networking infrastructure exposed by different cloud providers reinforce the importance of the virtual network to allow flexibility in how end nodes are connected. Specific network configurations for the clouds were as follows:

- Amazon EC2 provides static IP networking (public and private), no Ethernet connectivity, and no ability to reconfigure IP addresses for network. Currently, only the VN interface model is supported.

Table 2-3. WAN Results for inter-cloud networking. Stream is in Mbs and RR is in trans/s (The inverse of trans/s would be equal to the average latency).

	EC2 / UF	EC2 / GoGrid	UF / GoGrid
Stream Phys	89.21	35.93	30.17
Stream VN	75.31	19.21	25.65
RR Phys	13.35	11.09	9.97
RR VN	13.33	10.69	9.76

- GoGrid provides 3 interfaces (one public, statically configured, and two private, which can be configured in any manner); the 2 private interfaces are on separate VLANs supporting Ethernet connectivity. The VN interface, router, and hybrid models are supported.

This experiment narrows down the performance evaluation to focus on WAN and LAN performance of VNs in cloud environments and consider Netperf single client-server interactions only. Amazon only supports Interface mode, thus it is only evaluated in the WAN experiment. It has been observed that, within Amazon, the VN is able to self-organize direct overlay connections [42]. Each test was run 5 times for 30 seconds, the standard deviation for all results was less than 1. Because of this, only the average is presented in Table 2-3.

It can be seen in Table 2-3 that the VN adds little overhead in the Netperf-RR experiment. Between UF and GoGrid as well as between UF and Amazon EC2, the overhead for the Stream experiment was about 15%. This may be attributed to the additional per-packet overhead of the VN and the small MTU set for the VN interface (1200). The MTU, or maximum transmission unit, is the largest packet that is sent from an interface. IPOP conservatively limits the VN MTU to 1200 down from the default 1500 to allow for overlay headers and to work properly with poorly configured routers, which has encountered in practical deployments. A more dynamic MTU, which will improve performance, is left as future work. The EC2 / GoGrid experiment had greater overhead which could possibly be attributed to by the VM encapsulation of cloud resources.

Table 2-4. LAN results performed at GoGrid. Stream is in Mbs and RR is in trans/s. Interface and Physical used the eth0 NIC, while Router and Hybrid used eth1. Different VLANs may give different results.

	VN Interface	VN Router	VN Hybrid	Physical
Stream	109	325	324	327
RR	1863	2277	2253	3121

Table 2-4 shows that some of the performance expectations for the different models in a LAN were accurately predicted while others were not so clear. Stream results match the expectation that VN models hybrid and router bypass virtualization and get near physical speeds, whereas interface does not. Interestingly, RR had rather poor results for Router and Hybrid though further testing seems to indicate that this is an issue of using the VLAN connected network interfaces as opposed to the public network connected interface.

Table 2-5. Virtual Network Comparison

	Overlay	Routing	Configuration	Miscellaneous
IPOP	Structured P2P overlay with $O(\log N)$ routing hops, where N is the size of P2P network. Self-optimizing shortcuts and STUN-based NAT traversal.	Mapping stored in DHT resolves virtual IP address to P2P address. Virtual network packets are routed to corresponding P2P address.	Each machine runs P2P VPN software with a dynamic IP address in a common subnet. Common configuration shared amongst all hosts.	Supports encrypted P2P links and end-to-end VPN tunnels (unpublished work). Migration possible; routes self-configure without user intervention, product of the P2P overlay.
N2N	Unstructured P2P network, super nodes provide control paths, forms direct connections for data.	Broadcast for discovery and overlay for control. No organization, no guarantees about routing time.	Requires N2N software at each host, must connect to a super node. Supports layer 2 Ethernet network.	Supports shared secrets to create private tunnels between edges. Migration not discussed, but potentially free due to layer 2 approach.
OCALA	Not tied to any specific overlay, layer 3 middleware.	Based upon chosen overlay.	Requires OCALA stack, overlay configuration, and IP to overlay mapping.	Security is overlay based or SSH tunnels. Migration not mentioned.

Table 2-5. Continued

	Overlay	Routing	Configuration	Miscellaneous
SoftUDC VNET	Decentralized with explicitly configured overlay routes.	Broadcast for discovery.	Requires software on each host and one proxy per site. Layer 2 networking.	Security is not discussed nor is wide-area migration.
ViNe	ViNe authority configures global network descriptor table (GNDT) explicitly at each router. Supports proxying to one location through another and NAT traversal.	GNDT provides overlay routes for all routers in overlay.	Each subnet is allocated a single router. Each host must be configured for regular and ViNe networks, but no VN software needed on host.	Supports encrypted tunnels between ViNe routers, migration not discussed.
Violin	Decentralized network with statically configured overlay routes.	Broadcast discovery for Ethernet, static routes for IP subnet.	Virtual hosts connect VMs to the VN. Hosts connect to virtual switches or proxies (gateways). Switches connect to proxies. Sites are typically allocated an IP address space.	Security potentially through the use of SSH Tunnels. Migration possible; requires reconfiguration of switches.
Virtuoso VNET	Decentralized with explicitly configured overlay routes.	Broadcast for discovery. Bridging learns paths after initial discovery. Virtual network packets are routed between VNET proxies. Can be configured manually.	Each site runs a proxy providing Ethernet bridge to other proxies. VM hosts forward packets to local proxy. Proxies configured to connect to other proxies.	Security through the use of SSL and SSH Tunnels. Layer 3 migration, product of layer 2 virtualization.
OpenVPN	Centralized	Central server	Servers manually configured to connect with each other. Clients randomly select server from pre-shared list	All communication traverses central server, end to end traffic by default is not protected from central server
Tinc and CloudVPN	Decentralized with explicitly configured overlay routes	Broadcast for discovery, messages traverse overlay	Manual configuration	NAT traversal through relays only
Hamachi	Centralized Discovery, P2P links	Peers establish security links and end point information from a central server, attempt to form direct connections, if fails, relay through central server	Select a network to join or create and specify a password, communicates with a centralized server to manage the VPN	Lacks portability, Linux version out of date, inability to run external relay servers, UDP NAT traversal

Table 2-5. Continued

	Overlay	Routing	Configuration	Miscellaneous
GBridge	Centralized Discovery, P2P links	Peers establish security links and end point information from a central server, attempt to form direct connections, if fails, relay through central server	Select a network to join or create and specify a password, communicates with a centralized server to manage the VPN	Lacks portability and inability to run external relay servers, uses TCP NAT traversal
Wippien	Centralized Discovery, P2P links	Peers discover and authenticate each other through XMPP chat server, security provided unknown, peers attempt to form direct connections with each other, if that fails, no communication	All peers must be members of associated XMPP chat rooms and be connected to the chat	Requires a GUI, difficulty penetrating NATs, claims to be open source though most of the code is unavailable, Linux client out of date and does not support NAT traversal
P2PVPN	Centralized Discovery, P2P links	Peers discover each other through a BitTorrent tracker and attempt to form direct links with each other, attempts to form all-to-all connectivity, if direct links are unavailable, indirect links can be used to forward packets	Peers must join the same tracker and use common shared secret	Work in progress to make more unstructured, currently a cross between centralized and decentralized

CHAPTER 3

BOOTSTRAPPING PRIVATE OVERLAYS

While P2P overlays provide a scalable, resilient, and self-configuring platform for distributed applications, their adoption rate for use across the Internet has been slow outside of large-scale systems, such as data distribution and communication. General use of decentralized, P2P applications targeting homes and small/medium businesses (SMBs) has been limited in large part due to difficulty in decentralized discovery of P2P systems, the bootstrap problem, further inhibited by constrained network conditions due to firewalls and NATs (network address translators). While these environments could benefit from P2P, many of these users lack the resources or expertise necessary to bootstrap private¹ P2P overlays particularly when the membership is unsteady and distributed across wide-area network environments where a significant amount of (or all) peers may be unable to initiate direct communication with each other due to firewalls and NATs.

Examples of large-scale P2P systems include Skype, BitTorrent, and Gnutella. Skype is a voice over P2P system, whereas BitTorrent and Gnutella are used for file sharing. The bootstrapping in these systems typically relies on overlay maintainers using high availability systems for bootstrapping, bundling their connection information with the application distribution. The application then uses these servers during the initialization phase to connect with other peers in the system. Alternatively, some services constantly crawl the network and place peer lists on dedicated web sites. A new peer wishing to join the network queries the web site and then attempts to connect to the peers on this list.

¹ In the context of this chapter, private implies that the overlay's purpose is not for general use. Once established, such overlays can support privacy in communication; however, overlay security is beyond the scope of this chapter and covered in more depth in chapter 4.

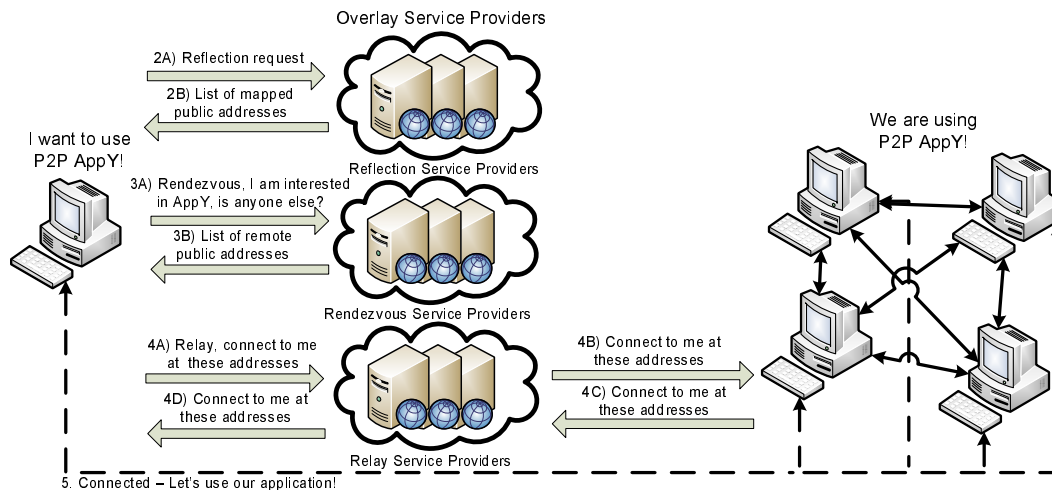


Figure 3-1. Bootstrapping a P2P system using an existing (generic) overlay.

In smaller-scale systems, P2P interests focus on decentralization. For example, users may desire to run an application at many distributed sites, but the application lacks dedicated central servers to provide discovery or rendezvous service for peers. In contrast, dedicated, centralized P2P service providers, such as LogMeIn's Hamachi, a P2P VPN, may collect usage data, which the users may wish to remain private, or are not free for use.

Many applications make sense for small-scale overlay usage, including multiplayer games, especially those that lack dedicated online services; private data sharing; and distributed file systems. Clearly, a small P2P system could be bootstrapped by one or more users of the system running on public addresses, distributing addresses out-of-band, instructing their peers to add that address to their P2P application, and then initiate bootstrapping; but these types of situations are an exception and not the norm. Ultimately, the users would be enhanced significantly through approaches that can make decentralized bootstrapping transparent through minimal and intuitive interaction with the P2P component.

The basic bootstrapping process can be broken down into two components: finding and connecting to an active peer in the system. When a node starts, it contacts various bootstrap servers, until it successfully connects with one, upon which they exchange

information. The bootstrap server may inquire into the overlay for the best set of peers for the new peer and respond with that information or it may respond with its existing neighbor set. At which point, the peer attempts to connect with those peers. This process continues aggressively until the peer arrives at a steady state, either connecting with a specific set of or a number of peers. Afterwards, the P2P logic becomes passive, only reacting to churn from new incoming or outgoing peers.

Overlay support for constrained peers, i.e., those behind NATs and restrictive firewalls, requires additional features to support all-to-all connectivity for peers in the overlay. The instantiation of P2P systems for private use could become overly burdensome, potentially relying on significant human interaction to bootstrap them, for example, by relaying connection information through phone calls and e-mail. Even if this is feasible, this sort of interaction is undesirable. P2P systems should be self-discovering, minimizing the amount of work users need to do in order to take advantage of them, a feature stressed by ad-hoc systems. In addition, these approaches may rely on centralized components; if they become unavailable, which is a possibility since most users lack the expertise in configuring highly available systems, the system will not be accessible.

To address this, I have explored the possibility of using existing public overlays as a means to bootstrap private overlays. There are many existing public overlays with high availability, such as Skype, Gnutella, XMPP (Extensible Messaging and Presence Protocol), and BitTorrent; by leveraging these systems, system integrators can easily enable users to seamlessly bootstrap their own private P2P systems. In the preceding paragraphs, I have identified the components necessary for bootstrapping a homogeneous system; in the following, I will expand them for environments to support the bootstrapping of a private overlay from a public overlay with consideration for network constrained peers. The public overlay must support the following mechanisms as illustrated in Figure 3-1:

1. REFLECTION. A method for obtaining global application and IP addresses or identifier for a peer that can be shared with others to enable direct communication.
2. RELAYING. A method for peers to exchange arbitrary data, when a direct IP link is unavailable.
3. RENDEZVOUS. A method for identifying peers interested in the same P2P service.

This work motivates from the belief that while small-scale P2P systems are attractive for decentralized systems, the overheads relating to creating and maintaining bootstrap services make them unfeasible. A public overlay can be used to transparently bootstrap a private overlay with minimal user interaction.

The requirements are presented and verified in the context of two prototype implementations: a XMPP (Jabber) [96] and Brunet [13]. XMPP-based overlays are commonly used as chat portals, such as GoogleTalk and Facebook Chat. XMPP also supports an overlay amongst servers forming through the XMPP Federation, which allows inter-domain communication amongst chat peers, so that users from various XMPP servers can communicate with each other. Brunet provides generic P2P abstractions as well as an implementation of the Symphony structured overlay. I present the architecture for these systems, the lessons learned in constructing and evaluating them, and provide an analysis of the latency to establish peer connectivity in a small-scale private Brunet overlay with NAT-constrained nodes.

The organization of this chapter follows. Section 3.1 overviews existing solutions to the bootstrapping problem, and NAT challenges in P2P systems. Section 3.2 presents a survey of overlays, applying the requirements for private overlay bootstrapping to them, and then show in detail how they can be applied to Brunet and XMPP. My implementation is described in Section 3.3. In Section 3.4, I perform a timing evaluation of bootstrapping overlays using my prototype on PlanetLab and discuss experiences in deploying the system.

3.1 Current Bootstrap Solutions

As described in the introduction, the simple case of bootstrapping is limited to one peer attempting to find an active peer in the overlay in order for itself to become a member. The large-scale providers have resources not readily available to small-scale overlays. This section reviews existing techniques and those being developed and describes their application to small-scale systems.

When using dedicated bootstrap overlays, a service provider hosts one or more bootstrap resources. Peers desiring to join the overlay query bootstrap nodes, until a successful connection is made to one. The bootstrap server will then assist in connecting the peer to other nodes in the P2P system. Bootstrap nodes are either packaged with the application at distribution time or through a meta data file, such as in BitTorrent. Drawbacks to this approach for small, ad-hoc pools include that the same server would have to be used every time to bootstrap the system, or users would have to reconfigure their software to connect to new bootstrap servers over time; at least one peer must have a publicly accessible address; and a bootstrap server can become a single point of failure.

Another commonly used approach for large-scale systems is the use of a host cache [27]. Clients post current connection information to dedicated web services, a host cache, that in turn communicate with other host caches. For small, ad-hoc networks, a host cache acts no differently than a centralized rendezvous point, requiring that at least one peer has a publicly accessible address.

“P2P VPN’s” [46] use of a BitTorrent tracker is similar to the host cache concept. The tracker hosts file meta data and peers involved in sharing. For the VPN, the peer registers a virtual file used to organize the peers, a form of rendezvous. Each peer in the VPN queries the tracker regarding the file, registers its IP address, and receives other active “sharers” IP addresses. Peers on public addresses or using UPNP are able to receive incoming connections from all other peers. The problem with this approach is

that it is heavily user-driven. A user must register with each BitTorrent tracker individually and maintain a connection with each of them, in order to handle cases where BitTorrent trackers go offline. In addition, this does not use the BitTorrent trackers in a normal fashion, so it may be banned by tracker hosts.

Research has shown that peers can use the locality properties of recent IP addresses in a large-scale P2P system to make intelligent guesses about other peers in the P2P system using an approach called random probing [25, 45]. The results show that, in a network of tens to hundreds of thousands of peers, a bootstrapping peer can find an active peer in 100 guesses to 2,000 guesses, depending on the overlay. The approach does not really apply well to small-scale systems, especially when peers are constrained by NATs and firewalls.

Rather than distribute an IP address, which points explicitly to some location in the Internet, a small P2P network can apply a name abstraction around one peer in the overlay using Dynamic DNS [62]. Peers share a DNS entry, which points to a bootstrap server. When the peers detect that the bootstrap server is offline, at random time intervals they will update the DNS entry with their own. The application of this approach is well-suited to small, ad-hoc groups, as the service could be distributed across multiple Dynamic DNS registrations. However, sharing a DNS entry requires trusting all peers in the overlay, making it easy for malicious peers to inhibit system bootstrapping. Also the approach requires that at least one peer be publicly addressable; if a non-publicly addressable peer updates the cache inadvertently, it could delay or permanently prevent peers from creating a P2P system. The results reported in [62] were simulation-based and did not determine how well a dynamic DNS handles rapid changing of name to IP mappings.

IP supports multicasting to groups interested in a common service. In the case of bootstrapping a P2P system [25, 94], all peers would be members of a specific group. When a new peer comes online, it queries the group for connection information and

connects to those that respond. The approach, by itself, requires that all peers are located in a multicast capable network, restricting this approach typically to local area networks.

A large-scale structured overlay [21, 24] could enable peers to publish their information into a dedicated location for their service or application and then query that list to obtain a list of online peers. Peers could search for other peers in their overlay and connect with them using their connection information. Since the service would be a large-scale system, it could easily be bootstrapped by a dedicated bootstrap or host caches. As it stands, the described works were position papers and the systems have not been fully fleshed out. The primary challenge in relationship to small, ad-hoc networks is that it lacks details bootstrapping of peers behind NATs into overlays as it provides only a means for rendezvous and not reflection nor relaying.

3.2 Core Requirements

As presented in the preceding sections, a solution to bootstrapping small P2P overlays must address several challenges, namely reflection, rendezvous, and relaying. This section presents a generic solution to this problem. The basis for my solution is reusing existing, free-to-join public overlay. In order to support these features the public overlay must have mechanisms for peers to obtain a public network identity (reflection); search for other peers that are bootstrapping the same P2P service (rendezvous); and send messages to peers through the overlay (relaying). These are the minimum requirements to bootstrap a decentralized, P2P system when all peers are behind NATs.

3.2.1 Reflection

Reflection provides a peer with a globally-addressable identifier for receiving incoming messages from other peers. Without reflection, peers on different networks with non-public addresses are unable to communicate directly with each other. Reflection is not limited to IP. For example, when a peer joins a service, such as a

chat application or a P2P system, the overlay provides a unique identifier, which also serves as a form of reflection.

In IP communication, reflection enables NAT traversal. The simplest method for NAT traversal relies on obtaining the public information for an existing UDP socket and then sharing that with other peers. This behavior can be supported through either local service or remote assistance. The local approach relies on having a router with a public IP address supporting either UPnP [110] or port forwarding / tracking. In many cases, UPnP is not enabled by default and in most commercial venues it will rarely be enabled. Port forwarding / tracking requires non-trivial router configuration, outside the comfort range of many individuals and is not uniform across routers. A peer using UPnP needs no further services, as UPnP enables a peer to set and obtain both public IP address and port mappings. Port forwarding and tracking mechanisms still require that the user obtains and inputs into the application their public IP address or use in-band assistance described next.

In the remotely assisted scenario, a peer first sends a message to a reflection provider, perhaps using STUN [91]. The response from the provider tells the peer from which IP address and port the message was sent. In the case of all cone NATs, this will create a binding so that the peer can then share that IP address and port with other peers behind NATs. When the two peers communicate simultaneously, all types of cone NATs can be traversed; the timing of messages needs to be carefully considered, however, since NAT mappings may change over time. So long as one peer is behind a cone NAT, NAT traversal using this mechanism is possible. The situation becomes complicated when both peers are behind symmetric NATs, or when either one of them have a firewall preventing UDP communication.

Peers behind symmetric NATs cannot easily communicate with each other, since there is no relation between remote hosts and ports and local ports. Further complicating the matter is that there are various types of symmetric NATs, having

behaviors similar to the various cone NAT types. In [89] the authors describe methods to traverse these NATs so long as there is a predictable pattern to port selection.

Unlike UDP, TCP NAT traversal is complicated by the state associated with TCP. In many systems, the socket API can be used to enable a peer to both listen for incoming connections and form outgoing connections using the same local addressing information. According to [78], this method works for various types of systems though the success rate on NATs is low, 40%. Other mechanisms rely on out-of-band communication [85], or use of complicated predictive models [11].

3.2.2 Relaying

NAT traversal services only deal with one aspect of the bootstrap problem: reflection. That is, peers are able to obtain a public address for receiving incoming connections with no means for to exchange addresses with other peers nor perform a simultaneous open to traverse restrictive NATs. To address this issue, many systems incorporate these NAT traversal libraries while using intermediaries to exchange addresses as a method of relaying. Another form of relaying exists when two peers are unable to form direct IP connections with each other and route data messages between a third-party.

The most common method for relaying in IP is the use of TURN [90], or Traversal Using Relay NAT. A peer using TURN obtains a public IP address and port that can be used as a forwarding address. When a remote peer sends to this address, the TURN server will forward the response to the peer who has been allocated that mapping. The lack of abstraction in TURN makes the system heavily centralized, making its application in small-scale systems complicated.

In overlays, peers typically have an abstracted identifier that does not associate them with a single server enabling more decentralized approaches to relaying. When a remote peer sends a message to the identifier, the overlay should translate the identifier into network level addresses and forward it to the destination. Because of this restriction,

messages sent by relaying cannot have expectations more than that of sending a packet by UDP. In other words, a packet will either be received in a reasonable amount of time or not at all. Support for reliability, streaming, and flow control, if necessary, must be provided in user-space.

Finally, the service should be asynchronous or event driven. The previous requirements would allow peers to relay through a message board or even by posting messages to a DHT. The problem with these two approaches is that peers may very well communicate for long periods of time using these services. That means the potential for posting large amounts of data to a service that will retain it and constantly querying the service to determine if an update is available. Both of these are highly undesirable and may be viewed as denial of service or spam attacks.

3.2.3 Rendezvous

A rendezvous service allows peers to discover the global identifier of peers interested in the same service. For any given overlay, a naive approach for rendezvous is the use of a broadcast query or random probing to determine if any other peers are using the same service. This approach is unreasonable, depending on the size of the bootstrap overlay compared to the destination overlay, it may be very difficult to find another peer, some or all peers may be behind NATs and unreachable without assistance, and in the worst case scenario a malicious attacker could be waiting for bootstrap requests into the system.

Rather than attempt to make a single unified rendezvous technique, each overlay style usually provide an efficient means for rendezvous, thus reducing the network and time overhead of finding another peer. For example, in the case of a DHT, peers can use a single DHT key to store multiple values, all of which would be addresses used to communicate with peers in the overlay. Alternatively, in a system like BitTorrent, peers could use the same tracker and become “seeds” to the same virtual file.

3.3 Implementations

Table 3-2 reviews various overlays, the majority of which are high availability, public, free-to-join overlays, though some research only overlays are included. From this list, I chose to extend Brunet and XMPP to support private overlay bootstrapping. Brunet provides a structured P2P infrastructure, though lacks an active, large-scale deployment outside of academic deployment (mine) due to being rooted in an academic project. XMPP, on the other hand, has support from a large contingency of private users and enables connections between friends with routing occurring across a distributed overlay.

My implementation makes heavy use of the transports incorporated into Brunet [13]. The key distinguishing feature of this library is the abstraction of sending over a communication link as it supports primitives similar to “send” and “receive” that enable the ability to create P2P communication channels over a variety of transports. In the next sections, I will describe how I extended Brunet to be self-bootstrapping as well as extensions to enable bootstrapping from XMPP.

The application of structured overlays as the basis private overlays focuses on the autonomous, self-managing property of the overlay network rather than the ability to scale to very large numbers. This has also been the motivation of related work which has employed structured overlays in systems in the order of 10s to 100s of nodes. For example, Amazon’s shopping cart runs on Dynamo [28] using a “couple of hundred of nodes” or less. Facebook provides an inbox search system using Cassandra [64] running on “600+ cores”. Structured overlays simplify organization of an overlay and provide each member a unique identifier abstracted from the underlying network. As mentioned in the cited works, they provide high availability and autonomic features that handle churn well. When used in small networks, most structured overlays (including Brunet and Pastry) in effect act as $O(1)$ systems, self-organizing links that establish all-to-all connectivity among peers. Brunet explicitly supports all-to-all connectivity, though in some cases may require constrained peers to route through relays. This can

further be ensured by setting the amount of near connections for the infrastructures, which in Brunet is configurable at run time.

3.3.1 Using Brunet

Prior to this work, Brunet bootstrapped using a recently online cache of peers and IP multicast. Brunet already supports behavior similar to STUN, such that, with every connection Brunet makes, peers inform each other of their view of the remote peers network state, a form of passive **reflection**. Peers also generate a unique 160-bit node identifier that can be used in the overlay as a directly receive packets regardless of the underlay conditions.

In a single overlay, Brunet supports **relaying** either through the overlay or pseudo direct connections called “Tunnels” [43], where peers route to each other through common neighboring connections. The relaying in this context is used either to maintain a necessary overlay connection, or to exchange intentions to connect with each other through “ConnectToMessage” messages. Thus when a peer desires a connection to another, both peers simultaneously attempt to connect to each other after exchanging endpoints discovered through reflection using the overlay relay mechanisms, dealing with the issue of more restrictive cone NATs and the case when the peer is behind a non-traversable NAT.

To support **relaying** within the scope of a private overlay, I have further extended Brunet’s transport library to support treating an existing overlay as a medium for point-to-point communication. This is called a “Subring” transport, because it supports the abstraction of multiple private sub-rings within a common large structured ring. When the private overlay transmits data across the public overlay, the private overlay packet is encapsulated (and possibly encrypted) in a packet that ensures it will be delivered to the correct private destination usually by means of greedy routing on the public overlay. In order to instruct peers to establish “Subring” links, they exchange an identifier of the form “brunet://P2P_ID”.

Peers store their “Subring” identifiers into the DHT for **rendezvous**. The DHT provides a scalable and self-maintaining mechanism for maintaining a bootstrap, so long as the DHT supports multiple values at the same key, as Brunet does. The key used for the DHT rendezvous is a hash of the services name and its version number, which I call a namespace. Peers can then query this entry in the DHT to obtain a list of peers in the private overlay. Since DHTs are soft-state, or lease systems, where data is released after a certain period of time, an online peer must actively maintain its DHT entry. In the case that a peer goes offline, the DHT will automatically remove the value after its lease has expired.

To support reflection in the private overlay, there were two potential paths. The first would have been to extend Brunet to support STUN in each of the remote servers and then have a private node query them for their public information. The problem with this approach is that it would require maintaining additional state in order to discern which of the remote peers are on public addresses and can provide STUN services.

Instead, I opted to multiplex the socket used for the public overlay as it already had gone through the process of “reflection”. The multiplexing of a single socket for multiple overlay is called “Pathing”. In this context, the public and private overlays are given a virtual transport layer that hooks into an another transport layer, thus not limited purely to socket transport layers. When peers exchange identifiers, instead of transmitting a simple identifier like “udp://192.168.1.1:15222”, the “Pathing” library extends it to “udp://192.168.1.1:15222/path”, where each path might signify a unique overlay.

The completed approach is illustrated in Figure [3-2](#). The approach of “Subring” and “Pathing” enabled the reuse of the core components of Brunet. Using “Subring” enables peers to form bootstrap connections to then exchange “ConnectToMessage” messages. If the direct connections failed, then the “Subring” connections could be used as permanent connections. The use of “Pathing” meant reuse of existing NAT traversal techniques and limited the amount of system resources required to run multiple




Figure 3-2. Bootstrapping a P2P system using Brunet.

overlays. In terms of total lines of code, these abstractions enabled a recursive overlay bootstrapping with a relatively small code footprint — less than 1000 lines of code.

3.3.2 Using XMPP

In addition to supporting recursive bootstrapping of private overlays, the techniques described above can be extended to use a different public overlay — an XMPP-based federation — to support the bootstrapping of private overlays. The key features that make XMPP attractive are the distributed nature of the federation and the openness of the protocol. As of December 2009, there are over 70 active XMPP servers in the XMPP Federation [119]. These include GoogleTalk, Jabber.org, and Live Journal Talk.

In XMPP, each user has a unique identifier of the form “username@domain”. Where the domain specifies the client’s XMPP server and the username uniquely identifies a single individual. XMPP supports concurrent instances for each user by appending a resource identifier to the user ID: “username@domain/resource”. A resource identifier can either be provided by the client or generated by the server. For users in the same domain, the server forwards the message from source to destination. When two users

are in different domains, the sender's server forwards the message to the receiver's server, who then relays it to the receiver.

XMPP allows for sending arbitrary binary messages called "IQ". While peer relationships are maintained by the server, they are initiated between peers using "IQ". Once peers have established a connection or subscription, they are informed through a "Presence" notification that the peer has come online, this include the full user identifier.

The first form of **reflection** in XMPP is the unique client identifier. Another is an IP reflection service available from some XMPP service providers called "Jingle" [69]. "Jingle" uses "IQ" to determine available STUN and TURN servers. Fortunately, these services are provided free of charge through GoogleTalk. In Brunet, I extended the UDP transport to support querying STUN servers to obtain and maintain and open an address mapping. STUN packets are easily distinguished from other packes as the first two bits are set to 0 as well as a static cookie found in all messages.

In order to support the situation where two peers are unable to communicate through the exchanged addresses, I have extended XMPP "IQ" as a transport to support **relaying**. Once peers has formed a connection through XMPP, they are able to route connection information to each other and attempt to form a direct connection. In the case that this is unsuccessful, they are able to fall back to this link as a means to transmit P2P data. This approach also has the benefit that, if a XMPP server does not support "Jingle", the two peers can still form links with each other. Since Brunet internally supports IP reflection, eventually, if one of the peers in the system has a public address, it will automatically assist the other peers into forming direct links with each other.

Rendezvous uses a two step approach. First peers advertise their use of private overlay in the resource identifier. The name is hashed to ensure that the users complete identifier does not extend past 1,023 bytes, the maximum length for these identifiers. In addition, a cryptographically generated random number is appended to the resource

identifier to distinguish between multiple instances of the users application in the same private overlay. Once a peer receives a presence notification from a remote peer and the base components match, that is the hash of the service, the peer adds it to a list of known online peers. If the peer lacks connections, the system broadcasts to that list a request for addresses. The peers respond with a list of addresses including UDP, TCP, and XMPP addresses, concluding rendezvous.

Ideally, peers would not need to create XMPP connections with each other; if they are on a public address, the rendezvous phase alone will suffice. When peers do not have a public address, they can obtain a mapping through STUN, then form an XMPP connection with each other, and finally perform simultaneous connection attempts. If NAT traversal fails, the peers can continue routing through the XMPP connection. Due to the abstractions employed by the transport library, the additional support for XMPP-based bootstrapping required only an additional 700 lines of code to Brunet and no modification to the core system.

3.4 Evaluating Overlay Bootstrapping

This section presents a qualitative evaluation of this system prototype bootstrapping a small-scale network as well as some of the experiences in deploying bootstrapping overlays.

3.4.1 Deployment Experiments

These experiments verify that the techniques work and determine expected overheads in using Brunet and XMPP to bootstrap an overlay. Rather than an extensive experiment overly focused on overheads of Brunet and XMPP, this experiment is primarily focused on the feasibility of forming small-scale overlays among network-constrained peers. The experiment represents 5 peers desiring all-to-all direct connectivity, a feature transparently available to them if they bootstrap into a private Brunet overlay. The experiments were run on peers deployed on 5 distinct virtual machines — each virtual

Table 3-1. Time in seconds for various private overlay operations

	Reflection	Rendezvous	Relaying	Connected
XMPP	0.035	0.110	0.243	20.3
Brunet	3.05	0.330	0.533	23.22

machine had its own separate NAT, and thus peers were unable to communicate directly without assistance.

The public Brunet overlay used in this experiment consisted of over 600 nodes running on PlanetLab. PlanetLab [23] is a consortium of research institutes sharing hundreds of globally distributed network and computing resources. GoogleTalk provided the XMPP overlay used in this experiment. Though this experiment does not take into advantage the features of the XMPP Federation, this aspect is presented in more detail in the next section reviewing experiences deploying overlays using XMPP.

In the experiment, 5 P2P nodes were started simultaneously, while measuring the time spent for reflection, rendezvous, reflection, and connection. The results are presented in Table 3-1. For XMPP, these are translated as follows: reflection measures the time to obtain IP addresses from the STUN server, rendezvous is the time to receive a presence notification, relaying is the time to receive a message across XMPP, and connected is once all nodes in the private overlay has all-to-all connectivity. For Brunet, these are translated as follows: reflection measures the time to connect to the public overlay, rendezvous is the time to query the DHT, relaying is the average time to send a message across the overlay, and connected is the time until the private overlay has all-to-all connectivity. The results are highly correlated to timeouts in Brunet, which employs a mixture of events and polling to stabilize the overlay, as well as the latency between the client and GoogleTalk. As this was more of a qualitative experiment, the results are clear: private overlays providing all-to-all connectivity among NATed nodes can bootstrap within a very reasonable amount of time.

3.4.2 Deployment Experiences

Recently, Facebook announced that they would be supporting XMPP as a means to connect to Facebook chat. This was rather exciting and further motivated this work, as Facebook has over 400 million active users, which would have made their XMPP overlay, potentially, the largest free-to-join overlay. Unfortunately, Facebook does not employ a traditional XMPP setup, instead it provides a proxy into their chat network, preventing features like arbitrary IQs and other forms of out-of-band messages to be exchanged between peers. User identifiers are also translated, so a peer cannot obtain a remote peers real identifier. Thus there exists no out-of-band mechanism for rendezvous. Peers could potentially send rendezvous messages through the in-band XMPP messaging, but this may be viewed by most recipients as spam as it would arrive as normal chat messages. Unfortunately, the realization is that not all XMPP servers, especially those unrelated to the Federation, support features necessary to bootstrap.

During initial tests in verifying the workings of the XMPP code base, I bootstrapped a private Brunet overlay on PlanetLab through various XMPP service providers. Unfortunately, some servers (GoogleTalk) ignored clients on PlanetLab. Another server crashed after 257 concurrent instances of the same account logged in. Because the provider had no contact information, I was unable to ascertain the reason for the crash. Though there did exist some servers that had no trouble hosting over 600 concurrent instances running on PlanetLab.

Once the system was running on PlanetLab, more tests were performed to determine the ability to bootstrap across the XMPP Federation. For this purpose, several friendships, or subscriptions, were formed between users across various XMPP service providers. In the most evaluated case, a single peer on GoogleTalk along with 600 peers on PlanetLab system using *jabber.rootbash.com*, the GoogleTalk peer would not always receive presence notifications for all peers online, though always would receive some. When a peer began the relaying mechanism, it would broadcast to every

peer from whom it received a presence notification. When performing this between GoogleTalk and *rootbash*, the GoogleTalk peer would not receive a response. Though in reducing the broadcast to a random selection of 10 peers, every 10 seconds until the GoogleTalk peer was connected, the peer received responses. The behavior indicates that the XMPP servers may have been filtering to prevent denial of service attacks.

Peers on the same XMPP server seem to be connected very quickly, though peers on different services can take significantly longer. For example, when bootstrapping a single peer from GoogleTalk into the *rootbash* system, it always took 1 minute for the node to become fully connected to the private overlay. When the peer used *rootbash*, the peer always connected within 30 seconds. It seems as if the communication between XMPP servers was being delayed for some reason. The same behavior was not experienced, when chatting between the two peers.

Table 3-2. Public and research overlays

	Description	Reflection	Rendezvous	Relay
BitTorrent	Default BitTorrent implementations rely on a centralized tracker to provide the initial bootstrapping. Peers can establish new connections through information obtained from established connections. This relegates the tracker as a means of monitoring the state of the file distribution. BitTorrent specifies a protocol, though each client may support additional features not covered by the protocol.	The current specification does not support NAT traversal, though future versions may potentially use UDP NAT traversal. At which point, BitTorrent may support a reflection service.	Peers can register as seeds to the same file hash, thus their IP address will be stored with the tracker.	Peers receive each other's IP addresses from the tracker, there is no inherent relaying.

Table 3-2. Continued

	Description	Reflection	Rendezvous	Relay
Gnutella	Gnutella is a large-scale unstructured overlay with over a million peers; primarily, it is used for file sharing. Gnutella consists of a couple hundred thousand ultra (super) peers to provide reliability to the overlay. Gnutella is free-to-join and requires no registration to use.	Work in progress. Peers attempt to connect to a sharer's resource, though a "Push" notification reverses this behavior. Thus a peer behind a NAT can share with a peer on a public address.	Peers can perform broadcast searches with TTL up to 2; when networks consist of millions of peers, small overlays will most likely not be able to discover each other.	Not explicitly, could potentially utilize ping messages to exchange messages.
Skype	Skype is a large-scale unstructured overlay, consisting of over a million active peers, and primarily used for voice over P2P communication. Skype, like Gnutella, also has super peers, though the owners of Skype provide authentication and bootstrap servers. Though Skype is free-to-join, it requires registration to use.	Skype APIs provide no means for reflection.	Skype supports applications, or add-ons, which can be used to transparently broadcast queries to a user's friend to determine if the peer has the application installed. Thus Skype does support rendezvous.	Skype applications are allowed to route messages via the Skype overlay, but because Skype lacks reflection, all communication must traverse the Skype overlay.
XMPP	XMPP consists of a federation of distributed servers. Peers must register an account with a server, though registration can be done through XMPP APIs without user interaction. XMPP is not a traditional P2P system, though it has some P2P features. XMPP servers on distinct servers are able to communicate with each other. Links between servers are created based upon client demand. During link creation, servers exchange XMPP Federation signed certificates.	While not provided by all XMPP servers, there exist extensions for NAT traversal. GoogleTalk, for example, provides both STUN and TURN servers.	Similar to Skype, XMPP friends can broadcast queries to each other to find other peers using the same P2P service. Thus XMPP supports rendezvous.	The XMPP specification allows peers to exchange arbitrary out-of-band communication with each other. Most servers support this behavior, even when sent across the Federation. Thus XMPP supports relaying.

Table 3-2. Continued

	Description	Reflection	Rendezvous	Relay
Kademlia [72]	There exists two popular Kademlia systems, one used by many BitTorrent systems, Kad, and the other used by Gnutella, called Mojito. Kademlia implements an iterative structured overlays, where peers query each other directly when searching the overlay. Thus all resources of a Kademlia overlay must have a publicly addressable network endpoint.	Existing implementations of Kademlia do not support mechanisms for peers to determine their network identity.	Peers can use the DHT as a rendezvous service, storing their connectivity information in the DHT at key location: <i>hash(SERVICE)</i> .	An iterative structured overlay has no support for relaying messages.
OpenDHT [86]	OpenDHT is a recently decommissioned DHT running on PlanetLab. OpenDHT is built using Bamboo, a Pastry-like protocol [94]. Pastry implements recursive routing, peers route messages through the overlay.	Existing implementations of Bamboo and Pastry do not support mechanisms for peers to determine their network identity. Though this is ongoing work.	Peers can use the DHT as a rendezvous service, storing their connectivity information in the DHT at key location: <i>hash(SERVICE)</i> .	Because Pastry uses recursive routing, it can be used as a relay. Furthermore, extensions to Pastry have enabled explicit relays called virtual connections [73].
Brunet [13]	Brunet like OpenDHT is a freely available DHT running on PlanetLab, though still in active development. Brunet creates a Symphony [71] overlay using recursive routing.	Brunet supports inherent reflection services, when a peer forms a connection with a remote peer, the peers exchange their view of each other.	Peers can use the DHT as a rendezvous service, storing their connectivity information in the DHT at key location: <i>hash(SERVICE)</i> .	Like Pastry, Brunet supports recursive routing and relays called tunnels [43].

CHAPTER 4

FROM OVERLAYS TO SECURE VIRTUAL PRIVATE NETWORKS

In this chapter, I take the results from Chapter 3 and apply them to IPOP [41] in order to construct a fully decentralized P2P VPN. While sharing overlays in IPOP makes for simplified use of the system, in reality, it introduces significant security challenges. For example, a misconfigured or malicious peer could potentially disable the entire overlay, rendering all VNs useless. If security and hence isolation is important, prior to VN deployment, a user would need to deploy a secure overlay and configure their VPN to bootstrap from it, given the complexity many users may reconsider the P2P approach and use a simple centralized VPN.

To address this challenge and to make a fully decentralized P2P VPN, I have extended the IPOP concept to support bootstrapping from public infrastructures and overlays into private and secure P2P overlays whose membership is limited to an individual VPN user base. Chapter 3 focused on a small scale feasibility of bootstrapping decentralized overlays. This chapter further extends into performance overheads of recursive Brunet overlays and larger network sizes. I then consider security in the overlay and present the first implementation and evaluation of an overlay with secure communication both between end points in the P2P overlay (e.g. VPN nodes) as well as between nodes connected by overlay edges. Security requires a means for peer revocation; however, current revocation techniques rely on centralized systems such as certificate revocation lists (CRLs). The proposed approach allows revocation using scalable techniques provided by the P2P overlay itself. I call the completed system and the interface used to administrate it **GroupVPN**, a novel decentralized P2P VPN.

The rest of this chapter is organized as follows. Throughout the chapter, there are two techniques used to evaluate my approaches, simulation and real system deployments; these are described in Section 4.1. Section 4.2 describes techniques

that allow users to create their own private overlays from a shared public overlay in spite of NATs. Use of security protocols has been assumed in many P2P works, though without consideration of implementation and overheads. I investigate implementation issues and overheads of security in P2P with emphasis on P2P VPNs in Section 4.3. Without revocation, use of security is limited, and in decentralized systems, the use of centralized revocation methods is not sufficient, I present novel mechanisms for decentralized revocation in Section 4.4. The complete system, GroupVPN, is presented in Section 4.5. Section 4.7 compares and contrasts this work with related work.

4.1 Experimental Environment

Throughout this paper, my quantitative evaluation environment uses both real deployments on PlanetLab and simulation. The evaluation requirements dictate the environment used. When the perspective of a single node is useful, PlanetLab's overloaded nature makes complex system analysis challenging, especially when attempting to simulate an instantaneous behavior on a system, which has random outage and delays in access.

IPOP uses Brunet as the underlying P2P infrastructure for connectivity. Brunet has been in active development for the past 5 years and is routinely run on PlanetLab [23] for experiments and tests. PlanetLab consists of nearly 1,000 resources distributed across Earth. In practical applications, though, roughly 40% of the resources are unavailable at any given time and the remaining behave somewhat unpredictably.

PlanetLab deployment takes approximately 15 minutes for all resources to have Brunet installed and connect to the overlay and then much more time to observe certain behaviors, making regression and verification tests complicated. To address this, I have extended Brunet to support a simulation mode. The simulator inherits all of the Brunet P2P overlay logic but uses simulated virtual time based upon an event-driven scheduler instead of real time. Furthermore, the simulation framework uses a specialized transport layer to avoid the overhead of using TCP or UDP on the host system, both of which are

limited resources and can hamper the ability to simulate large systems. The specialized transport uses datagrams to pass messages between nodes, thus from the node's perspective, it is very similar to a UDP transport and can simulate both latency and packet dropping. Latency between all node pairs is set to 100 ms by default.

Both simulation and real system evaluation provide unique advantages. Simulations allow faster than real time execution of reasonable sized networks (up to a few thousand) using a single resource, while enabling easy debugging. In contrast, deployment on real systems, in particular PlanetLab, presents opportunities to add non-deterministic, dynamic behavior into the system which can be difficult to replicate, such as network glitches and long CPU delays on processing.

4.2 Towards Private Overlays

Many users of IPOP begin by using the public shared overlay and, once comfortable, move towards hosting their own infrastructure. Some are successful without assistance, while a majority are not. Network configuration issues tend to be the most common issue preventing users from hosting their own independent IPOP systems. While users were able to easily join the shared overlay, similar attempts to construct their own were hindered and ultimately only successful after receiving feedback.

Prior work in IPOP [44] enabled many VPNs to share a single P2P overlay by storing IP address into the DHT at the key *hash(Namespace : IP)*. Unfortunately, this approach is fraught with security issues. In the previous chapter, I established methods that enabled bootstrapping private Brunet overlays as easily as connecting to a public P2P overlay. This chapter begins by focusing on the integration of the methodologies employed in recursive Brunet overlays as applied to IPOP.

To bootstrap from an existing Brunet overlay, peers first insert their public overlay node address into the key represented by *hash(\$PrivateOverlayNamespace)* and continue to do so regularly until they disconnect, so as to not let the entry become stale and disappear. Peers attempting to bootstrap into the private overlay can then query this

key and obtain a list of public overlay nodes that are currently acting as proxies into the private overlay. By using the public overlay as a transport, similar to UDP or TCP, the private overlay node forms bootstrapping connections via the public overlay. At which point, overlay bootstrapping proceeds as normal. The entire process is represented in Figure 3-1.

As mentioned in the previous chapter, small overlays may have no members with a public address, making it difficult to provide overlay based NAT traversal. To avoid having a special case for NAT traversal in private overlays, in my model, the private overlay share TCP and UDP sockets with the public overlay. This mechanism, referred to as “*pathing*”, allows multiplexing a single UDP socket and listening TCP socket by many overlays. This is only possible due to the generic transports library of the Brunet P2P overlay, which does not differentiate UDP, TCP, or even relayed links. Pathing works as a proxy, intercepting a link creation request from a local entity, mapping that to a path, and then requesting from the remote entity a link for that path. The underlying link is then wrapped by pathing and given to the correct overlay node, resulting in a completely transparent multiplexing of a TCP and UDP sockets, thereby enabling the NAT traversal in one overlay to benefit the other. Once a link has been established, the pathing information is irrelevant, limiting the overhead into the system to a single message exchange during link establishment.

4.2.1 Time to Bootstrap a Private Overlay

This experiment focuses on the overheads in bootstrapping a private overlay using the techniques mentioned in the previous section. The time to bootstrap can be derived analytically by considering the minimum steps for a node to join the public overlay, obtain private overlay peers from the public overlay DHT, and then connect to the private overlay. In Brunet, peers begin by forming leaf or bootstrapping connections and use these to communicate with the neighbor or peer in the P2P network nearest to their P2P address. The process to form a connection can be done in as few as 4 messages and

up to 6, if the peers only know each other's P2P address, which is the case for neighbor connections.

Assuming a peer already has IP address information for another, a connection can be initiated by the peer sending a message to the remote peer expressing the desire for a connection. The remote node responds by either rejecting the request or committing to the connection. In the next exchange, the initiating peer commits to forming the connection and the remote peer acknowledges. The two phase commit process is used to handle the complexity that ensues when multiple simultaneous connection attempts occur in parallel. All these messages take 1 hop, since they are direct links between peers.

When peers only have each other's P2P address and/or the initiating peer is behind a NAT, it may take fifth and sometimes a sixth message. These messages are requests for the remote peer's IP addresses as well as asking the peer to connect with the initiating peer, addressing the case where the remote peer is behind a NAT and cannot handle inbound messages. These messages are routed over the overlay taking $\log(N)$ hops, where N is the network size of the public overlay.

Private overlay bootstrapping follows a similar process, though, first, the peer acquires P2P addresses of other participants through the public DHT, an operation taking $2 * \log(N)$ hops. In the private overlay, the leaf connections do not communicate directly; rather, they use the public overlay, causing some of the 1 hop operations above to take $\log(N)$ hops. Finally, finding the nearest remote peer in the private overlay takes $\log(N) + \log(n)$, where n is the network size of the private overlay.

Given this model, each operation takes the following hop counts: public overlay bootstrapping $\Rightarrow 8 + \log(N)$, DHT operations $\Rightarrow 2 * \log(N)$, and private overlay bootstrapping $\Rightarrow 4 + 5 * \log(N) + \log(n)$. The cumulative operation takes $12 + 8 * \log(N) + \log(n)$ hops. The dominating overhead in bootstrapping the private overlay is the time it takes to perform overlay operations on the public overlay ($\log(N)$). For

instance, assuming a network size of 512 public and 8 private, a node should be connected within 87 hops.

To evaluate my implementation for GroupVPN, I used both PlanetLab and the simulator. 100 tests were run for various network sizes. Though due to difficulty in controlling network sizes in PlanetLab, I set each PlanetLab node to randomly decide if it would connect to the private overlay. The network sizes were then used in the simulator and the analytical model. The average public network size for each of these tests was 600. The results are presented in Figure 4-1¹.




Figure 4-1. CDF of the time to bootstrap a private overlay node in a private overlay of the size stated in the legend using a public overlay consisting of 600 nodes. Using a 100 ms delay like the simulator results in 9.2 and 9.3 seconds for the analytical model for private network sizes of 68 and 147, respectively.

Based upon the results presented in Figure 4-1, the bootstrapping time for the implementation performs better than the analytical model, due to the simplicity of the analytical model and the small network sizes. It is of interest that while the simulator

¹ I performed measurements for many more private network sizes, but all the results were so similar that it did not introduce anything of interest and are omitted from the plots to improve clarity.

results tend to be in a well defined range, the PlanetLab results have a few outliers with long bootstrap times. Some of the expected causes for this are churn in the system and state machine timeouts in Brunet, though I have not considered this in much depth.

4.2.2 Overhead of Pathing

Much like the previous experiment, this verifies that the pathing technique has negligible overheads for VPN usage. To determine the overheads, two GroupVPNs are deployed on resources on the same gigabit LAN. To measure latency and throughput, netperf experiments are run for 30 seconds, 5 times each on an unutilized network switch. Other specifications of the machine are ignored as the system without pathing is used as the baseline. The results, Table 4-1, indicate that the use of pathing presents negligible overhead for both throughput and latency, justifying the use of this approach to transparently deal with NAT and firewall traversal.

Table 4-1. Pathing overheads

	Latency (ms)	Throughput (Mbit/s)
Standard	0.303	225.27
Pathing	0.308	224.36

4.3 Security for the Overlay and the VPN

Structured overlays are difficult to secure and a private overlay is not secure if it provides no means to limit access to the system. Malicious users can pollute the DHT, send bogus messages, and even prevent the overlay from functioning, rendering the VPN useless. To address this in means that make sense for VPNs and common users, I have employed a public key infrastructure (PKI) to encrypt and authenticate both communication between peers as well as communication across the overlay, called point-to-point (PtP) and end-to-end (EtE) communication, respectively.

Use of a PKI motivates from the ability to authenticate without a third party, ideal for P2P use, unlike a key distribution centers (KDC) used by other VPNs. A PKI can use either pre-exchange public keys or a certificate authority (CA) to sign public keys,

i.e., certificates. Thus peers can exchange keys and certificates without requiring a third-party to be online.

The reasons for securing PtP and EtE are different. Securing PtP communication prevents unauthorized access to the overlay, as peers must authenticate with each other for every link created. Though once authenticated, a peer can perform malicious acts and since the overlay allows for routing over it, the peer can disguise the origination of the malicious acts. By also employing EtE security, the authenticity of messages transferred through an overlay can be verified. Though EtE security by itself, will not prevent unauthorized access into the overlay. By employing both PtP and EtE, overlays can be secured from uninvited guests from the outside and can identify malicious users on the inside. Implementing both leads to important questions: what mechanisms can be used to implement both and what are the effects of both on an overlay and to a VPN on an overlay.

4.3.1 Implementing Overlay Security

There are various types of PtP links, such as TCP and UDP sockets and relays across individual nodes and the overlay. EtE communication is datagram-oriented in IPOP. Traditional approaches of securing communication such as IPsec are not convenient due to complexity, i.e., operating system specific, portability constraints, and lack of common APIs. Security protocols that rely on reliable connections, such as SSL or TLS are undesirable as well as they would require a userspace implementation of reliable streams (akin to TCP). As such, I have implemented an abstraction called a security filter as presented in Figure 4-2, which enables nearly transparent use of security libraries and protocols. To this date, I have implemented both a DTLS [83] filter using the OpenSSL implementation of DTLS as well as a protocol that reuses cryptographic libraries provided by .NET that behaves similarly to IPsec.

A security filter has two components: the manager, and individual sessions or filters. While the individual sessions could act as filters by themselves, by combining with a




Figure 4-2. An example of the security filter abstraction used by senders and receivers through an EtE secured chat application. Each receiver and sender use the same abstracted model and thus the chat application requires only high-level changes, such as verifying the certificate used is Alice's and Bob's, to support security.

manager, they can be configured for a common purpose and security credentials. This approach enables the use of security to be transparent to the other components of the system as the manager handles session establishment, garbage collection of expired sessions, and revocation of peers.

Certificate embed identity of the owner, thus a signed certificate states that the signer trusts that the identity is accurate. In network systems, the certificate uses the domain name to uniquely identify and limit the use of a certificate. When a CA signs the certificate, by including the domain name, it ensures that users can trust that a certificate is valid, while used to secure traffic to that domain. Communication with another domain using the same certificate will raise a flag and will result in the user not

trusting the certificate. In environments with NATs, dynamic IP addresses, or portable devices, typical of P2P systems, assigning a certificate to a domain name will be a hassle as it constrains mobility and the type of users in the system. Furthermore, most users are unaware of their IP address and changes to it. Instead, a certificate is signed against the user's P2P address and unique user name as delegated by the CA. The purpose of the former is for efficiency of revocation as discussed in Section 4.4. During the formation of PtP links or while parsing EtE messages, the two nodes discover each other's P2P addresses. If the addresses do not match the address on the verified certificate, the communication need not proceed further.

Prior to trusting the security filter, the core software or the security filter must ensure that the P2P address of the remote entity matches that of the certificate. In my approach, I did this by means of a callback, which presents the underlying sending mechanism, EtE or PtP, and the overlay address stored in the certificate. The receiver of the callback can attempt to cast it into known objects. If successful, it will compare the overlay address with the sender type. If unsuccessful, it ignores the request. If any callbacks return that the sender does not match the identifier, the session is immediately closed. Thus the security filter need not understand the sending mechanism and the sending mechanism need not understand the security filter.

The last consideration comes in the case of EtE communication that provides an abstraction layer. For example, in the case of VPNs, where a P2P packet contains an IP packet and thus a P2P address maps to a VPN IP address, a malicious peer may establish a trusted link, but then hijack another user's IP session. As such, the application must verify that the IP address in the IP packet matches the P2P address of the sender of the P2P packet. In general, an application address should be matched against a P2P address.

4.3.2 Overheads of Overlay Security

When applying an additional layer to a P2P system, there are overheads in terms of time to connect with the overlay. Other less obvious effects are throughput, latency, and processing overheads, assuming that the P2P system will be used over a wide area network, where the latency and throughput limitations between two points will make the overhead of security negligible. Though bootstrapping will be affected due to additional round trip messages used for forming secure connections.

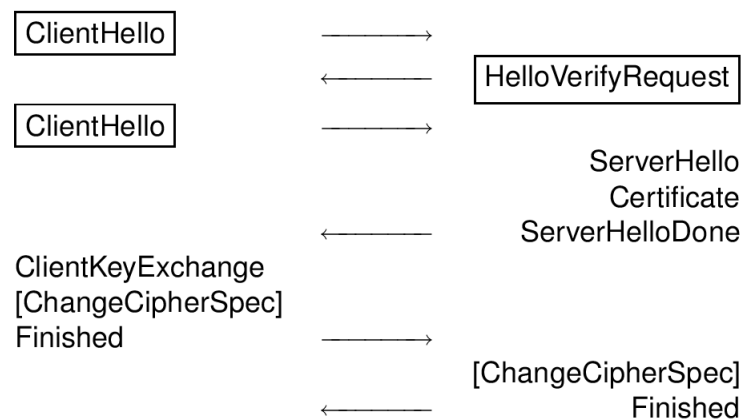


Figure 4-3. DTLS handshake

The DTLS handshake as presented in Figure 4-3, which consists of 6 messages or 3 round trips. PtP security may very well have an effect on the duration of overlay bootstrapping. There even exists a possibility that with more messages during bootstrap, the probability one drops is higher, which could, in turn, also have an effect, though possibly negligible, on time to connect. To evaluate these concerns, I have employed both simulation and real system experiments.

The following experiments use both simulation and PlanetLab deployment to evaluate time to connect a new node to an existing resource. Then another experiment is performed to evaluate how long it takes to bootstrap various sized overlays if all nodes join at the same time. This experiment is only feasible via simulation as attempting

to reproduce in a real system is extremely difficult due to how quickly the operations complete.

4.3.2.1 Adding a Single Node

This experiment determines how long it takes a single node to join an existing overlay with and without DTLS security. The experiment is performed using both simulation and PlanetLab. After deploying a set of nodes without security and with security on PlanetLab, the network is crawled to determine the size of the network. In both cases, the overlay maintained an average size of around 600 nodes. At which point, I connected a node 1,000, each time using a new, randomly generated P2P address, thus connecting to a different point in the overlay. The experiment concludes as soon as the node has connected to the peers in the P2P overlay immediately before and after it in the P2P address space. In the simulation, a new overlay is created and afterward a new node joins, this is repeated 100 times. The cumulative distribution functions obtained from the different experiments are presented in Figure 4-4.




Figure 4-4. Time in seconds for a single node to join a secure (dtls) and insecure (nosec) structured overlay, using both PlanetLab (plab) and the Simulator (sim).

4.3.2.2 Bootstrapping an Overlay

The purpose of this experiment is to determine how quickly an overlay using DTLS can bootstrap in comparison to one that does not given that there are no existing participants. Nodes in this evaluation are randomly given information about 5 different nodes in the overlay and then all attempt to connect with each other at the same time. The evaluation completes after the entire overlay has all nodes connected and in their proper position. For each network size, the test is performed 100 times and the average result is presented in Figure 4-5.




Figure 4-5. Time in seconds for a secure (dtls) and insecure (nsec) structured overlay to bootstrap, given that all nodes bootstrap simultaneously.

4.3.3 Discussion

Both evaluations show that the overhead in using security is practically negligible, when an overlay is small. In the case of adding a single node, it is clear that the simulation and deployment results agree, as the difference between bootstrapping into an overlay with and without security remains nearly the same. Clearly this motivates the use of security if time to connect is the most pressing question.

The time to bootstrap a secure overlay was not significantly more than that of an insecure overlay. What I realized is that complex connection handshaking, as

implemented in Brunet, seems to dominate connection establishment time. For example, in Brunet, two peers must communicate via the overlay prior to forming a connection, and the system differentiates between bootstrapping connections and overlay connections. Thus even though a peer may have a bootstrapping connection, it will need to go through the entire process to form an overlay connection with a peer. While this may lead to inefficiencies, this simplification keeps the software more maintainable and easier to understand.

4.4 Handling User Revocation

Unlike decentralized systems that use shared secrets, in which the creator of the overlay becomes powerless to control malicious users, PKIs enable their creators to effectively remove malicious users. Typical PKIs either use a certificate revocation list (CRL) or online certificate verification protocols such as Online Certificate Status Protocol (OCSP). These approaches are orthogonal to decentralized systems as they require a dedicated service provider. If the service provider is offline, an application can only rely on historical information to make a decision on whether or not to trust a link. In a decentralized system, these features can be enhanced so not to rely on a single provider. In this section, I present two mechanisms of doing so: storing revocations in the DHT and performing overlay broadcast based revocations.

4.4.1 DHT Revocation

A DHT can be used to provide revocation similar to that of OCSP or CRLs. Revocations, a hash of the certificate and a time stamp signed by the CA, are stored in the DHT at the key formed by the hashing of the certificate. In doing so, revocations will be uniformly distributed across the overlay, not relying on any single entity.

The problem with the DHT approach is that it does not provide an event notification for members currently communicating with the peer. While peers could continue to poll the DHT to determine a revocation, doing so is inefficient. Furthermore, a malicious

peer, who has a valid but revoked certificate could force every member in the overlay to query the DHT, negatively affecting the DHT nodes storing the revocation.

4.4.2 Broadcast Revocation

Broadcast revocation uses a structured overlay based broadcast approach as described in Appendix 8. The form of broadcast can be used to perform to notify the entire overlay immediately about a new revocation. It is important to note, that the message needs to be delivered locally prior to forwarding, so that peers who have a connection to the malicious peer, will end the connection prior to accidentally forwarding the message to the peer by receiving and acting upon the revocation prior to forwarding the message.

4.4.3 Evaluation of Broadcast

I performed an evaluation on the broadcast using the simulation to determine how quickly peers in the overlay would receive the message. The tested network sizes ranged from 2 to 256 in powers of 2. The tests were evaluations were performed 100 times for each network size. The CDF of hops for each node are presented in Figure 4-6. The results make it quite clear that the broadcast can efficiently distribute a revocation much more quickly than $\log(N)$ time.

4.4.4 Discussion

In contrast to the DHT solution, broadcast revocation occurs only once and leaves no state behind. Thus the broadcast is not a complete solution, as new peers connected to the overlay or those who missed the broadcast message will be unaware of a revocation. Furthermore, if an overlay is shared by many VPNs, it may prevent overlay broadcasting or itself may be inefficient.

The DHT solution by itself may also not sufficient as revocations may be lost over time as the entries must have their leases renewed in the DHT. To address this condition, each peer maintains a local CRL and the owner of the overlay can occasionally send updates to the CRL through an out of band medium, such as e-mail.




Figure 4-6. Overlay broadcast time CDF.

A better long term solution may be the use of a gossip protocols so that peers can share their lists with each other during bootstrapping phases.

A key assumption in using these is that a Sybil [30], or collusion attack, is difficult in the secured overlay. If a Sybil attack is successful, both a DHT and broadcast revocation may be unsuccessful, though peers could fix this problem by obtaining the CRL out of band. In addition, previous work [20] has described decentralized techniques to limit the probability of such attacks from occurring. In my approach, the use of central authority to review certificate requests can be used to limit a single user from obtaining too many certificates as well as ensuring uniform distribution of that user's P2P addresses, further hampering the likelihood of a Sybil attack. The ability to automate this is left as future work.

One way to mitigate Sybil attacks using the broadcast approach is to bundle colluding offenders into a single revocation message. That would prevent those from colluding together to prevent each other's revocations. Furthermore, while not emphasized above, revocation in my system revokes by user name and not individual certificates. Combined these two components limit Sybil attacks against broadcast.

4.5 Managing and Configuring the VPN

While the PKI model applies to P2P overlays, actual deployment and maintenance of security credentials can be too complex to manage, particularly for non-experts. Most PKI-enabled systems require the use of command-line utilities and lack methods for assisting in the deployment of certificates and policing users. My solution to facilitate use of PKIs for non-experts is a partially-automated PKI reliant on a group-based Web interface distributable in forms of Joomla add-ons as well as a virtual machine appliance. In this environment, groups can share a common Web site, while each group has their own unique CA. Although this does not preclude other methods of CA interaction, experience has shown that it provides a model that is satisfactory for many use cases.

Group-based Web 2.0 sites enable low overhead configuration of collaborative environments. The roles in a group environment can be divided into administrators and users. Users have the ability to join and create groups; whereas administrators define network parameters, can accept or deny join requests, remove users, and promote other users to administrators. By applying this to a VPN, the group environment provides a simple to use wrapper around PKI, where the administrators of the group act as the CA and the members have the ability to obtain signed certificates.

Elaborating further, when a user joins a group, the administrator can enable automatic signing of certificates or require prior review; and when peers have overstayed their welcome, an administrator can revoke their certificate by removing them from the group. Revocations are handled as described in [Section 4.4](#). In the context of GroupVPN systems, a user revocation list as opposed to a CRL simplifies revocation, since users and not individual certificates will be revoked.

Registered users who create groups become administrators of their own groups. When a user has been accepted into a group by its administrator, they are able to download VPN configuration data from the Web site. Configuration data is loaded by the GroupVPN during its configuration process to specify IP address range, namespace,

and security options. The configuration data also stores a shared secret, which uniquely identifies the user, enabling the Web site to automatically sign the certificate (or enqueue it form manual signing, depending on the group's policy). Certificate requests consist of sending a public key and a shared secret over an HTTPS connection to the web server. Upon receiving the signed certificate, peers are able to join the private overlay and GroupVPN, enabling secure communication amongst the VPN peers. The entire bootstrapping process, including address resolution and communication with a peer, is illustrated in Figure 4-7.

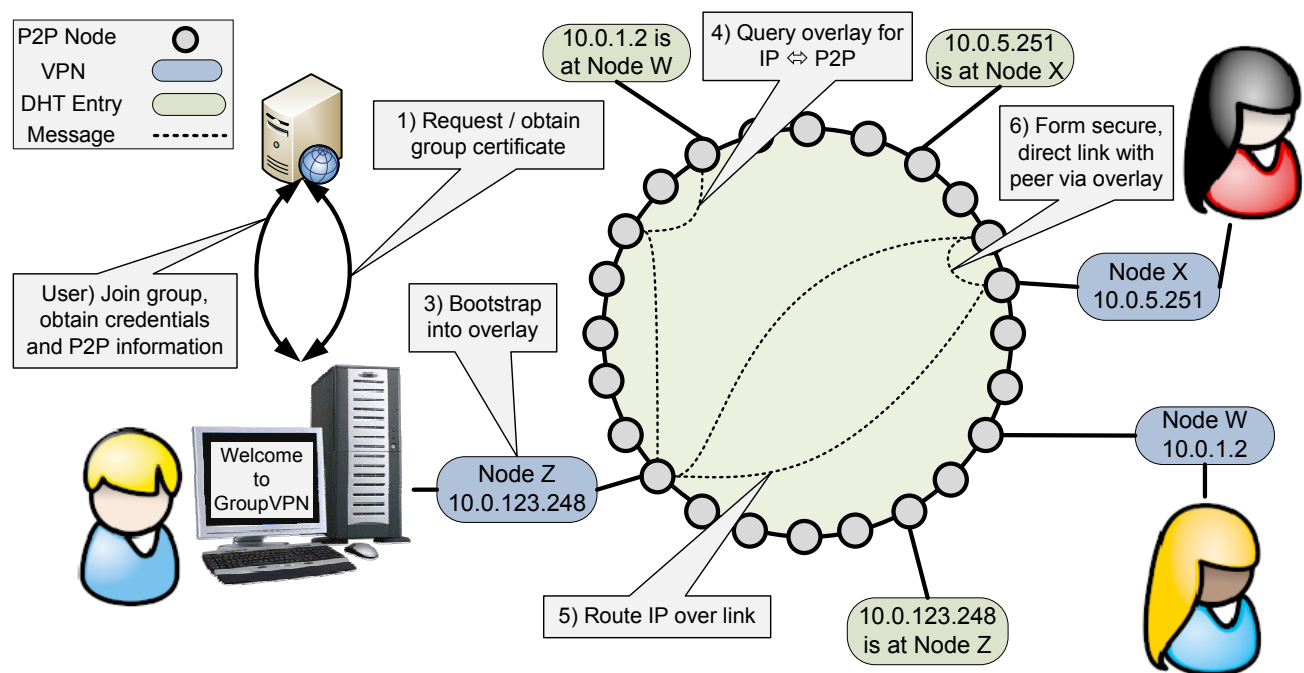


Figure 4-7. Bootstrapping a new GroupVPN

There are many ways of implementing and hosting the Web site. For example, Google offers free hosting of Python web applications through Google Apps, an option available if the user owns a domain. Alternatively, the user could host the group site on a public virtual network. In this case, peers interacting with the GroupVPN would need to connect with the public virtual network in order to create an account, get the configuration data, and retrieve a signed certificate, at which point they could disconnect from it. This does not preclude the use of other social mediums nor a central site

dedicated to the formation of many GroupVPNs. Many GroupVPNs can share a single site, so long as the group members trust the site to host the CA private key.

4.6 Leveraging Trust from Online Social Networks

Groups are very useful for coordinating a set of individuals when a subset of them can be used to establish trust amongst them all. However, groups can lack clear and concise individuality and limit independence from the collective. As noted in [56], trust can be leveraged from existing social networks to create trust in other domains. Much like that of a social network, a VPN consists of trusted links, tying the two together produced the SocialVPN [56]. In this work, I along with my fellow researchers implemented a prototype for SocialVPN, which exemplifies the utility of my approaches in handling security both in terms of trust and session establishment as well as endpoint configuration of the VPN.

Besides the content described in the following subsections, namely establishing identity and trust as well as address allocation and discovery, SocialVPN reuses existing components already provided in IPOP, such as secure link establishment, endpoint configuration, and packet handling. Even the functionality used by SocialVPN for address allocation and discovery builds upon existing abstractions already provided by IPOP.

4.6.1 Architecture

SocialVPN leverages the online social network to establish trust and exchange certificates. Thus a social network must provide a means for an external application to determine friendships and store arbitrary data into the social network. The arbitrary data in this case would be the certificate that can be used to find a peer in the network overlay and verify its identity. The certificate consists of the peer's social network information and P2P address. Thus once a peer has connected to a social network the first time, it need only be repeated to obtain the latest information. Existing certificates

remain valid until the friendship has ended for the certificate cache has been explicitly flushed.

Once a peer has a certificate, connections are immediately established with all “friends” that are currently connected to the overlay. As new peers come online, they establish connections with those already there. Due to potential network problems, this may not occur, and so all members of SocialVPN will occasionally check the liveness of peers not connected to the overlay. Because online peers have an active connection, there is no need to explicitly monitor their state. When they go offline, the connection will be broken and can be represented to the user appropriately.

The motivation in establishing connections immediately comes from two purposes: no overhead in bootstrapping on demand connections and better ability to distribute IP multicast and broadcast packets. Using the traditional IPOP style to establish a direct connection can take somewhere between several hundreds of milliseconds up to several seconds, which may be disappointing to users who have used centralized VPNs that have much faster connection establishment. Because there currently exists no support for efficient broadcast / multicast message distribution inside SocialVPN, maintaining an active link to all peers allows a peer to push that message to all their peers without having to establish a new trusted link first.

4.6.2 Leveraging Trust From Facebook

Trust or friendships already established in Facebook used a now deprecated technology that allowed desktop applications access to Facebook. Certificate exchange relied upon a web-based data store component provided by Facebook, which was presented as a database. When SocialVPN first contacts Facebook, it would add the current certificate, if it did not exist and then download certificates for all friends that it did not already have. Because a user might have more than one instance of SocialVPN running, the database was designed to allow the user to store multiple certificates and to clear their certificates. As mentioned earlier, each certificate contains the friends P2P

address, which allows a peer to discover a remote peer and establish a trusted, direct VPN link with them.

Unfortunately, Facebook's interface was poorly constructed and no longer exists. Each application had to embed in itself private key information to authenticate itself with Facebook. A malicious attacker could easily discover this and change the stored data to suit their needs. While Facebook never gave a reason for shutting down the desktop application component of their system, this is a probable reason.

As an alternative, I developed a web application, which was used for a short period of time to replace the desktop application; however, this was fraught with problems. Unlike the desktop application that only required trust with Facebook, the web application required hosting a third-party web site to support the system. The trust model is not significantly different as the administrator for the application has access to the trusted material regardless, it simply meant another centralized component. These complications led to the development of an XMPP-based SocialVPN.

4.6.3 Leveraging Trust from XMPP

Unlike Facebook, XMPP is a well standardized, open system with many individual members contributing compatible services. Thus if one of them decides to break from the XMPP specification, users can easily migrate to another service provider. After the unfortunate incident with Facebook, this open aspect was much more attractive to SocialVPN as a research project

As discussed in Section 3.3.2, XMPP is primarily used as an open protocol instant messenger service. Though it has support for exchanging binary messages through "IQ." When a peer using SocialVPN connects to XMPP, they are informed of their friends that are online. Each friend has a unique name in the form "username@domain/resource." When a peer receives this message, they can determine if that friend is using SocialVPN by the resource name. If the peer is discovered to be using SocialVPN they will exchange certificates and proceed to establish trusted links in the overlay.

4.6.4 Address Allocations and Discovery

In creating a VPN where links are defined by social networking relationships, the mechanisms for IP address allocation via the DHT do not apply well. A social networking based VPN will form an overlay that does not need to rely on a structured overlay and so to do without requires a new addressing scheme. Additionally, attempting to place all peers inside a social network within an IP address range, especially IPv4, is fraught with problems [92]. Namely that it can be difficult to find a common address space for all peers inside a VPN, which can be made even more difficult if those peers use another VPN product.

The concept employed in SocialVPN is to place each user inside their own private address space independent of other users. Each friend of that peer has a unique IP address unbeknownst to them inside this address space. The IP is mapped to the users P2P address. Then through the use of packet translation, IP addresses are transparently changed as the packet is transferred between peers. Prior to delivery, the packet's destination address is converted to the peers own pre-defined IP address and the source address is based upon a mapping stored inside a hashtable that maps P2P address to IP address. Only trusted peers will have a mapping like this.

4.7 Related Work

4.7.1 VPNs

Hamachi [67] is a centralized P2P VPN provider using the web site for authentication, peer discovery, and connection establishment. While the Hamachi protocol claims to support various types of security [68], the implementation appears to only support the key distribution center (KDC) requiring that all peers establish trusted relationship through the central website. The Hamachi approach makes it easy for users to deploy their own services, but places limitations on network size, uses a proprietary security stack, and does not allow independent VPN deployments. In contrast, my approach presents a completely decoupled environment allowing peers to start using the

shared system to bootstrap private overlays and migrate away without cost if need be. Furthermore my approach relies only on a central server to obtain the certificate otherwise, it is decentralized. In Hamachi, if the central server goes offline, no new peers can join the VPN.

Campagnol VPN [12] provides similar features to Hamachi: a P2P VPN that relies on a central server for rendezvous or discovery of peers. The key differences between Hamachi and Campagnol is that Campagnol is free and does not provide a service; users must deploy their own rendezvous service. The authors of Campagnol also state that the current approach limits the total number of peers sharing a VPN to 100 so not to overload the rendezvous service. The current implementation does not support a set of rendezvous nodes, though doing so would make the approach much more like ours. In addition, the system relies on traditional distribution of a CRL to handle revocation.

Tinc [100] is a decentralized VPN requiring users to manually organize an overlay with support for finding optimal paths. In comparison to my approach, Tinc does not automatically handle churn in the VPN. If a node connecting two separate pieces of the VPN overlay goes offline, the VPN will be partitioned until a user manually creates a link connecting the pieces. Furthermore, Tinc does not form direct connections for improved latency and throughput reasons, thus members acting as routes in the overlay incur the price of acting as packet forwarders.

The last VPN, I discuss is the most similar to IPOP, its called N2N [29]. N2N uses unstructured p2p techniques to form an Ethernet based VPN. While their approach, like ours, has built-in NAT traversal, it requires that users deploy their own bootstrap and limits security to a single pre-shared key for the entire VPN, thus users cannot be revoked. Since N2N provides Ethernet, users must provide their own mechanism for IP address allocation, while discovery utilizes overlay broadcasting. Thus there are concerns that as systems get larger, N2N may not be very efficient.

4.7.2 P2P Systems

BitTorrent [10], a P2P data sharing service, supports stream encryption between peers sharing files. The purpose of BitTorrent security is to obfuscate packets to prevent traffic shaping due to packet sniffing. Thus BitTorrent security uses a weak stream cipher, RC4, and lacks peer authentication as symmetric keys are exchanged through an unauthenticated Diffie-Hellman process.

Skype [99] provides decentralized audio and video communication to over a million concurrent users. While Skype does not provide documentation detailing the security of its system, researchers [35, 48] have discovered that Skype supports both EtE and PtP security. Though similar to Hamachi, Skype uses a KDC and does not let users setup their own systems.

As of December 2009, the FreePastry group released an SSL enabled FreePastry [94]. Though relatively little is published regarding their security implementation, the use of SSL prevents its application for use in the overlay and for overlay links that do not use TCP, such as relays and UDP. Thus their approach is limited to securing environments that are not behind NATs and firewalls that would prevent direct TCP links from forming between peers.

CHAPTER 5

EXTENSIONS TO P2P OVERLAYS AND VIRTUAL NETWORKS

This chapter contains components that extend the VPN software to provide additional important features. Many of these components derive from experiences and demands that have arisen as a result of the deployment of the VPN software in real systems. Deployment experiences include but are not limited to usage on PlanetLab, resources including personal computers and clusters in residential and academics environments, virtual machines, and cloud resources. Each of these environments exposes a different set of requirements to the design and implementation of a practical P2P VPN.

5.1 Built-in Self-Simulation

Software systems are complex and involve many moving parts. Traditionally, system design begins by considering the goals of the system, choosing algorithms and data structures that can achieve those goals, and simulating or modeling the system. Those results then translate into a real system that consists of a new code based upon the concept in the simulation. In this process, simulation is applied primarily to validate a design concept but not its implementation. Then the entire software base must be independently checked for bugs and other issues that may have already appeared in the simulation code, doubling developer efforts.

To reduce efforts in development and evaluation, I have investigated and implemented mechanisms for distributed systems and in particular Brunet to support built-in self-simulation using event-driven simulation techniques. In other words, even though Brunet is written for real system deployment, the same code can run using simulated communication links and simulated time, allowing many nodes to run on the same resource and potentially faster than wall clock time. This approach allows transitioning features from simulation directly into deployment, hastening development cycles. Furthermore, interesting discoveries in the real system can be modeled in simulation to

make the simulation behavior more accurate. Because a simulation can run on a single computer, scaling up to a significantly large system, new features can be constructed and evaluated locally, removing many bugs and reusing and applying test cases already present in the simulation environment significantly reducing testing overheads.

The concept can be applied to networking / distributed systems in general. Distributed system software can usually be divided into many pieces, such as network communication, state, time-based events, user actions, and so on. Simulation of these systems focuses on three aspects handling of time-based events, communication between the various members of the system, and the injection and handling of user actions. In the rest of this section, I will discuss these in more depth and discuss how I addressed them in the context of Brunet.

5.1.1 Time-Based Events

Events or actions cause changes in a system. Some are due to external stimuli, such as hardware or software interrupts in a processors or user input, others are a result of timers, which may be a subset of hardware interrupts. In the context of simulation, timers and external stimuli can be viewed as two different components. The external stimuli may be delivered based upon a timed event or an action initiated by a remote party. If time is ignored, then a node will run in a loop until its steady state has been achieved and then constantly verifying that it still is in steady state. Timing allows a node to delay this behavior, such as establishing more connections or verifying its connectivity, behavior which produces more efficient systems.

A system could be made entirely without timers and run on external events alone. In this case, timing is still required to model the communication delay between peers. A message sent from one peer should not instantaneously arrive at another peer. As will be described in [Section 5.1.2](#), peers can use timers to simulate latency between peers.

Events in a simulation are stored in a timer with delays specified in terms of a virtual clock. Methods in order to retrieve the current time should be based upon the same

clock in both simulation and deployment systems. In a real system, this would then reveal the actual current time, whereas in a simulated environment, this would reveal the current simulated time. By virtualizing time retrieval calls, the caller can be directed to the appropriate clock depending on whether the system is running in simulation mode or not. How this is implemented depends on the language the software is written in. For example, languages with namespaces can easily replace the clock functions with their own. Languages like C may require pre-processor macros to specify real or virtual time.

As events are queued into the system, they must be stored in ordered. The structure should be such that the event to execute next is always available in minimum time, while optimizing for inserting and removing events from the timer. For this application, a minimum heap works well. A minimum heap provides constant seek time for the smallest value as well as $\log(N)$ insert and deletion time. In Brunet, this has been implemented as a binary heap.

After the system has initialized, it may add one or more events into the timer to cause an action to occur. The simulator will then advance the virtual clock to the time the next event is supposed to occur, execute all events that occurred up to that point including the next event and then repeat. The running of events should not stop until the next event to execute should be run in the future, because one event may cause another event to occur immediately requiring no delay. It should be noted that events may want to execute other events. These events should not be executed in-line and instead should be added into the queue to be executed at the same virtual clock time. If this is not done, there is a potential for stack overflows due to extremely deep calls into the code.

5.1.2 Network Communication

Using communication models or transports that rely on limited resources such as the number of open sockets or interactions with the operating system can severely hinder the usability and functionality of a simulator. A system using sockets will quickly hit a wall, Linux, for example, limits the amount of open file descriptors to 1,024, which

means that in a UDP system a simulation would be limited to having as few as 1,024 peers in the system whereas a TCP system could be unable to proceed with more connections than 32 (32 peers with all-to-all connections would result in 1,024 active TCP sockets). Furthermore, each interaction with a socket requires at least one if not more transitions between user-space and kernel-space.

So while existing transports could be used for simulated communication, the overhead in doing so is undesirable as it would limit large-scale simulations. Assuming that the system is modularly written, it is possible for various forms of transport layers to be used for network communication. Thus for scalability peers could exchange buffers or pointers to messages with each other. This would remove any restrictions on OS resources and would not require that each communication pathway pass through a system call.

Brunet supports a generic transports framework that provides the method for sending a message and the ability to register a callback when a message is received. This concept is built into an “edge.” Each edge is associated with a remote peer, and when sending or receiving a packet, the destination or source would be the peer associated with that edge. Edges come in pairs, if they are connected, thus a simulated edge consists of two components: knowledge of the remote edge and timing measurement of the latency between the two peers. When a peer sends a message to the remote peer, the simulated edge enqueues the message into the timer with a callback into the remote edge’s receive handler.

TCP and UDP use IP addresses and ports to locally and, potentially, globally uniquely distinguish themselves and that, more importantly, can be shared with others. In other words, the concept of addresses is key to transports. Since the simulated transports are all running in the same address space, there does not need to be a multilevel naming scheme as provided by IP addresses and ports. Instead, simulated

transports use a single integer, which can then be used as the key into a hashtable, whose value is the node matched to the integer.

When a peer wants to connect to a specific node, instead of connecting to a peer at a remote IP, port pair, it seeks the remote node in the hash table. If no peer exists, depending on the protocol simulated, the result will be a broken link or a connection error. If it finds an entry, the two peers create edges associated with each other. At which point, the peers can easily communicate with each other using the timer and the exchange of buffers.

5.1.3 User Actions

A user in a distributed system does not necessarily imply a human, but rather, an external input from either an application, a user, a sensor, or by some other means. In a simulated environment, these types of behaviors should be properly modeled. That is, if a user requests information from another user using the simulator, it should be delivered when available, not after polling some entry point after some period of time.

To effectively model this behavior requires the use of asynchronous interfaces. After the initiation action is triggered, a registered callback will be triggered upon completion. A synchronous call can be inefficiently turned into an asynchronous call through the use of a thread or by polling. Though for performance purposes, it is best to use an asynchronous interface that only gets invoked upon completion of the task. Fortunately, it is very easy to make synchronous interfaces from asynchronous, so if designed properly, this is not difficult to implement for system designers. If asynchronous handlers are not available, the interface can be made asynchronous through polling at the cost of overhead.

In Brunet, this behavior has been modeled in interactions with the DHT and sending messages through the overlay. A common abstract class contains a method to start the action. It notes the starting time when this method is executed. It then waits for the

asynchronous response from the underlying component to inform that the user action has completed. Optionally, it will call another user-specified callback upon completion.

5.1.4 The Rest of the System

The other components of the system may have an impact on the speed of simulation, but in general should not affect the ability of the system to be simulated. Thus the key to making a system self-simulating is modularity and support for asynchronous interfaces. In the following section, I discuss optimizations that can be made to these components and others to improve simulation.

5.1.5 Optimizations

Simulations can be slow for a number of reasons and that only increases by attempting to simulate software that was not intended to be simulated. Overlay software, for example, typically uses very large addresses (16 bytes or larger) just to represent another node, whereas in a simulation or model this is typically represented as an integer or not at all. Additionally, due to the fact that the lifetime of various buffers in the system can be hard to predict, when interacting with incoming messages and even outgoing messages, many data structures either stay in scope for a long time or there is heavy churn on memory in the heap. For managed languages, this can result in significant overhead due to garbage collection. Finally, since the entire system depends on ordered time, the mechanism ordering timing events plays a key role.

While the typical address in a P2P system may be large in order to allow nodes to obtain addresses independently of each other through random number generation, in a simulation, this large address space is unnecessary, because there is no need to generate the addresses without knowledge of other addresses. One condition that may need large address spaces is extremely large simulations, but given that a 32-bit number allows for 4 billion nodes, this should not be an issue.

Many common data structures are generated in a distributed system even inside a single node. This includes transport addresses, P2P addresses, and common

strings inside the system. By caching these values, the system can reduce its memory consumption and be nicer to the garbage collector. A cache in this sense consists of a hashtable, whose key is the object of interest and the value is a singleton or a value that is identical to the key in every way except but potentially they refer to different locations in memory. Thus when a peer constructs another peer's address, it can check the hashtable for a singleton. If one exists, it uses the singleton and no additional memory is required besides a pointer to this singleton. If one does not exist, this new value is stored as a singleton into the system. There are various means to limiting the entries in a hashtable, such as only keeping the last N entries, keeping track of the last access time, counting the number of references, or using a concept known as weak references. Weak references provide an attractive option as it requires no additional state in the cache, a garbage collector will remove an object when there are no references to it besides weak references. Thus stale entries in the hashtable will return null objects. So a cache using weak references will need to iterate through the entire cache occasionally to remove these stale references.

Messages are usually assembled from a set of memory blocks. Prior to transferring them, they must be placed into a contiguous buffer. Unfortunately, this can lead to significant memory allocations and garbage collections. To address this, I have utilized memory heaps, which can be used to create multiple memory blocks. The concept is to allocate a large memory block. When assembling a message, it is written to an offset into this block. The block can then be shared with others by providing a reference to the memory block and the offset and length of the message inside this block. When the block is no longer in scope, it is garbage collected. This approach significantly reduces dynamic allocation of data and in turn significantly improves the performance of the simulation.

5.2 Efficient Relays

Sometimes NAT traversal using STUN [93] fails due to restrictive firewalls and NATs. Occasionally there are other, harder to diagnose, connectivity issues. Some P2P VPNs [66, 67] support relaying, similar to Traversal Using Relay NAT (TURN) [90] provided by a managed relay infrastructure. Centralized and decentralized VPNs do not suffer from this problem as all traffic passes through the central server or managed links. To address the management and overhead concerns in these systems, I propose the use of distributed, autonomic relaying system based upon previous work [43, 73]. This previous work involved the use of triangular routing that allowed peers next to each other in the node ID space to communicate despite being unable to communicate directly because of firewall, NAT, or Internet fragmentation issues.

The process for forming local relays or "tunnels" [43] begins with two nodes discovering each other via existing peers and determining the need to be connected. If a direct connection attempt fails, the peers exchange neighbor sets through the overlay. Upon receiving this list, the two peers use the overlap in the neighbor sets to form a two-hop connection. In this work, I have further extended this model to support cases when nodes do not have an overlap set. This involves having the peers connect to each other's neighbor sets proactively creating overlap, as represented in Figure 5-1.

Additionally, I have added the feature to exchange arbitrary information along with the neighbor list. Thus far, I have implemented systems that pass information about node stability (measured by the age of a connection) and proximity (based upon ping latency to neighbors). Furthermore, when overlap changes, another mechanism can determine which subset of the peers to use; for example, a peer may only route through the fastest or more stable overlap in the set.

To verify the usefulness of two-hop over overlay routing, I performed experiments and share the results in Section 5.2.1. In a live system, I have verified the accuracy and usefulness of the latency-based relay selection algorithm in Section 5.2.2.




Figure 5-1. Creating relays across the node address space, when direct connectivity is not possible. Two members, 0000 and ABCD, desire a direct connection but are unable to directly connect, perhaps due to NATs or firewalls. They exchange neighbor information through the overlay and connect to one of each other's neighbors, creating an overlap. The overlap then becomes a relay path (represented by dashed lines), improving performance over routing across the entire overlay.

5.2.1 Motivation for Relays in the Overlay

The purpose of this experiment is to quantify the performance benefits of autonomic relays. For this experiment I used the MIT King data set [49], which contains all-to-all latencies between 1,740 well-distributed Internet hosts. Various sizes of networks up to 1,740 nodes were evaluated 100 times each. The experiments were executed by running the Brunet in simulated mode. Once at steady state, I then calculated the average all-to-all latency for all messages that would have taken two overlay hops or more, the average of the low latency relay model, and the average of single hop communication. In the low latency relay model, each destination node form a connection to the source node's physically closest peer as determined via latency (in a live system




Figure 5-2. A comparison of the average all-to-all overlay routing, two-hop relay, and direct connection latency in a Structured P2P environment, Brunet, using the King data set.

by application level ping). Then this pathway is used as a two-hop relay between source and node. I only look at two overlay hops and more, as a single hop would not necessarily benefit from the work and would be the cause of a triangular inequality.

The results are presented in Figure 5-2. The initial starting size for the network was set to 25, because network sizes around 20 and under tend to be fully connected due to the connectivity requirements of the system. It is not until the network size expands past 100 and towards 200 nodes that relays become significantly beneficial. At 100 nodes, there is approximately a 54% performance increase, whereas at 200 there is an 87% increase and it appears to grow proportionately to the size of the pool. The key take away is that latency-bound applications using a reasonably sized overlay would significantly benefit from the use of two-hop relays.

5.2.2 Comparing Relay Selection

In this experiment, I share my experiences of testing the use of latency-aware relays using the public P2P pool running on Planet-Lab as well as Hamachi-Free and Hamachi-Pro relays. Due to Hamachi not supporting relays in Linux, this experiment was performed in Windows Vista 64-bit. Hamachi is discussed in greater depth in

Table 5-1. Results of the evaluation comparing latency and bandwidth of Hamachi relays and IPOP latency-aware autonomic relay selection.

	Latency		Bandwidth	
	(ms)	stdev	Kbit/s	stdev
Hamachi-Free	60.8	2.54	40.2	0.87
Hamachi-Pro	60.2	1.68	1000	1.29
Latency-aware	58.1	35.5	2245	1080

Chapter 2. The testing platform consists of two virtual machine located on the same host with a firewall preventing them from establishing direct connections. All experiments were repeated 5 times using a clean configuration each time. In Hamachi, this meant that the server would need to re-evaluate NAT traversing capabilities and the optimal relay to use. In Brunet, this meant a new node ID and establishing relays with peers in different regions of the overlay. The results are presented in Table 5-1.

As Hamachi was started and figured out that NAT traversal was not possible, it began using multiple different relays as evident by several different ping times. Eventually Hamachi settled on a relay server and it appeared to be the same one every time, for both Hamachi-Free and Hamachi-Pro. The only difference between Hamachi-Pro and Hamachi-Free is that in Pro there is a bandwidth cap of approximately 1 Mbit/s whereas Free is limited to 40 Kbit/s.

Brunet has nodes both on Planet-Lab but also dedicated systems for Archer [37]. These machines are at Universities and thus have a high bandwidth and low latency connection to the testing site. As witnessed by the results, it appears that in most if not all these experiments peers had a low latency connection to a University compute resource and it was chosen ahead of Planet-Lab.

The two take aways are the benefit of being able to dynamically deploy relay servers and reuse compute nodes as relay systems. As the network grows, there may be need to implement some form of bandwidth limit at relay nodes.

5.3 Policies for Establishing Direct Connections

Routing through a ring-structured overlay using a greedy routing algorithm takes $\log(N)$ time and adds $\log(N)$ overall bandwidth for a single message. Therefore, sending messages frequently between two peers through the overlay is not cost effective. What is not apparent through the algorithmic complexity analysis is the fact that many paths in an overlay can be inefficient due to peers routing through distant parts of the world or having limited bandwidth. Less frequently, packets routed via the overlay can just disappear due to nodes disconnecting or packet drops across the Internet. To address this, Ganguly et al. [42] made a system for creating adaptive shortcuts.

Adaptive shortcuts enable peers to establish direct links with each other using Brunet's builtin NAT traversal capabilities. The approach taken by Ganguly et al. was to monitor incoming packets from remote peers and after a certain threshold was passed the system automatically makes a direct connection to the remote peer. As a result of this transparency, software using the overlay could simply start this service without making any additional changes to the application.

5.3.1 Limitations

Unfortunately this approach comes with limitations. There was never any systematic understanding of behaviors that should signify the creation of a direct link or when a direct link should be closed. Thus applying the layer naively could result in connection churn, which would have ramification on the routability of the network. Thus a compromise was made to have it enabled only for selected traffic, which in the case as described by Ganguly et al. [42] was IP or VN traffic.

Two issues made the approach no longer feasible: the increase in size of the overlay network and the securing of IP links. The Internet drops packets at approximately .00835% of the time according to the data provided by iPlane [70]. This compounded with the fact that an overlay message may take $\log(N)$ hops significantly increases the likelihood of a packet dropping before arriving at its destination. Adding a security model

makes this even more complicated, because a trusted link must be established before routing any messages between the end points. In the case of DTLS, this can result in 6 messages traversing the overlay 4-3, prior to the first IP packet. If peers must first establish a trusted link and then transmit a certain amount of packets in a given time, there is a reasonable chance that they may never naturally trigger the creation of an adaptive link.

In practice, it was quite common for this not to succeed and, in fact, security links were often times not even formed. To verify this, I implemented a network profiling tool and deployed it on to PlanetLab. The monitoring tool measured the delay and success of sending messages between the node and every other node in the overlay. The drop rate and latency for a round trip message per hop distance between two peers are presented in Figure 5-4 and Figure 5-3, respectively. The data in those figures is compared to data retrieved from iPlane, which makes it clear that PlanetLab exacerbates the situation. While this is a well-known issue, it is a very important conclusion since many of the public systems provided by my research group rely on PlanetLab.

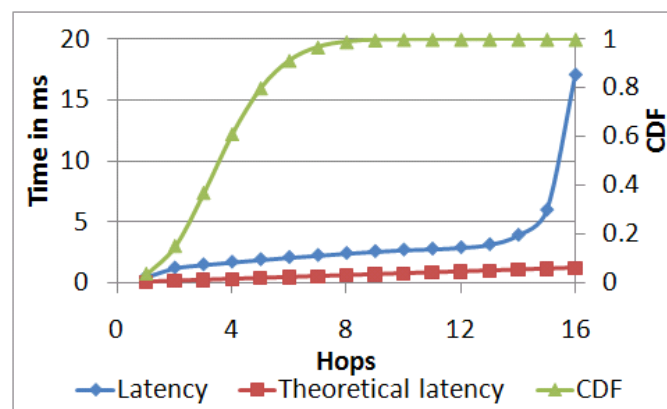


Figure 5-3. Average latency in a round trip message between two nodes specified by the amount of hops between them. The PlanetLab average is compared against data available from iPlane. The CDF for node distance is also plotted for reference.

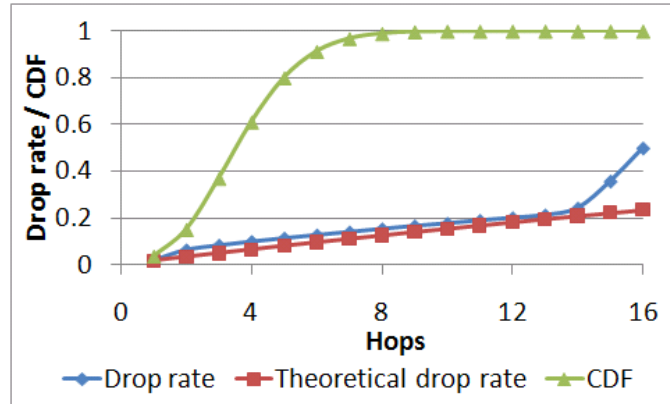


Figure 5-4. Average drop rate in a round trip message between two nodes specified by the amount of hops between them. The PlanetLab average is compared against data available from iPlane. The CDF for node distance is also plotted for reference.

5.3.2 On-Demand Connections

Before defining a new architecture, I measured the network traffic of active and idle applications that were of interest to typical P2P VPNs, Condor [65] and data transfers. Data transfers tend to be simple, if they are TCP driven, first a TCP link must be established, then data transferred, and finally the link is closed. Condor is a job schedule management tool, which is discussed in more depth in Chapter 6. The important aspect for this section is understanding that all nodes in a Condor pool have a relationship with a manager node. Their behavior is to initially send a registration message containing the details of the node and thereafter to send a one-way message stating their presence every 5 minutes or so.

The next step was determining the cost of creating a connection versus routing via the overlay. This cost needs to consider that before a single IP packet can be routed, a security link must be established. In Brunet, the creation of a link requires 1 round trip message across the overlay, whereas the security link, as mention earlier, takes 3. So it is intuitive that sending a single packet secured through an end-to-end channel via a direct link is far more efficient than doing so via the overlay. So instead of having a meter determining when to create connections, connections should be made as soon

as there is interest in communication, in other words, “on-demand.” In a DHT system, this may be done prior to sending or retrieving data from the DHT. In a VPN, this may be during the mapping of IP to P2P as described in Section 2.2.2, which occurs before secure link establishment. Unfortunately, this has the side affect of not being transparent to applications using the P2P software’s interface at the cost of being more responsive.

The creation of on-demand connections results in a higher frequency of connection establishment. As a result, better heuristics are necessary in order to determine when to close unused connections. Using the profiling information retrieved before with regards to Condor’s one-way “heartbeat” messages, it was important that a connection was only closed if it was unused in both directions. Otherwise the manager would be constantly closing connections and peers would randomly disappear from Condor for periods of time. An algorithm that seems to have worked so far is based upon something I call a time-based cache. Initially, entries are stored in a hashtable; after a certain period of time, they are moved to a second hashtable, and those in the second hashtable are lost. If an entry is accessed while in the second hashtable or not at all, it is added to the first hashtable and if applicable removed from the second hashtable. When an entry is removed from the cache, it causes an eviction notice, which results in the connection being removed. The timer is based upon a 7.5 minute timer, so that an inactive connection will be closed within 7.5 to 15 minutes.

The applicability of on-demand connections compared to Chota was evaluated using the simulator with 1,024 nodes and a drop rate of 0.00925% as found on PlanetLab. The On-demand connections were established using the exact semantics of the On-demand protocol; however, the behavior of establishing Chota connections is a little complicated, since the traffic behavior of successful and unsuccessful connection attempts is not identical. To address this, I simulated an ideal Chota situation: the nodes optionally establish a security connection via the overlay, then they exchange a round trip message, and finally establish a direct connection between each other. The

On-demand approach involved optionally creating a security connection followed by the direct connection.




Figure 5-5. CDF of the connection establishment time for secure and security-free Chota and On-demand connections in an overlay of 1,024 nodes and a drop rate of .00925%.

While this evaluation model creates a highly ideal situation for Chota establishment. The results in Figure 5-5 make it clear that Chota is not ideal for this type of application and that security only makes the issues worse. The On-demand connections show a significant improvement, but there still exists an obvious issue with connection establishment that will only be made worse as the system expands. Perhaps using multiple paths on the overlay can improve this situation packet drops on the overlay may have high correlation rather than being uniformly random.

In application, this modification made the Archer, which uses both secure links and Condor, significantly more stable. As the system expanded, nodes were constantly appearing and disappearing, users jobs were being lost due to disconnectivities, and users were complaining about being unable to even submit jobs into the system. Since the change, the issues have been resolved.

5.4 Broadcasting IP Broadcast and Multicast Packets Via the Overlay

The use of a private virtual overlay enables a new method for sending multicast and broadcast packets. In the original approach to IPOP, broadcasting a packet to the entire overlay is not suitable because the overlay could consist of peers from other

Figure 5-6. IP Broadcast using the DHT compared to overlay broadcast

VPNs and those not even involved with VPN operations, while approaches that generate unicast messages when a broadcast or multicast packet arrives at the VPN (e.g. by querying a DHT key where all peers in the VPN would place their overlay address so that they could receive the packets) do not scale well. The abstraction of a private virtual overlay enables scalable broadcasting within a VPN because the only peers in the private overlay are peers for a single VPN. Like the broadcast revocation discussed earlier, IP broadcasting and multicasting use the method described in Appendix 8 to efficiently distribute messages. Though in VPN situations, many peers may already have connections to most if not all of their VPN peers, thus the broadcast algorithm has been modified to allow a peer to select how many peers they would like to forward the message to. Otherwise in many cases, this algorithm will degenerate into one similar to the previous approach.

A validation of the overlay broadcast method for IP broadcast and multicast communication comparing it to the original DHT method is presented in Figure 5-6. Metrics measured include bandwidth used by the packet originator and time required for all peers to receive the packet. System bandwidth remains the same as bounded-broadcast requires exactly N messages to broadcast the message to the entire overlay. By using a broadcast algorithm to route multicast and broadcast IP packets instead of the unicast packets, the sending node has much less used bandwidth. The time required by the DHT approach involves the time to retrieve all entries from the DHT and the latency to the peer farthest away. Whereas the broadcast approach has no DHT lookup and is somewhere around the average latency in the system multiplied by $\log^2(N)$.

5.5 Full Tunnel VPN Operations

The configuration detailed so far describes a split tunnel: a VPN connection that handles *internal VPN traffic only, not Internet traffic*. Prior to this work, only centralized




Figure 5-7. An example of both full and split tunnel VPN modes. In both, packets for the server are sent directly to the server. In split tunnel mode, Internet packets bypass the VPN and are routed directly to the Internet. In full tunnel mode, Internet packets are first routed to the VPN gateway, and then to their Internet destination.

VPNs currently support full tunnel: providing the features of a split tunnel in addition to securely forwarding *all their Internet traffic* through a VPN gateway. A full tunnel provides network-layer privacy when a user is in a remote, insecure location such as an open wireless network at a coffee shop by securely relaying all Internet traffic through a trusted third party, the VPN gateway. Both models are illustrated in Figure 5-7.

Central VPN clients use full tunneling through a routing rule swap, setting the default gateway to be an endpoint in the VPN subnet and traffic for the VPN server is routed explicitly to the LAN gateway. This rule swap causes all Internet packets to be routed to the VN device and the VPN software can then send them to the remote VPN gateway. At the VPN gateway, the packet is decrypted and delivered to the Internet. A P2P system encounters two challenges in supporting full tunnels: 1) P2P traffic must not be routed to the VPN gateway and 2) there may be more than one VPN gateway. I address these issues and provide a solution to this problem in Section 5.5.

The challenges faced in a decentralized P2P VPN are providing decentralized discovery of a VPN gateway and supporting full tunnel mode in a P2P environment such that all P2P traffic is sent to the intended receiver directly instead of through the

gateway. The remainder of this section covers gateway and client solutions to address these challenges.

5.5.1 The Gateway

A gateway can be configured through NAT software, like masquerading in IPtables or Internet Connection Sharing with Windows. This automatically handles the forwarding of packets received on the NAT interface to another interface bringing the packet closer to its destination. Similarly, incoming packets on the outgoing interface must be parsed in order to determine the destination NAT client.

Following from the original design of the VPN state machine in Figure 2-2, if a VPN is a gateway, the VPN state machine no longer rejects packets, when the destination is not in the VPN subnet, though when the VPN gateway mode is disabled these packets are still rejected. When enabled, all Internet and non-VPN based traffic is written to the TAP device setting the destination Ethernet address to the TAP device. The remaining configuration is identical to other members of the system as packets from the Internet will automatically have the clients IP as the destination as a product of the NAT. To provide for dynamic, self-configuring systems, VPN gateways announce their availability via an entry in the DHT. As future work, this approach can be explored to provide intelligent selection and load balancing of gateways.

5.5.2 The Client

VPN Clients wishing to use full tunnel must redirect their default traffic to their VN device. In the prototype VPN model, a virtual IP address is allocated for the purpose of providing distributed VN services DHCP and DNS. This same address is used as the default gateway's IP. Because this IP address never appears in a Internet bound packet, only its Ethernet address does, as shown in Figure 5-8, this approach enables the use of any and multiple remote gateways.

To support full tunnel mode, the VPN's state machine has to be slightly modified to handle outgoing packets destined for IP addresses outside of the VPN, only rejecting

Ethernet Header	Src: Host PN	Dst: LAN Gateway
IP Header	Src: Host PN	Dst: VPN Gateway
Data	Src: Host VN	Dst: Internet
	Data	

Figure 5-8. The contents of a full tunnel Ethernet packet. PN and VN are defined as physical and virtual network, respectively.

them when full tunnel client mode is disabled. When enabled, the VPN software sends packets to the remote peer acting as a full tunnel gateway. Likewise, incoming packets that have a source address outside the subnet should not be rejected but instead the overlay address should be a certified VPN gateway prior to forwarding the packet.

To select a remote gateway, peers query the DHT. As there may be multiple gateways in the system, the peer randomly selects one, forwarding packets to that node. To ensure reliability, when the client has not heard from the gateway recently, the client sends a liveness query to the gateway. If the gateway is down, the taken pessimistic approach finds a new gateway when the next Internet packet arrives.

The real challenge in applying full tunnel VPN mode to P2P VPNs is the nature of the P2P system, namely dynamic connections. Peers do not know ahead of time what remote peer connections will be thus a simple rule switch does not work. The original approach was to watch incoming connection requests and adding additional routing rules on demand, though this is only reasonably feasible with UDP as a TCP handshake message would need to be intercepted and potentially replayed by the local host in order to enable the rule and allow proper routing. The real drawback of the approach though is that UDP messages can easily be spoofed by remote peers enabling unsecured Internet packets to be leaked in the public environment. Even if the connections are secured, it could take some time for the peers to recognize a false connection attempt and delete the rule.

A solution to the security problem is to have all traffic directly routed to the VN device with no additional routing rules. The VN is then responsible for filtering P2P traffic and forwarding it to the LAN's gateway via Ethernet packets. In the VPN application, outgoing IP packets' source ports are compared to VPN application's source ports. Upon a match, the VPN application directs the packet to the LAN's gateway. The three steps involved in this process are 1) translating the source IP address to match the physical Ethernet's IP address, 2) encapsulating the IP packet in an Ethernet packet with a randomly source address [116] and the destination the LAN's gateway, and 3) sending the packet via the physical Ethernet device. Sending an Ethernet packet is not trivial as Windows lacks support for this operation and most Unix systems require administrator privilege. An alternative, platform independent solution uses a second TAP device bridged to the physical Ethernet device, allowing Ethernet packets to be sent indirectly through the Ethernet device via the TAP device. Because the solution results in incoming packets to arrive at a different IP address than the actual original source IP address TCP does not work in this solution. This method has been verified to work on both Linux and Windows using OS dependent TAP devices and bridge utilities.

5.5.3 Full Tunnel Overhead

While the full tunnel client method effectively resolves the lingering problem of ensuring that all packets in a full tunnel will be secure, it raises an issue: could the effect of having all packets traverse the VPN application be prohibitively expensive. Analysis of this approach compares it with one that uses the traditional routing rule switch. Figure 5-2 present the ping time from a residential location to one of Google's IP addresses using a gateway located at the University of Florida when the VPN is in split tunnel mode, full tunnel using the routing rule switch, and full tunnel using Ethernet forwarding. The results express that there is negligible difference between the full tunnel approaches. One interesting result is the latency to gateways public address in the

Table 5-2. Latency results comparing full tunnel approaches measured in ms. Legend:
 GW Pri - gateway's VPN address, GW Pub - gateway's VPN address,
 Ethernet - full tunnel Ethernet packet method, Routing - full tunnel routing rule
 switch, None - split tunnel or no VPN.

	Google	GW Pri	GW Pub
Ethernet	70.6	12.9	13.9
Routing	71.4	13.2	11.0
None	66.1	N/A	10.9

routing test, which most likely is a result of the ping being sent insecurely avoiding the VPN stack completely.

CHAPTER 6

AD-HOC, DECENTRALIZED GRIDS

“Give a man a fish, feed him for a day. Teach a man to fish, feed him for a lifetime” –
Lau Tzu

Large-scale grid computing projects such as TeraGrid and Open Science Grid provide researchers vast amounts of compute resources but with requirements that could limit access, results delayed due to potentially long job queues, and environments and policies that might affect a user’s work flow. In many scenarios and in particular with the advent of Infrastructure as a Service (IaaS) cloud computing, individual users and communities can benefit from less restrictive, dynamic systems that include a combination of local resources and on-demand resources provisioned by one or more IaaS provider. These types of scenarios benefit from flexibility in deploying resources, remote access, and environment configuration.

Grid computing presents opportunities to combine distributed resources to form powerful systems. Due to the challenges in coordinating resource configuration and deployment, researchers tend to either become members of existing grids or deploy their own private resources. The former approach is limited by lack of flexibility in the environment and policies, while the latter requires expertise in systems configuration and management. Though there exists a wealth of middleware available, including resource managers such as Condor [65], Torque (PBS) [84], and Sun Grid Engine [105], many see the cost of installing and managing these systems as being greater than their usefulness and as a result turn to inefficient ad hoc resource discovery and allocation. To combine resources across multiple domains solutions there exist solutions such as the Globus Toolkit [40] or gLite [9]; however, these tool sets come with their own challenges that require the level of expertise most researchers in fields outside of information technology lack.

With the recent advent of cost-effective on-demand computing through Infrastructure as a Service “clouds”, new opportunities for user-deployed grids have arisen; where, for example, a small local computer cluster can be complemented by dynamically provisioned resources that run “cloud-burst” workloads. However, while cloud-provisioned resources solve the problem of on-demand instantiation, the problem of how to *configure* these resources to seamlessly and securely integrate with one’s infrastructure remains a challenge. In particular, considering that users may provision resources from multiple IaaS providers, the configuration demands are similar to a distributed grid: while a cloud image can be encapsulated with a grid computing stack, it still needs configuration in terms of allocating and distributing the appropriate certificates, network configuration to establish end-to-end connectivity, and proper configuration of the middleware to establish worker, submit, and scheduler nodes.

In this chapter, I present techniques that reduce the entry barrier in terms of necessary expertise and time investment in deploying and extending ad hoc, distributed grids. To verify this assertion, I have implemented a system supporting these ideas in the “Grid Appliance,” which as will be demonstrated, allows users to focus on making use of a grid while minimizing their efforts in setting up and managing the underlying components. The core challenges solved by my approach include:

- decentralized directory service for organizing grids,
- decentralized job submission,
- grid single sign on through web services and interfaces,
- sandboxing with network support,
- and all-to-all connectivity despite network asymmetries.

The “Grid Appliance” project and concepts have been actively developed and used in several projects for the past six years. Of these projects, Archer, a distributed grid for computer architecture research, has demonstrated the feasibility and utility




Figure 6-1. The “Grid Appliance” connects to other resources over a common network. Both the grid middleware and the VPN use the P2P overlay to configure and connect the user to other members of the grid. The process uses configuration data provided by the web interface to self-configure the system using information available in the P2P network.

of this approach by deploying a shared collaborative infrastructure spanning clusters across six US universities, where the majority of the nodes are constrained by network address translation (NAT). Every resource in Archer is configured in the same, simple manner: by deploying a “Grid Appliance” that self-configures to join a wide-area grid. Researchers interested or desiring the ability to access both grid resources and specialized commercial simulation tools (such as Simics) can easily use and contribute resources from this shared pool with little effort by joining a website, downloading a configuration image and a virtual machine (VM), and starting the VM inside a VM manager (VMM). Upon completion of the booting process, users are connected to the grid and able to submit and receive jobs.

At the heart of my approach lies a P2P infrastructure based upon a distributed hash table (DHT) useful for decentralized configuration and organization of systems. Peers are able to store key, value pairs into the DHT and to query the DHT with a key and potentially receive multiple values efficiently. The DHT provides discovery and coordination primitives for the configuration of a decentralized P2P virtual private network (VPN), which supports unmodified applications across a network overlay. The

DHT is also used for the decentralized coordination of the grid. Users can configure their grid through a web interface, which outputs configuration files that can be used with the “Grid Appliance.”

The techniques described in this paper have many applications. The basic system supports the creation of local grids by starting a virtual machine on the computers intended for use within the grid and using LAN multicast for discovery. It allows users to seamlessly combine their dedicated grids with external resources such as workstations and cloud resources. The level of familiarity with security, operating systems, and networking is minimal as all the configuration details are handled as components of the system. Management of the system including users and network configuration utilizes a social networking like group interface, while deployment uses pre-built virtual machine images. A graphical overview of the system is illustrated in Figure 6-1.

These techniques simplify the tethering of resources across disparate networks. The setup of security, connectivity, and their continuous management imposes considerable administrative overhead, in particular when networks are constrained by firewalls and NAT devices that prevent direct communication with each other, and which are typically outside the control of a user or lab. Our approach integrates decentralized systems behind NATs in a manner that does not require the setup of exceptions and configuration at NAT/firewall by system administrators.

The rest of the paper is as follows. Section 6.1 highlights of my research groups previous work to provide background for my contributions in this paper. In Section 6.2, I describe the components of the “Grid Appliance” WOW. Section 6.3 provides a case study of a grid deployment using standard grid deployment techniques compared to our “Grid Appliance,” describing qualitatively the benefits and evaluating quantitatively the overheads of this approach. I share my experiences from this long running project in Section 6.4. Finally, Section 6.5 compares and contrasts other solutions to these problems.

6.1 WOWs

This work furthers the vision began by myself and my research lab in earlier described as work wide-area overlay of virtual workstations [42] (WOW). The WOW paper established the use of virtualization technologies, primarily virtual networking and virtual machines, to support dynamic allocation of additional resources in grids that span wide area networks. For reference, the extensions made in this paper to the WOW concept are means for the dynamic creation of grids with support for security, decentralized access, and user-friendly approaches to grid management. This section covers the development of WOWs over the years as it relates to other publications and as means to distinguish the contributions made by me and in this chapter.

6.1.1 P2P Overlays

Peer-to-peer or P2P systems create environments where members have a common functionality. P2P systems are often used for discovery in addition to some user-specific service, such as voice and video with Skype or data sharing with BitTorrent. Many forms of P2P have autonomic features such as self-healing and self-optimization with the ability to support decentralized environments. As I will show, this makes their application in the system very attractive.

For the “Grid Appliance,” I have chosen to use Brunet [13], a type of structured overlay. Structured overlays tend to be used to construct distributed hash tables (DHT) and in comparison to unstructured overlays provide faster guaranteed search times ($O(\log N)$ compared to $O(N)$, where N is the size of the network). The two most successful structured overlays are Kademlia [72], commonly used for decentralized BitTorrent, and Dynamo [28], to support Amazon’s web site and services.

Brunet support for NAT traversal makes it unique from other structured overlays. Originally in the WOWs [42], Brunet facilitated the dynamic connections amongst peers in the grid. Since then, it has been extended to support DHT with atomic operations [44],

efficient relays when direct NAT traversal fails [115], resilient overlay structure and routing [43], and cryptographically secure messaging [115].

6.1.2 Virtual Private Networks

A common question with regards to this work is “why VPNs?” The core reason is connectivity. IPv4 has a limited address space, which has been extended through the use of NAT allowing a single IP to be multiplexed by multiple devices. This creates a problem; however, as it breaks symmetry in the Internet limiting the ability for certain peers to become connected and which peers can initiate connections. With the advent of IPv6, the situation might improve, but there are no guarantees that NATs will disappear nor can users be certain that firewalls will not be in place that inhibit symmetry. A VPN circumvents these issues, so long as the user can connect to the VPN, as all traffic is routed through a successfully connected pathway.

The problem with traditional VPN approaches is management overhead including maintaining resources on public IP addresses and establishing links amongst members in the VPN. The VPN used in the system is called IPOP [41, 115]. IPOP (IP over P2P), as the name implies, uses a P2P overlay (Brunet) to route IP messages. By using P2P, maintaining dedicated bootstrap nodes have less overhead, my approach with IPOP allows an existing Brunet infrastructure to bootstrap independent Brunet infrastructures in order to isolate IPOP networks in their own environments [117].

Once IPOP has entered its unique Brunet overlay, it obtains an IP address. IP address reservation and discovery relies on Brunet’s DHT. Each VPN stores its P2P identifier into the DHT at the generated by the desired IP address, such that the key, value pair is $(hash(IP), P2P)$. In order to ensure there are no conflicts, the storing of this value into the DHT uses an atomic operation, which succeeds only if no other peer has stored a value int $hash(IP)$.

The process for creating connections begins when IPOP receives an outgoing message. First it parses the destination address and queries the DHT for the remote

peers P2P address. The peer then attempts to form a secure, direct connection with the remote peer using Brunet's secure messaging layer. Once that has formed, packets to that IP address are directed over that secure link.

In my original design [113], the virtual network was secured through a kernel-level IPsec stack, a model kept through the first generation Archer deployment. This approach only secures virtual network links between parties and does not secure the P2P layer; furthermore, in IPsec configuration each peer requires a unique rule for every other peer, which limited the maximum number of peers in the VPN. Securing the P2P layer is important, otherwise malicious users could easily derail the entire system, but securing with IPsec would practically negate the benefits of the P2P system, because of network configuration issues related to NATs and firewalls. In modern deployments, I have employed the security layer at the P2P layer, which in turn also secures virtual networking links.

For grids that rely upon VPNs to connect resources and users, this can impose the need for a certificate for the VPN and one for the grid. Though in our approach, I avoid this problem by using a VPN that allows a user to verify the identity of a remote peer and obtain its certificate, and have taken advantage of hooks in grid software that are called to verify a remote peers authenticity. In other words, user access is limited by the VPN and identity inside the grid is maintained by that same certificate. This might not be possible if all users were submitting from the same resources but is feasible in the system since each user submits from their own system.

6.1.3 Virtual Machines in Grid Computing

Earlier work [39] advocated the use of virtual machines (VMs) in grid computing for improved security and customization. Others since [7, 58, 97] have been established VMs as means for sandboxing, that is environments that allow untrusted users to use trusted resources in a limited fashion. VMs run as a process on a system, where processes running inside the VM have no access to the host operating system.

Furthermore, VMs can have limited or no networking access as controlled by the host, which effectively seals them in a cage or sandbox protecting the hosts environment. VMs are also useful for customization and legacy applications, since a developer can configure the VM and then distribute it as an appliance, with the only requirement on the end user being that they have a VM software or manager. Quantitatively, previous work has shown that CPU-bound tasks perform fairly well running with no more than 10% overhead and in some cases 0%, which is the case with VMs like Xen.

While not a direct correlation to grid computing, clouds have benefited significantly from VMs. VMs are the magic behind cloud infrastructures that provide IaaS, such as EC2. In these environments, users are able to create customized instances, or packaged operating systems and applications, inside of cloud environments, share with each other, and dynamically create or shutdown them as necessary. While the application of clouds is generic, it can easily be applied towards grids. A user can create push excess jobs into the cloud, when there is overflow, high demands, or the user does not want to maintain their own hardware. One challenge, however, is the dynamic creation of a grid as well as extension of an existing grid using the cloud, challenges that are addressed in this paper.

6.2 Architectural Overview

My approach attempts to reuse as many available components to design a grid middleware generic enough that the ideas can be applied to other middleware stacks. As a result, my contribution in this chapter and in particular this section focuses primarily on the following key tasks: making grid construction easy, supporting decentralized user access, sandboxing the users environment, limiting access to the grid to authorized identities, and ensuring priority on users own resources.

Table 6-1. Grid Middleware Comparison

	Description	Scalability	Job queue / submission site	API Requirements
Boinc	Volunteer computing, applications ship with Boinc and poll head node for data sets	Not explicitly mentioned, limited by the ability of the scheduler to handle the demands of the client	Each application has a different site, no separation from job queue and submission site	Boinc API and middleware bundling required
BonjourGrid	Desktop grid, use zeroconf / Bonjour to find available resources in a LAN	No bounds tested, limits include multicasting overheads and processing power of job queue node	Each user has their own job queue / submission site	None
Condor	High throughput computing / on demand / desktop / etc / general grid computing	Over 10,000 ¹	Global job queue, no limit on submission sites, submission site communicates directly with worker nodes	Optional API to support job migration and checkpointing
PastryGrid	Use structured overlay Pastry to form decentralized grids	Decentralized, single node limited by its processing power, though collectively limited by the Pastry DHT	Each connected peer maintains its own job queue and submission site	None
PBS / Torque [84]	Traditional approach to dedicated grid computing	up to 20,000 CPUs ²	Global job queue and submission site	None
SGE	Traditional approach to dedicated grid computing	Tested up to 63,000 cores on almost 4,000 hosts ³	Global job queue and submission site	None
XtremWeb	Desktop grid, similar to Condor but uses pull instead of push, like Boinc	Not explicitly mentioned, limited by the ability of the scheduler to handle the demands of clients	Global job queue, separate submission site, optionally one per user	No built-in support for shared file systems

¹ http://www.cs.wisc.edu/condor/CondorWeek2009/condor_presentations/sfiligoi-Condor_WAN_scalability.pdf

² <http://www.clusterresources.com/docs/211>

³ http://www.sun.com/offers/docs/Extreme_Scalability_SGE.pdf

6.2.1 Web Interface and the Community

Before deploying any software or configuring any hardware, a grid needs organization including certificate management, grid access, user account management, and delegation of responsibilities. These are complex questions, which can be challenging to address, though for less restrictive systems, like a collection of academic labs sharing clusters, they may be very easy. One of the professors could handle the initial authorization of all the other labs and then delegate to them the responsibility of allowing their affiliates, such as students and scholars access.

For academic environments, grids become more challenging when the professor or worse yet students must maintain the certificates, handling certificate requests, and placing signed certificates in the correct location. Our solution to this potentially confusing area was a group interface, akin to something like Facebook's or Google's groups. Albeit, those types of groups are not hierarchal, which is a necessity in order to have delegated responsibilities. Thus I have a two layer approach, a grid group for members of the grid trusted by the grid organizers and user groups for those who are trusted by those in the grid group. Members of the grid group can create their own user groups. A member of a user group can gain access to the grid by downloading grid configuration data available within the user group web interface. This configuration data comes in the format of a disk image, when added to a "Grid Appliance" VM, it is used to obtain the user's credentials and enabling them to connect to the grid.

To give an example, consider the computer architecture grid, Archer. Archer was seeded initially by the University of Florida, so my group and I are the founders and maintainers of the Archer grid group. As new universities and independent researchers have joined Archer, they request access to this group. Upon receiving approval, they then need to form their own user group so that they can allow others to connect to the grid. So a trusted member might create a user group titled "Archer for University X" and all members of university X will apply for membership in that group. The creator can

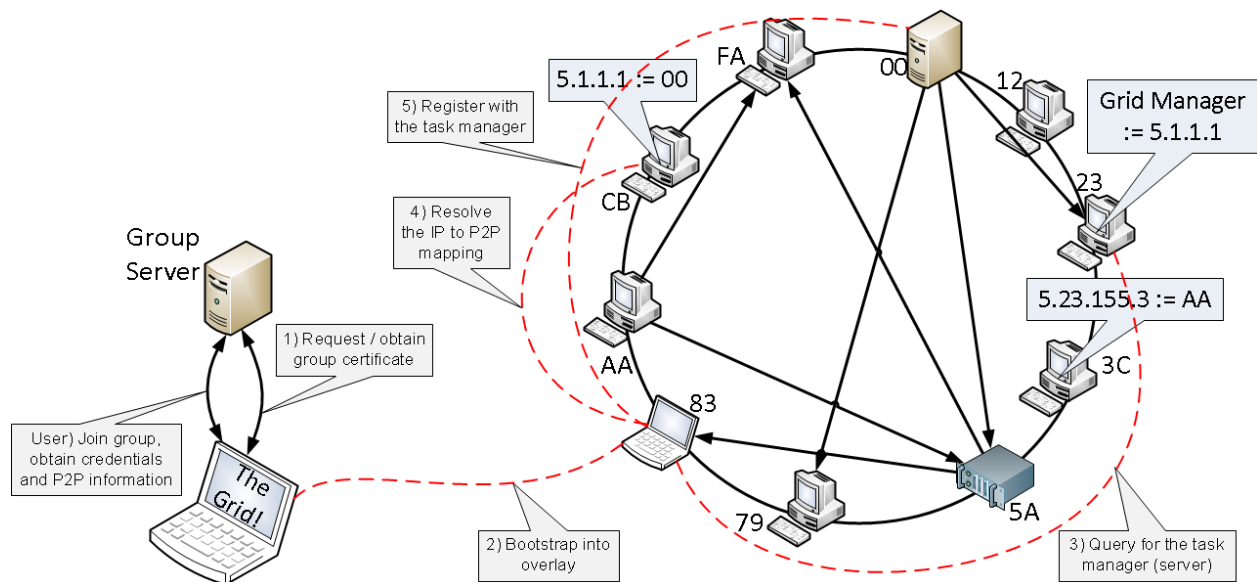


Figure 6-2. An example deployment scenario: obtaining configuration files, starting the appliance, and connecting with a resource manager.

make decisions to either accept or deny these users. Once the user has access, they will download their configuration data formatted as a virtual disk image and the “Grid Appliance” VM and start the “VM.” After starting the VM, the user will be connected to the grid and able to submit and receive jobs.

Joining is easy; a grid requires a user to sign onto a website and download a configuration data, which can then be used on multiple systems. To support this process, the configuration data contains cryptographic information that facilitates acquisition of a signed certificate from the web interface through XML-RPC over HTTPS. The process begins by either booting the “Grid Appliance” or restarting a “Grid Appliance” service. When starting the service will detect if there is new configuration data, and if there is, it contacts the web interface with the cryptographic information and a public key. The web interface verifies the user’s identity, retrieves their profile from its database and binds that information with the public key to create a certificate request, which will then be signed and returned to the user.

With a public web interface, I have been able to create a variety communities. One of particular interest is not the grid itself but rather a bootstrapping community for grids.

The web interface has been designed to support many grid groups, so too has the P2P infrastructure as it supports bootstrapping into unique private overlays for individual grids by means of Brunet's ability to support recursive bootstrapping. By using the public interface, users have an opportunity to reuse a public bootstrap infrastructure and only need to focus on the configuration of their VPN and grid services, which has been trivialized to accepting or denying users access to a group and turning on resources. We would like to note that there is no need to make an explicit public grid community through the web interface, since all "Grid Appliances" come with a default configuration file that will connect them to an insecure public grid.

6.2.2 The Organization of the Grid

The previous section focused facilitation of grid configuration using the web interface and skirted the issues of detailed configuration and organization. The configuration of the grid mirrors that of the connection process. The first tier group maps to a common grid and each grid maps to a VPN. Thus when a user creates a new grid group, they are actually configuring a new VPN, which involves address range, security parameters, user agreements, and the name of the group. The system provides defaults for address range and security parameters, so users can focus on high level details like the user agreement and the grid's name.

As mentioned earlier, the second tier of groups enables members in the grid group to provide access to their community. It is also the location that users download their configuration data. The configuration files come in three flavors: submission, worker, or manager. Worker nodes strictly run jobs. Submission nodes can run jobs as well as submit jobs into the grid. Manager nodes are akin to head nodes, those that manage the interaction between worker and submission nodes.

While the configuration details are handled by the web interface and scripts inside the "Grid Appliance," organization of the grid, more specifically the linking of worker and submission nodes to manager nodes, relies on the DHT. Managers store their IP

addresses into the DHT at the key *managers*. When workers and clients join the grid, they automatically query this key, using the results to configure their grid software. Managers can also query this key to learn of other managers to coordinate with each other.

6.2.2.1 Selecting a Middleware

My grid composition is largely based upon a desire to support a decentralized environment, while still retaining reliability and limiting documentation support efforts. As there exist many middlewares to support job submission and scheduling, I surveyed available and established middleware to determine how well they matched my requirements. My results are presented in Table 6-1, which covers most of the well established middleware and some recent research projects focused on decentralized organization.

Of the resource management middlewares surveyed, I chose to use Condor as it matches closest with my goals due to its decentralized properties and focus on desktop grids. Condor allows multiple submission points, a non-trivial obstacle in some of the other systems. Additionally, adding and removing resources in Condor can be done without any configuration from the managers. Conversely, in SGE and Torque, after resources have been added into the system, the administrator must manually configure the manager to control them. Most scheduling software assumes that resources are dedicated, while Condor supports opportunistic cycles, by detecting the presence of other entities and will suspend, migrate, or terminate a job, thus enabling desktop grids. A common drawback to established middlewares is the requirement of a manager node; having no manager in an ad hoc grid would be ideal.

6.2.2.2 Self-Organizing Condor

While the requirement of a central manager may be undesirable, they can easily be run inside a VM and Condor supports the ability to run many in parallel through the use of “flocking [33].” Flocking allows submission sites to connect to multiple managers.

This serves two purposes: 1) to provide transparent reliability by supporting multiple managers and 2) users can share their resources through their own manager. Flocking allows each site to run its own manager or share the common manager.

To configure Condor, manager IP addresses are stored into the DHT using the key *managers*. Joining peers query the DHT to obtain a list of managers, selecting one randomly to use as its primary manager with the result used for flocking. If the system prefers managers from its group, it will randomly contact each manager in an attempt to find a match, selecting one at random if no match is found. Until a manager is found, the process repeats every 60 seconds. Upon finding a manager, the state of the system is verified every 10 minutes and new managers are added to the flock list.

6.2.2.3 Putting It All Together

The following summarizes the configuration and organization of the grid. Minimally a grid will constitute a manager, some workers, and a submitter. Referencing Figure 6-2 step “1,” during system boot, without user interaction, each machine contacts the group website to obtain a valid VPN certificate. Whereupon, it connects to the P2P overlay whose bootstrap peers are listed inside the configuration file, “step 2.” At which point, the machine starts the VPN service running on top of the P2P overlay, also part of step “2.” The self-configuring VPN creates a transparent layer hiding from the user and administrators the complexity in setting up a common fabric that can handle potential network dynamics. Machines automatically obtain a unique IP address and find their place inside the grid. For a manager machine, this means registering in the DHT (not shown), while clients and workers search for available managers by querying the DHT, step “3,” IPOP translates the IP to a P2P address, step “4;” and then client contacts the manager directly, step “5.”

6.2.3 Sandboxing Resources

As tasks can run on worker and potentially submission nodes, I have devised means to sandbox the environments that do not limit user interactions with the system.

While more traditional approaches to sandboxing emphasize a separation between worker and submission machine, in actual deployments, very few users explicitly deploy worker machines, most are submission machines. Thus I developed sandboxing techniques to limit the ability of submitted jobs on systems that are simultaneously being used for submission. So these sandboxing technique considers more than just locking down the machine but also ensuring a reasonable level of access.

6.2.3.1 Securing the Resources

The core of my sandboxing approach is to limit attacks to software in the system and not poorly configured user space, such as poorly chosen passwords or resources external to the “Grid Appliance.” All jobs are run as a set of predefined user identities. When the jobs are finished executing, whether forcibly shutdown or completed successfully, all processes from that user are shutdown, preventing malicious trojan attacks. Those users only have access to the working directory for the job and those with permission for everybody. Escalation of privilege attacks due to poor passwords are prevented by disallowing use of “su” or “sudo” for these users. Finally, network access is limited to the VPN, thus they are unable to perform denial of service attacks on the Internet.

Additionally, systems can be configured such that the only network presented to them is that of the virtual network. To support this, IPOP has been enhanced to support a router mode, which can be bridged to a virtual machine adapter running on the host machine that connects to the network device running inside the VM. Not only does this improve performance, due to reduced I/O overhead, the same virtual network router can be used for multiple VMs.

To ensure that submit machines still have a high level of functionality without risking the system to external attacks even from users on the same network, user services are run only on a “host-only” network device within the virtual machine. This includes an SSH server and a Samba or Windows File Share. The user name matches that from the

website, while the password defaults to “password.” I would like to note that file sharing services work the opposite to that of host to guest as most VMs already have in place. Instead users can access their files on the VM from the host. This was done to limit potential attacks on submission machine.

6.2.3.2 Respecting the Host

Another aspect of sandboxing is respecting the usage of the host. While Condor can detect host usage on a machine it is running, when run inside a VM it cannot detect usage on the host. Thus it is imperative to support such a configuration otherwise my approach would be limited in that it can only be run during idle times. In the “Grid Appliance”, this is addressed by running a light-weight agent on the host that communicates to the VM through the second Ethernet interface. The agent discovers a VM through multicast service discovery executed only on “host-only” virtual network devices. When a user accesses the host, the agent notifies a service in the VM, which results in running tasks being suspended, migrated, or terminated. The machine remains off limits until there has been no user activity for 10 minutes.

6.2.3.3 Decentralized Submission of Jobs

From the administrator’s perspective, not requiring a submission machine is also a form of sandboxing. Maintaining a worker machine requires very low overhead, since jobs and their associated files are removed upon the completion of a job and corrupted workers can be deleted and redeployed. Maintaining a submission machine means user accounts, network access, providing data storage, and trusting users to play nicely on a shared resource. So having users be able to submit from their own resources reduces the overhead in managing a grid. It does come with a consequence, most grids provide shared file systems, which are statically mounted in all nodes. In a dynamic grid that might have multiple shares, this type of approach may not be very feasible.

All is not lost, for example, Condor provides data distribution mechanisms for submitted jobs. This can be an inconvenience, however, if only a portion of the file is

necessary, as the entire file must be distributed to each worker. This can be particularly true with disk images used by computer architecture simulations and applications built with many modules or documentation. To support sparse data transfers and simplify access to local data, each “Grid Appliance” has a local NFS share exported with read-only permission. To address the issue of mounting a file system, there exists a tool to automatically mount file systems, autofs. autofs tool works by intercepting file system calls inside a specific directory, parsing the path, and mounting a remote file system. In the “Grid Appliance,” accessing the path `/mnt/ganfs/hostname`, where hostname is either the IP address or hostname of an appliance, will automatically mount that appliance’s NFS export without the need for super-user intervention. Mounts are automatically unmounted after a sufficient period of time without any access to the mounted file system.

6.3 Deploying a Campus Grid

I now present a case study exploring a qualitative and quantitative comparison in deploying a campus grid and extending it into the “Cloud” using traditional techniques versus a grid constructed by “Grid Appliance.” One of the target environments for the “Grid Appliance” is resources provided in distributed computer labs and many small distributed clusters on one or more university campus as shown in Figure 6-3. The goals in both these cases are to use commodity software, where available, and to provide a solution that is both simple but creates an adequate grid. In both cases, Condor is chosen as the middleware, which is a push scheduler and by default requires that all resources be on a common network thus a VPN will be utilized. Additionally, in this section, I cover details of the “Grid Appliance” that did not fit in the context of previous discussions in the paper.

6.3.1 Background

In this case study, I will compare and contrast the construction of two types of grids: a static grid configured by hand and a dynamic grid configured by the “Grid Appliance.”




Figure 6-3. A collection of various computing resources at a typical university.

Each grid is initially constructed using resources at the University of Florida and later extended to Amazon’s EC2 and Future Grid at India University using Eucalyptus. Each environment has a NAT limiting symmetric communication: University of Florida resources are behind two layers, first an “iptables” NAT and then a Cisco NAT; EC2 resources have a simple 1:1 NAT; and the Eucalyptus resources appear to have an “iptables” NAT.

6.3.2 Traditional Configuration of a Campus Grid

A VPN must be used to connect the resources due to the lack of network symmetry across the sites. There exists a wealth of VPNs available [67, 100, 120] and some explicitly for grids [53, 106, 108]. For simplicity sake, OpenVPN was chosen due to the simplicity in its configuration. In reality, OpenVPN makes a poor choice because it is centralized, thus all traffic between submitter and worker must traverse the VPNs server. Whereas others in the list are distributed and thus allow nodes to communicate directly, but in order to do so, manual setup is required, a process, that would overwhelm many

novice grid deployers. In all these cases, the VPN requires that at least a single node have a public address, thus I had to make a single concession in the design of this grid, that is, the OpenVPN server runs on a public node.

In order to connect to OpenVPN, it must know the server's address and have a signed certificate. While typically, most administrators would want a unique private key for each machine joining the grid, in my case study and evaluation, I avoided this process and used a common key, certificate pair. In doing so, there are potential dangers, for example, if any of the machines were hijacked, the certificate would have to be revoked and all machines would be rendered inoperable. To create a properly secured environment, each resource would have to generate or be provided a private key, a certificate request submitted to the certificate authority, and a signed certificate provided to the resource.

With the networking and security components in place, the next step is configuring grid middleware. Prior to deploying any resources, the manager must be allocated and its IP address provided to other resources in the system. Submission points are not a focus on this case study, though in general most systems of this nature have a single shared submission site. The challenges in supporting multiple submission points in this environment include creating certificates same as worker nodes, requiring users to configure OpenVPN and Condor, and handling NFS mounts. Whereas having a single submission point creates more work for the system administrator as mentioned earlier. Both approaches have their associated costs and neither is trivial. The evaluation assumes a single user submitting from a single resource.

To address potential heterogeneity issues. An administrator would need to collaborate with others to ensure that all resources are running a common set of tools and libraries. Otherwise an application that works well on one platform could cause a segmentation fault on another, through no fault of the user, but rather due to library incompatibilities.

To export this system into various clouds, an administrator starts by running an instance that contains their desired Linux distribution and then installing the grid utilities like Condor and OpenVPN. Supporting individualization of the resources is challenging. The simplest approach is to store all the configuration in that instance including the single private key, certificate pair as well as the IP address of the manager node. Alternatively, the administrator could build an infrastructure that receives certificate requests and returns a certificate. The IP address of the manager node and of the certificate request handler could be provided to the cloud via user data, a feature common to most IaaS clouds that allows users to provide either text or binary data that is available via a private URL inside a cloud instance.

6.3.3 Grid Appliance in a Campus Grid

All these configuration issues are exactly the reasons why “Grid Appliance” and its associated group Web interface are desirable for small and medium scale grids. The first component is deciding which web interface to use, public (www.grid-appliance.org) or private hosted on their own resources. Similarly, users can deploy their own P2P overlay or use the shared overlay.

The web interface enforces unique names for both the users and the groups. Once the user has membership in the second tier of groups, they can download a file that will be used to automatically configure their resources. As mentioned earlier, this handled obtaining a unique signed certificate, connecting to the VPN, and discovering the manager in the grid. Configuration of the VPN and grid are handled seamlessly, the VPN automatically establishes direct links with peers on demand and peers configure based upon information available in the P2P overlay dynamically allowing for configuration changes.

Heterogeneity is a problem that will always exist if individuals are given governance of their own resources. Rather than fight that process, the “Grid Appliance” approach is to provide a reference system and then include that version and additional programs

in the resource description exported by Condor. Thus a user looking for a specific application, library, or computer architecture can specify that in their job description. Additionally, by means of the transparent NFS mounts, users can easily compile their own applications and libraries and export them to remote worker nodes.

Extending the “Grid Appliance” system into the clouds is easy. The similarity between a VM appliance and a cloud instance are striking. The only difference from the perspective of the “Grid Appliance” system is where to check for configuration data. Once a user has created a “Grid Appliance” in a cloud, everyone else can reuse it and just supply their configuration data as the user data during the instantiation of the instances. As I describe in Section 6.4.2, creating “Grid Appliance” from scratch is a trivial procedure.

As described in detail earlier, an administrator needs to install the necessary software either by deploying VMMs and VM appliances or installing “Grid Appliance” packages on Debian / Ubuntu systems. Additionally, these systems need to be packaged with the configuration files or floppy disk images. At which point, the systems will automatically configure and connect to the grid. An administrator can verify this by monitoring Condor. Additional resources can be added seamlessly, likewise resources can be removed by shutting them off without direct interaction with the “Grid Appliance” or manager node.

6.3.4 Comparing the User Experience

In the case of a traditional grid, most users will contact the administrator and make a request for an account. Upon receiving confirmation, the user will have the ability to SSH into a submission site. Their connectivity to the system is instantaneous, their jobs will begin executing as soon as it is their turn in the queue. User’s will most likely have access to a global NFS. From the user’s perspective, the traditional approach is very easy and straightforward.

With the “Grid Appliance,” a user will obtain an account at the web interface, download a VM and a configuration file, and start the VM. Upon booting, the user will be able to submit and receive jobs. To access the grid, users can either SSH into the machine or use the consoles in the VM. While there is no single, global NFS, each user has their own unique NFS and must make their job submission files contain their unique path. For the most part, the user’s perspective of the “Grid Appliance” approach has much of the same feel as the traditional approach. Although users have additional features such as accessing their files via Samba and having a portable environment for doing their software development.

6.3.5 Quantifying the Experience

The evaluation of these environments focuses on the time taken to dynamically allocate the resources, connect to the grid, and submit a simple job to all resources in the grid. In both systems, a single manager and submission node were instantiated in separate VMs. In the traditional setup, OpenVPN is run from the manager node. Each component in the evaluation was run three times. Between iterations, the submission node and the manager node were restarted to clear any state.

The times measured include the time from when the last grid resource was started to the time it reported to the manager node, Figure 6-4, as well as the time required for the submit node to queue and run a 5 minute job on all the connected workers, Figure 6-5. The purpose of the second test is to measure the time it takes for a submission site to queue a task to all workers, connect to the workers, submit the job, and to receive the results; thus a stress test on the VPN’s ability to dynamically create links and verifying all-to-all connectivity. The tests were run on 50 resources (virtual machines / cloud instances) in each environment and then on a grid consisting of all 150 resources with 50 at each site.

In the previous section, I qualified why the approach was easier than configuring a grid by hand, though by doing so I introduce overheads related to configuration and




Figure 6-4. Comparison of times to construct a grid in various environments using both a statically configured grid and a grid constructed by the “Grid Appliance.” Legend: EC2 - Amazon’s EC2, Euca - Indiana University’s Eucalyptus, UF - University of Florida’s ACIS Lab resources, Static - OpenVPN, GA - Grid Appliance.




Figure 6-5. Comparison of times to run a 300 second job on each resource in various grids configured statically and through the “Grid Appliance.” Legend: EC2 - Amazon’s EC2, Euca - Indiana University’s Eucalyptus, UF - University of Florida’s ACIS Lab resources, Static - OpenVPN, GA - Grid Appliance.

organization. The evaluation verifies that these overheads do not conflict with the utility of my approach. Not only do resources within a cluster install the VMs and connect to the grid quickly, the clouds do as well. While the results were similar, it should be noted that the time required to configure the static approach was not taken into effect. A process that is difficult to measure and is largely reliant on the ability of the administrator and the tools used. Whereas the time for the “Grid Appliance” does include many of these components.

It should be stated that the evaluation only has a single submission node. In a system with multiple submitters, the OpenVPN server could easily become a bandwidth bottleneck in the system as all data must pass through it, which can be avoided using IPOP. Additionally, the current “Grid Appliance” relies on polling with long delays, so as to not have negative effects on the system. Either shrinking those times or moving to an event based system should significantly improve the speed at which connectivity occurs.

6.4 Lessons Learned

This section highlights some the interesting developments and experiences, we have had that do not fit the topics discussed so far.

6.4.1 Deployments

A significant component of my experience stems from the computational grid provided by Archer [37], an active grid deployed for computer architecture research, which has been online for over 3 years. Archer currently spans six seed universities contributing over 600 CPUs as well as contributions and activities from external users. The Archer grid has been accessed by hundreds of students and researchers from over a dozen institutions submitting jobs totaling over 500,000 hours of job execution in the past two years alone.

The Grid Appliance has also been utilized by groups at the Universities of Florida, Clemson, Arkansas, and Northwestern Switzerland as a tool for teaching grid computing. Meanwhile the universities of Clemson and Purdue are using the Grid Appliance’s VPN (GroupVPN / IPOP) to create their own grid systems. Over time, there have been many private, small-scale systems using the shared system available at www.grid-appliance.org with other groups constructing their own independent systems. Feedback from users through surveys have shown that non-expert users are able to connect to the public Grid appliance pool in a matter of minutes by simply downloading and booting a plug-and-play VM image that is portable across VMware, VirtualBox, and KVM.

6.4.2 Towards Unvirtualized Environments

Because of the demands put on Archer in terms of avoiding the overheads of virtualization and the perceived simplicity of managing physical resources as opposed to virtual resources running on top of a physical resources, many users have requested the ability to run Grid Appliances directly on their machine. Unlike clouds with machine images such as AMIs or VM appliances, physical machines images cannot be easily exported. Most physical OS installed on physical machines will need some some custom tailoring to handle environment specific issues.

With this in mind, I moved away from stackable file systems and towards creating repositories with installable packages, such as DEB or RPM. The implications of packages mean that users can easily produce “Grid Appliances” from installed systems or during system installation. With the VPN router mode, mentioned earlier, resources in a LAN can communicate directly with each other rather than through the VPN. That means if they are on a gigabit network, they can full network speeds as opposed to being limited to 20% of that due to the VPN, overheads discussed in [116].

6.4.3 Advantages and Challenges of the Cloud

I have had the experience of deploying the “Grid Appliance” on three different cloud stacks: Amazon’s EC2 [5], Future Grid’s Eucalyptus [76], and Future Grid’s Nimbus [60]. All of the systems, encountered so far, allow for data to be uploaded with each cloud instance started. The instance can then download the data from a static URL only accessible from within the instance, for example, EC2 user data is accessible at <http://169.254.169.254/latest/user-data>. A “Grid Appliance” cloud instances can be configured via user-data, which is the same configuration data used as the virtual and physical machines, albeit zip compressed. The “Grid Appliance” seeks the configuration data by first checking for a physical floppy disk, then in specific directory (/opt/grid_appliance/var/floppy.img), followed by the EC2 / Eucalyptus URL, and finally the Nimbus URL. Upon finding a floppy and mounting it, the system continues on

with configuration. Clouds have been also very useful for debugging. Though Amazon is not free, with Future Grid, grid researchers now have free access to both Eucalyptus and Nimbus clouds. Many bugs can be difficult to reproduce in small system tests or booting one system at a time. By starting many instances simultaneously, I have been able to quickly reproduce problems and isolate them, leading to timely resolutions, and verification of those fixes.

Beyond the use of extending into clouds for on-demand resources, they are also very convenient for debugging. Doing so on Amazon though is not free. Fortunately, grid researchers now can have free access to Future Grid with both Eucalyptus and Nimbus style clouds. I did have to do some tinkering to get these systems to work. First, because the user data is binary data and the communication exchange uses RPC, which may have difficulty handling binary data, it must be converted to base64 before transferring and converted back into binary data afterward. EC2 handles this transparently, if using command-line tools. Unfortunately, Eucalyptus and Nimbus do not, even though Eucalyptus is supposed to be compatible with EC2.

Furthermore, when starting an EC2 instance, networking is immediately available, whereas with Eucalyptus and Nimbus, networking often times takes more than 10 seconds after starting to be available. Thus a startup script must be prepared for networking not to be ready and hence unable to immediately download user data. The best approach to deal with this in a distribution independent manner is to wait until the primary Ethernet interface (eth0) has an IP and then continuing.

6.4.4 Stacked File Systems

Configuring systems can be difficult, which makes it important to have the ability to share the resulting system with others. The approach of actually creating packages can be overly complicated for novices. To address this concern, the original “Grid Appliance” supported a built-in mechanism to create packages through a stackable file system using copy-on-write, as describe in [113]. In this environment, the VM used 3 disks:




Figure 6-6. Example of a stackable file system from the previous “Grid Appliance.” A file will be read from the top most file system in the stack and all writes are directed to Home.

the “Grid Appliance” base image, the software stack configured by us; a module; and a home disk. In normal usage, both the base and module images are treated as read-only file systems with all user changes to the system being recorded by the home image, as depicted in Figure 6-6.

To upgrade the system, users replaced their current base image with a newer one, while keeping their module and home disks. While the purpose of the module was to allow users to extend the configuration of the “Grid Appliance.” To configure a module the system would be booted into developer mode, an option during the boot phase, where only the base and module images are included in the stacked file system. Upon completing the changes, a user would run a script that would clean the system and prepare it for sharing. A user could then share the resulting module image with others.

Issues with this approach made it unattractive to continue using. First, there exists no kernel level support for stackable file systems, I had to add UnionFS [118] to the kernel, adding the weight of maintaining a kernel unto my shoulders. While FUSE (filesystem in userspace) solutions exist, they require modifications to the initial ram disk, which is reproduced automatically during the installation of every new kernel,

furthermore, our experience with them suggests they are not well suited for production systems. Additionally, the approach was not portable to clouds or physical resources. So while I have deprecated the feature for now, I see it as a potential means to easily develop packages like DEB and RPM.

6.4.5 Priority in Owned Resources

In Archer, seed universities should have priority on the resources at their university. Similarly, users should have priority on their contributions. Otherwise, users will remove their resources from the grid, when they want guaranteed access. To support user and group based priorities, Condor has mechanisms that can be enforced at the server that allow for arbitrary means to specify user priority for a specific resource. So the configuration specifies that if the resource's user or group matches that of the submitter, the priority is higher than otherwise. This alone is not sufficient as malicious users could easily tweak their user name or group to obtain priority on all resources. Thus whenever this check is made the user's identity in the submission information is verified against their P2P VPN certificate. Failed matches are not scheduled and are stored in a log at the manager for the administrator to deal with later.

To support this behavior, the following statements have been added to the respective system's Condor configuration file:

The format for this configuration is as follows:

Job queue (server):

```
NEGOTIATOR_PRE_JOB_RANK = 10 * (MY.RANK)
```

Worker:

```
GROUP_RANK = TARGET.Group == MY.Group
```

```
USER_RANK = TARGET.User == My.User
```

```
RANK = GROUP_RANK || USER_RANK
```

Worker and Submitter:

```
Group = "Group's Name"
```

User = "User's Name"

6.4.6 Timing in Virtual Machines

Certain applications, particularly license servers, are sensitive to time. Because of the nature of grids, there exist possibilities of having uncoordinated timing, such as improperly specifying the time zone or not using a network time protocol (NTP) server. With regards to VMs, VMWare [112] suggests synchronizing with the host's time and to avoid using services like NTP, which may have adverse affects on timing inside the virtual machine. While NTP might have some strange behavior, relying on host time may produce erratic jumps in time that some software cannot handle. My experiences recommends the use of NTP to address these concerns, which has resolved many issues with strange software behavior and frustration from users when their jobs fail due to being unable to obtain a license due to a timing mismatch.

6.4.7 Selecting a VPN IP Address Range

One challenge in deploying a VPN is ensuring that the address space does not overlap with that over the environments where it will be used. If there is overlap, users will be unable to connect to the VPN. Doing so will confuse the network stack, as there will be two network interfaces connected to the same address space but different networks. A guaranteed, though not necessarily practical solution is to run the resource on a VM NAT or a cluster NAT that does not overlap the IP address space of the VPN.

Users of the "Grid Appliance" should not have to concern themselves with this issues. Prior work on the topic by Ala Rezmerita et al. [85] recommends using the experimental address class E ranging between 240.0.0.0 - 255.255.255.254, unfortunately this requires Linux kernel modifications. With the amount of bugs and security fixes regularly pushed into the kernel, maintaining a forked kernel requires a significant amount of time, duplicating the work already being performed by the OS distribution maintainers. This would also limit the ability to easily deploy resources in physical and cloud environments. Additionally, users that wanted to multipurpose a

physical resource may not want to run a modified kernel, while in most cloud setups the kernel choice is limited.

I have since moved towards using the 5.0.0.0 -

5.255.255.255 address range. Like the class E address space it is unallocated, but it requires no changes to any operating systems. The only limitation is that some other VPNs also use it, thus a user would not be able to run two VPNs on the same address space concurrently. This approach is much better than providing kernels or dealing with network address overlaps. Interestingly, even with this in place, we still see some “GroupVPNs” using address ranges in normal private network address ranges for the VPN, like 10.0.0.0 - 10.255.255.255 and 192.168.0.0 - 192.168.255.255.

6.4.8 Administrator Backdoor

While most administrators will agree that most problems that users encounter are self-inflicted, there are times, when the system is at fault. Debugging systems faults in a decentralized system can be very tricky, since it is very difficult to track down a resource in order to gain direct physical access. Additionally, having a user bring their resource to an administrator may be prohibitively complicated, as the user would need to relocate their “Grid Appliance” instance and have network connectivity in order to connect to the grid and show the problem to the administrator. To address this and other concerns that only appear after running the system for long periods of time, we have supplied an administrator backdoor into all resources by installing our public ssh key, though users are informed of this and are free to remove it for privacy concerns. In typical configurations, this approach might not be feasible, but because the “Grid Appliance” ships with a decentralized VPN supporting all-to-all connectivity, any resource connected to the VPN is accessible for remote debugging by an administrator. Most users involved are extremely delighted with the process as it has an appearance that the system “just works.”

6.5 Related Work

Existing work that falls under the general area of desktop grids/opportunistic computing include Boinc [6], BonjourGrid [2], and PVC [85]. Boinc, used by many “@home” solutions, focuses on adding execute nodes easy; however, job submission and management rely on centralization and all tasks must use the Boinc APIs. BonjourGrid removes the need for centralization through the use of multicast resource discovery; the need for which limits its applicability to local area networks. PVC enables distributed, wide-area systems with decentralized job submission and execution through the use of VPNs, but relies on centralized VPN and resource management.

Each approach addresses a unique challenge in grid computing, but none addresses the challenge presented as a whole: easily constructing distributed, cross-domain grids. Challenges that I consider in the design of my system include allowing submission sites to exist any where without being confined to complex configuration or highly available, centralized locations; the ability to dynamically add and remove resources by starting and stopping a a resource; and the sharing of common servers so that no group in the grid is dependent on another. I emphasize these points, while still retaining the ease of use of Boinc, the connectivity of PVC, and the flexibility of BonjourGrid. The end result is a system similar to OurGrid [8]; however, OurGrid requires manual configuration of the grid and networking amongst sites, administration of users within a site, and limits network connectivity amongst resources, whereas “Grid Appliance” transparently handles these issues with a P2P overlay and VPN to handle network constraints and support network sandboxing and a web interface to configure and manage the grid.

With regards to clouds, there exists contextualization [59]. Users construct an XML configuration file that describes how a cloud instance should be configured and provide this to a broker. During booting of a cloud instance, it will contact a third-party contextualization broker to receive this file and configure the system. This approach

has been leveraged to create dynamic grids inside the Nimbus cloud [51]. While this approach can reproduce similar features of the “Grid Appliance,” such as creating grids inside the cloud, there are challenges in addressing cloud bursting, automated signing of certificates, and collaboration amongst disparate groups.

CHAPTER 7

SOCIAL PROFILE OVERLAYS

Online social networking has become pervasive in daily life, though as social networks grow so does the wealth of personal information that they store. Once information has been released on a social network, known as a user's profile, the user and the data are at the mercy of the terms dictated by the social network infrastructure, which today is typically third-party, centrally owned. If the social network engages in activities disagreeable to the user, due to change of terms or opt-out programs not well understood by users such as recent issues with Facebook's Beacon program [77], the options presented to the user are limited. The options include leaving the social network, surrendering their identity and features provided by the social network; accepting the disagreeable activities; or to petition and hope that the social network changes its behavior.

As the use of social networking expands to become the primary way in which users communicate and express their identity amongst their peers, the users become more dependent on the policies of social network infrastructure owners. Recent work [15] explores the coupling between social networks and P2P systems as a means to return ownership to the users, noting that a social network made up of social links is inherently a P2P system with the aside that they are currently developed on top of centralized systems. This chapter extends this idea with focus on the topic of topology; that is, how to organize social profiles that leverage the benefits offered by a structured P2P overlay abstraction.

Structured P2P overlays provide a scalable, resilient, autonomic platform for distributed applications. Structured overlays enable users to easily create their own decentralized systems for the purpose of data sharing, interactive activities, and other networking-enabled activities. This chapter is based upon my previous work [115, 117] discussed in chapters 3 and 4 to enable social network profile overlays. These works

address the challenges of bootstrapping secure, private overlays in environments constrained by network address translators (NATs) and firewalls through a public overlay used for discovery and as a relay or communication transport.

A typical social network consists of users and groups. Each user has a profile, a set of friends, and the ability to send and receive private messages; each group consists of one or more managers, users, and a messaging board. Profiles contain user's personal information, status updates, and public conversations, similar to a message board. Friends are individuals trusted sufficiently by a user to view the user's profile. Private messaging sends messages discretely between users without leaking the message to other members. Groups have similar features, though identity is shared by many users.

Using this social networking model, I have designed OverSoc. OverSoc uses a public overlay as a directory for finding and befriending peers or finding and accessing groups. Once group and profile access has been offered, the public overlay can be used to bootstrap connectivity to existing profile and group overlays. Security for a profile is provided by a public key infrastructure (PKI), where profile owners or group managers are the certificate authorities (CA) and all members have signed certificates. The overlay stores profile data or group information in its distributed data store, supporting decentralized access using scalable mechanisms regardless of the profile owner's online presence. In this chapter, I present the architecture of these overlays, as presented in Figure 7-1.

The rest of this chapter is organized as follows. Section 7.1 discusses related work. Section 7.2 describes OverSoc, explaining how to map social networks onto structured P2P overlays. Section 7.3 expresses expectations for user interaction in the system. In Section 7.4, I explore some of the remaining challenges introduced by this approach.

7.1 Related Works

In [16], Buchegger et al. describe how to use a DHT to store social networking profile. The DHT provides look-up services for storing meta-data pertaining to a peer's




Figure 7-1. An example OverSoc social overlay network. Alice has a friendship with Bob and Carol, hence both are members of her profile overlay. Bob has a friendship with Alice and Dave but not Carol; hence Alice and Dave are members of his profile overlay, while Carol is not. Each peer has many overlay memberships but a single root represented by dashed lines in various shades of gray. For clarity, overlay shortcut connections are not shown.

profile. Peers query the DHT for updated content from their friends by hashing their unique identifiers (e.g. friends' email addresses). The retrieved meta-data contains information for obtaining the profile data such as IP address and file version. Their work relies on a PKI system that provides identification, encryption, and access control. In contrast, OverSoc maps individual user profiles and groups to a private overlay secured by point-to-point encryption and authentication amongst all peers in the overlay. The private overlay provides a clean abstraction of access control, whereby once admitted to a private overlay, users can access a distributed data store which holds the contents of the owner's profile.

Shakimov et al. in [98] take a different approach by depending on virtual individual servers (VIS) hosted on a cloud infrastructure such as Amazon EC2. Friends contact

each other's VIS directly for updates. A DHT is used as a directory for groups and interest-based searches. Their approach assumes bidirectional end-to-end connectivity between each VIS, where a profile is only available during the up time of the VIS. Because of the demands on network connectivity and up time, the approach assumes a cloud-hosted VIS and has difficulty being used on user-owned resources. OverSoc allows peers to have asymmetric connectivity and does not require constant up time through the use of NAT traversal support and the ability to store the profile in the overlay's distributed data store.

The approach presented by Cutillo et al. in [26] relies on a central system to host identities and certificates that can then be used to query a DHT to discover an initial hop in a route to a specific peer through their circle of friends. The circle of friends consists of an unstructured overlay, where direct friends maintain direct connections with the peer, and outer circles consist of friends of friends and friends of friends of friends. The main goal of this work is to remove the private components of a profile from a central entity, whereas OverSoc makes a clean break from all centralization and enables scalability through distributed replication techniques.

Unlike the above approaches, the P2P social network presented by Abbas et al. in [1] uses an unstructured overlay without a DHT where peers connect directly to each other rather than through the overlay establishing unique identifiers to deal with dynamic IPs. Peers cache each other's data to improve availability, while helper nodes are used to assist with communication between peers behind NATs. The approach lacks security and access control considerations and lacks the guarantees and the simplicity of the abstraction offered by a structured overlay.

7.2 Social Overlays

In this section, I explain how OverSoc maps online social networking to virtual private overlays consisting of a public directory overlay with many private profile overlays. The directory overlay supports friend discovery and verification and stores




Figure 7-2. Alice requests and receives a friendship from Bob.

a lists of peers currently active in each profile overlay. Profile overlays support message boards, private messages, and media sharing.

7.2.1 Finding Friends

In a traditional social network, directories are used to search for users based upon public information, such as the user's full name, user ID, e-mail address, group affiliations, and friends. The resulting search returns zero or more matching directory entries. In OverSoc, directory entries are inserted into the DHT of a public overlay. Since the public information has many components, various subsets form DHT keys that all point to a common, complete listing of the matching public information. For example, a user can store a pointer at the DHT key *hash("alice")* or *hash("alicebob")*. The key here is that any subset of the user's public information in lower-case format can be hashed into a DHT index that would eventually direct the searching user to one or more users' public information. More explicit searches could sift through the results and present to the user only those peers matching all the search parameters. The amount of information shared publicly should be configurable by the user.

While looking for an individual, a peer may discover that many individuals have overlapping public information components, such as the user's name. Assuming all entries are legitimate, the overlay must have some method of supporting multiple, distinct values at the same key, requiring the application and user to parse the responses and determine the best match by reviewing the contents of each certificate. Alternatively, a technique like Sword [3], which supports attribute based searching, could be used to efficiently find peers in an overlay.

To address trust levels when searching for friends, a PGP certificate can be used to store user's public information and verify user's friends and groups. In OverSoc, the main portion of a PGP certificate contains information such as user name, full name, e-mail address, potentially other user-defined data, and signature packets from the user and those that trust the certificate including groups and individuals. These signature packets represent a list of verifiable friends and groups assisting to further uniquely identify a user. Each time a user befriends someone, they should exchange signature packets containing at a minimum the friend's PGP certificate ID, a signature expiration time, and a signature binding this information with the new friend's existing PGP certificate. This increases the trust level of individuals searching for others especially if they have common friendships or group membership. The use of a time stamp in the signature assists in deciding whether or not a friendship link is still active without accessing the profile overlay of either peers. Thus peers that maintain friendships need to periodically exchange signature packets.

7.2.2 Making Friends

In this example, Alice becomes friends with Bob, as illustrated in Figure 7-2. Once a user, *Alice*, has found a friend candidate, *Bob*, *Alice* can issue a friendship request and store it in the DHT using the hash of Bob's certificate as an index, this acts a public overlay mailbox. *Bob* can review the public information of *Alice* prior to making a decision. If *Bob* accepts the request, *Alice* and *Bob* exchange signature packets and




Figure 7-3. Alice, already a friend of Bob, connects to his social overlay.

are granted access to each other's profiles. Once profile access has been enabled, the *Alice* and *Bob* can learn more information, and if it turns out to be a mistake, either one of them can unilaterally end the relationship.

Alice's friendship request should contain a pointer to her certificate in the overlay, a time stamp, and Bob's certificate identifier. The friendship request is encrypted using Bob's public key and signed using Alice's private key for the purposes of anonymity and authenticity. When Bob receives the friendship request, he can verify that the request was made for Bob by Alice. Upon receiving the friendship request, he has three choices: a conditional accept, an unconditional accept, or a reject. During an unconditional accept, Bob signs Alice's PGP certificate and issues a request to befriend her. Alternatively, he could issue a request to befriend her and wait for her to sign his certificate and investigates her profile prior to signing hers.

Discovery of a user is not limited to the directory entries. Because users have a public overlay based mailbox, they are not required to discover each other only through the directory. Instead, they can use out of band discovery, using mechanisms like e-mail, chat, or personal websites to exchange certificates. Once a peer has received another peer's certificate, they can submit secure friendship requests using the public overlay. In

fact, this sort of system can leverage the trust established by an existing social network to sign and exchange OverSoc's certificates.

7.2.3 The Profile Overlay

In a traditional social network, the profile or user-centric portion consists of private messaging, data sharing, friendship maintenance, and a public message board for status updates or public messages. In this section, we explain how these components can be applied to a structured overlay dedicated to an individual profile.

Using the techniques such as those described earlier, it is feasible to efficiently multiplex a P2P system across multiple, virtual private overlays enabling each profile owner to have a profile overlay consisting of their online friends. For access control, OverSoc employs point-to-point encryption and authentication, peers bootstrap private connections by exchanging the base of the PGP certificate and the profile overlays signature packet obtained in the "making friends" stage. Because the profile owner also is the CA, control of which could be distributed across the users resources, for all members of the overlay, they can easily revoke users from access to the profile overlay. Chapter 4 describes efficient mechanisms for overlay revocation through the use of broadcasting for immediate revocation and the use of DHT for indirect and permanent revocation.

The message board of a profile can be stored in two ways: distributed within the profile overlay via a data store or stored on the profile owner's personal computing devices. The distributed data store provides the profile when the owner is offline and also distributes the load for popular profiles. For higher availability, each peer always stores and provides all data in their profile when they are online. To ensure authenticity and integrity, peers sign their messages and each peer's certificate is available in the overlay as well as stored by mutual friends for verification. Messages that are unsigned are ignored by all members of the overlay. An ideal overlay for this purpose should

support complex queries [50] allowing easy access to data stored chronologically, by content, by type, i.e., media, status updates, or message board discussions.

Private messaging in the profile overlay is unidirectional; only the profile owner can receive private messages using their overlay. To enforce this, a private message should be prepended with a symmetric key encrypted by the profile owners public key, the message should be appended by a signature of the message using the private key of the message sender, and the entire message encrypted by the symmetric key. This approach ensures that only the sender and the profile owner can decrypt the private message and verify the senders identity. The contents of the private message include the sender, time sent, and the subject. Messages are be stored in well known locations in the DHT, like “private messages for me”, so that the profile owner can either poll the location.

7.2.4 Event Based Message Notification

Both the directory and profile overlays have methods by which peers can receive messages. In the directory overlay, these take form by means of friendship requests and friendship accepts, certificate signature packets. The profile overlay supports private messages. While polling the location in the DHT occasionally will allow peers to receive the messages, polling has inherent delays and network costs. Alternatively, event enable peers to receive sent messages very quickly after they have been sent with minimal impact on network throughput.

A simple method for implementing an event notification system involves using the DHT. Each event would have an identification that would map to a list of peers wanting to know when an event occurred and the data associated with it. Thus mapping the (*eventid*, *listener*) to the DHT could be done by hashing a string such as “private messages for me” or taking a hash of the user’s certificate hash for public overlay messages and storing the profile owners active nodes into the list of listeners . When a message is inserted into the user’s mailbox, the sender could query this list and send to

each listener a notification of the new private message. Alternatively, if a higher degree of anonymity is required, the DHT server could be modified to forward the response to the listeners directly rather than returning a list of listeners. Of course, this does not prevent potential race conditions occurring, such as a situation where a peer recently joined their profile overlay, had already queried their mailbox and found it empty, while simultaneously a private message was sent to them yet they were not in the listeners list. Thus occasional polling is required, though can be minimized, the longer a node has been online.

7.2.5 Active Peers

The directory overlay should be used to assist in finding currently active peers in the profile overlays. By placing their node IDs at a well-known, unique per-profile overlay keys in the DHT, active peers can bootstrap incoming peers into the profile overlay. I implemented and evaluated this concept in Chapter 4. Because the profile overlay members all use PKI to ensure membership, even if malicious peers insert their ID into the active list, it would be useless as the peer would only form connections with peers who also have a signed certificate. Extending from the earlier example, where Alice became Bob's friends, Figure 7-3 presents in detail how she would join his private overlay.

7.2.6 Groups

Groups can be considered extensions of profile overlays. The fundamental difference between a group and a profile is that a group lacks private messaging and has shared ownership. So just as a peer can find a profile in the directory by hashing the name of the user and other identifiable information, so can the user find the group. Like the certificate of the user, the members of a group sign the group's certificate to represent their membership to that group. In OverSoc, users request membership to the group like they do friendship requests, in response a group manager can sign

their certificate allowing that member access to the group. Finally, the group can be bootstrapped in the same way as the profile overlay through the directory overlay.

The unique challenge presented by groups is the sharing of the CA task. A decentralized solution would be for all members of the group to be listed in the groups DHT and when a peer becomes a manager, they obtain a new signature packet that contains a user-defined component stating that they are managers. If an administrator loses their position, then all members who had their certificate signed by that administrator would need to obtain a new certificate. To avoid member churn, the owner could provide signature packets for all group members. Thus the managers just allow temporary access until the owner comes online and provides more permanent access.

7.3 User Interaction

OverSoc consists of many components that are transparent to the user, the user experience should appear to the user no differently than an existing online social network. The OverSoc could be a downloadable application or a browser based Flash or Silverlight application. If the user, Bob, had already created an account, Bob would be presented with an interface showing their friends profiles. Based upon Bob's configuration, the social application could retrieve profile updates as he navigates to individual profiles or as soon as the application joins an individual profile overlay, reactive versus proactive profile querying.

If this was Bob's first time starting OverSoc, he would be presented with screens asking for his privacy preferences, such as whether or not he wants his information in the directory overlay, if he felt comfortable enough with the idea of people knowing he was a member of the social network and who his friends are. Then OverSoc would ask for personal information to populate his profile and to generate his directory information. At which point, the OverSoc would join the overlay and create Bob's private overlay.

Bob could then start searching for friends, make friend requests, and respond to friend requests.

Recently, Bob had been thinking about his high school days and was curious if Alice was also a member of OverSoc, though Bob did not have Alice's e-mail address, just her first and last name. Bob enters Alice's name into the OverSoc search box and is presented by a list of Alice's. As Bob reviews each of the entries, he recognizes an Alice that is friend's with some of the same people Bob was in high school. Bob selects to become her friend. At which point, the OverSoc transparently inserts a friendship request to Alice and signs Alice's certificate so Alice can view Bob's profile. Of course that is because Bob has chosen to allow user-initiated friend requests access to his profile. Alice receives Bob's request, peruses his profile and feels fine becoming friends with Bob, which initiates a transparent process of signing Bob's certificate and placing the result in the public overlay. There is one problem though, when Bob receives Alice's signature and views her profile, he realizes that this is some other Alice. He quickly chooses to defriend her. This causes Bob's OverSoc instance to broadcast a revocation for Alice's signature and to store the revocation in the DHT. Alice, who was viewing Bob's profile, is notified of this sudden loss of trust and while she is able to view the contents of Bob's profile, which she has already accessed and obtained, she can no longer receive updates as members of Bob's overlay prevent her from accessing it.

In another instance, Bob bumped into Carol, who e-mailed Bob a copy of her certificate. Bob points OverSoc to the certificate, and OverSoc verifies that he wants to become friends with the identity associated with the certificate. When he accepts, OverSoc immediately submits a request to become Carol's friend. Carol receives notification and accepts Bob's friendship request. At this point, both Bob and Carol have transparently exchanged signed certificates and have mutual access to each other profiles. As Bob reads Carol's latest news, he remembers a funny personal story and that he would like to share with Carol. So he sends Carol a private message. Carol

is offline though. The next time Carol goes online, her social application discovers the message and presents it to her. In this scenario, OverSoc has taken the private message, secured it with her public key and a symmetric key and signed it with his private key. After which, it inserts the message into the DHT and sends a notice to the event notification system, which detects that there were no listeners. When Carol's application comes online, it queries the DHT receiving the message. Prior to presenting Carol the message, the OverSoc decrypts and verifies the message.

The OverSoc architecture can leverage existing social networks to bootstrap trust. For example, consider Bob and David are two friends on Facebook. Bob joins a Facebook application called "OverSoc/Facebook Bridge", which stores a copy of his OverSoc certificate in his personal profile. Bob has been bragging to David about OverSoc and mentions to him how easy it is to migrate from Facebook to OverSoc using this application. So David joins OverSoc as well as the application. When David accesses the application, it pastes his certificate to his profile, notifies him that he has a friend already using it, Bob, and that he can immediately sign Bob's certificate, and leaves a request for Bob to sign his certificate. Additionally, when David logs into OverSoc, he can leave a friend request there as well, so that the next time Bob accesses Facebook or OverSoc, he will receive David's request and can sign David's certificate. At which point, both will have access to each others OverSoc profile overlays.

7.4 Challenges

While structured P2P overlays have been well-studied in a variety of applications, their use in social profile overlays raises new interesting questions, including:

1) Handling small overlay networks - P2P overlay research typically focuses on networks larger than the typical user's friend count (Facebook's average is 130¹). Because social profile overlays are comparatively smaller, this can impact the reliability

¹ <http://www.facebook.com/press/info.php?statistics>

of the overlay and availability of profile data. A user can host their own profile; however when the user is disconnected it is important that their profile remains available even under churn. It is thus important to characterize churn in this application to understand how to best approach this problem. An optional of per-user deployment of a virtual individual server (VIS) and the use of replication schemes aware of a user's resources provide possible directions to address this issue.

2) Overlay support for low throughput, unconnected devices - devices such as smart phones cannot constantly be actively connected to the overlay and the connection time necessary to retrieve something like a phone number may be too much to make this approach useful. Similar to the previous challenge, this approach could benefit from using a VIS enabling users access to their social overlays by proxy without establishing a direct connection to the overlay network.

3) Reliability of the directory and profile overlay - Overlays are susceptible to attacks that can nullify their usefulness. While the profile overlay does have point-to-point security, in the public, directory overlay, the lack of any form centralization makes policing the system a complicated procedure. While the approach of appending friends list can assist users in making decisions on identity, it does not protect against denial of service attacks. For example, users could attempt create many similar identities in an attempt to overwhelm a user in their attempt to find a specific peer. Previous work has proposed methods to ensure the usability of overlays even while under attack. For the social overlay to be successful, one must identify which methods should be used. A possible approach is to replicate public information within a user's profile overlay thus providing an alternative directory overlay for querying prior to using the public directory overlay.

4) Social profile data storage - In previous works, DHTs have been used as the building blocks to form more complex distributed data stores as presented in Past [95] and Kosha [17]. Application of data stores will be heavily dependent on the churn rate

associated with the overlay. If the system lacks any reasonably stable membership, large data files may be corrupted while smaller data sets are completely lost. Ideally, the usage model would be similar to those of Skype and Twitter, which have active processes for the duration of the computers usage. In an environment like this, data storage would be limited only by the available bandwidth of the participants.

CHAPTER 8

CONCLUSIONS

This work brings significant advances to the usability of VPNs through understanding important practical applications and verification in both simulation and real deployments. The architecture explored herein provides a general framework for creating VPNs that has contributed to various end-point and overlay configurations useful for both large and small scale deployments for group or personal use.

In order to support a completely ad-hoc, decentralized VPN, users begin by connecting to a public overlay, such as XMPP or Kademlia, in order to discover other users. After exchanging their information via these mediums, peers can establish direct communication links with each other and with other peers already in the VPN. Peers can exchange or obtain trusted identities using established peer or group relationships in existing social networks.

Because most peers are behind NATs and firewalls, this dissertation covers methods, which allow peers to use existing overlays to bootstrap through NATs and firewalls. This can mean using a third-party system until a peer on a public IP address comes online, or more likely, using a service to obtain a public IP and port mapping for the peer's private address. When peers cannot directly establish direct links, the overlay still provides the ability to route messages between peers. Routing messages across the overlay can incur significant overhead and is not optimized for any specific purpose. To remedy this, I have established a mechanism to create two-hop links between peers emphasized on the latency between the peers.

This work describes two novel mechanisms for handling address assignments inside a VPN: using a DHT with atomic capabilities as well as independent networks with explicit links based upon social connections. The DHT approach allows for highly scalable systems in comparison to other approaches that require state to be manually spread across the system or through the use of broadcast mechanisms. By making the

network addresses dependent on social links and independent of the actual overlay, peers need not worry about address collisions. Furthermore, this work explores means to transparently migrate resources using DHT style addressing even when those resources are connected via the VPN router.

Existing approaches to VPN placement use either interface or router models. Interface models can easily be constructed to be transparent, whereas existing router models have no such features. Through the use of network protocols supported in network stacks found in common operating systems, mechanisms for support transparent configuration of both interface and routing models have been detailed. Where routers are desirable due to performance, and interface is attractive for security purposes, a hybrid model provides a middle ground that combines these two aspects to support high-performance, though secure virtual networking.

Because attempting to verify the system in every possible environment after making adding features or fixing bugs requires a significant time investment, I have employed a built-in self-simulating environment into Brunet. The application has reduced the time necessary to develop, evaluate, and debug new contributions as well as reproduce bugs and protect from having them reoccur. In the context of this dissertation, it has been used to verify and motivate the necessity for on demand as opposed to passive connections, as well as the relay, bootstrapping, and security work.

This work has been the corner stone in a real system used to provide ad-hoc grid computing named the Grid Appliance [114], which has been realized in a voluntary computing grid for computer architecture research called Archer [37]. Archer currently spans six universities with over 600 resources. Over hundreds of users have connected seamlessly to these resources from many locations. A PlanetLab back end distributed across over 600 resources provides near constant overlay uptime for Archer and external users. External users include classes and groups at other universities. Most recently, a grid at La Jolla Institute for Allergy and Immunology went live with minimal

communication with our group. Researchers at the Clemson University and Purdue have opted for this approach over centralized VPNs as the basis of their future distributed compute clusters and have actively tested networks of over 1000 nodes.

The majority of this dissertation focused on services to enable user-friendly VPNs. During the design and evaluation, it became apparent that there still exist significant deficiencies in the design described herein. The following are important research topics that are left as future research topics. During the evaluation of packet drop rates across the Internet and in particularly PlanetLab, it is apparent that recursive packet routing will not scale well especially when dealing with non-negligible traffic. Another aspect of limited scalability is presented when using a single UDP socket multiplexed across potentially many VPN connections. Each UDP socket has a small buffer associated with it. When that buffer is exhausted sends either become blocking or are thrown away, depending on usage. In terms of VPNs, there still exists a wide gap between IPv6 and IPv4. My work could be used to assist in deploying IPv6 to IPv4 tunnels, which unlike existing approaches, have natural fail over support and efficient paths between source and destination. Finally, the last major draw back to the Grid Appliance is the necessity for a centralized scheduler / management component. Ideally, this could be handled in decentralized means with trust, a user rank system for priority, ability to handle simultaneous scheduling of tasks, and limiting the reliance of a node being online to receive results for tasks.

APPENDIX: STRUCTURED OVERLAY BROADCAST




Figure A-1. Tree-based overlay broadcast

Broadcast revocation can be used to address the deficiencies of DHT revocation. As a topic of previous research works [32, 111], structured overlays can be used without additional state to perform efficient broadcasts from any point in the overlay to the entire overlay. In these papers, analysis and simulations have shown that the approach can be completed in a network size of n in $O(\log^2 n)$ time with n messages. The overlay broadcast algorithm used in this paper provides a complete overlay broadcast in $O(\log^2 n)$ time with n messages. When applied to Brunet, as illustrated in Figure A-1, it utilizes the organization of a structured system with a circular address space that requires peers be connected to those whose node addresses are the closest to their own, features typical of one-dimensional structured overlays including Chord [104], Pastry [94], and Symphony. Using such an organization, it is possible to do perform a broadcast with no additional state. To perform a broadcast, each node performs the following recursive algorithm:

BROADCAST(start, end, message):

RECEIVE(message)

```

for  $i$  in length(connections) do
    n_start  $\leftarrow$  ADDRESS(connections[ $i$ ])
    if n_start  $\notin$  [start, end) then
        continue
    end if
    n_end  $\leftarrow$  ADDRESS(connections[ $i + 1$ ])
    if n_end  $\notin$  [start, end) then
        n_end  $\leftarrow$  end
    end if
    msg  $\leftarrow$  (BROADCAST, n_start, n_end, message)
    SEND(connections[ $i$ ], msg)
end for

```

with “connections” as a circular list of connections in non-decreasing order from the perspective of the node performing the current recursive, broadcast step.

In this algorithm, the broadcast initiator uses its own address as the start and end, thus the broadcast will span the entire overlay after completing recursive calls at each connected node. A recursive end, “n_end”, must be inside the region between “start” and “end”, thus if the connection following the current sending connection, “connections[$i + 1$]”, is not in that region, it will only broadcast up to “end” and not the address specified by that connection. To summarize, the overlay is recursively partitioned amongst the nodes at each hop in the broadcast. By doing so, all nodes receive the broadcast without receiving duplicate broadcast messages.

REFERENCES

- [1] S. M. A. Abbas, J. A. Pouwelse, D. H. J. Epema, and H. J. Sips. A gossip-based distributed social networking system. In *Enabling Technologies, IEEE International Workshops on*, 2009.
- [2] H. Abbes, C. Cérin, and M. Jemni. Bonjourgrid: Orchestration of multi-instances of grid middlewares on institutional desktop grids. In *International Parallel and Distributed Processing Symposium (IPDPS)*, 2009.
- [3] J. Albrecht, D. Oppenheimer, A. Vahdat, and D. A. Patterson. Design and implementation trade-offs for wide-area resource discovery. In *ACM Trans. Internet Technol.*, 2008.
- [4] S. Alexander and R. Droms. *RFC 2132 DHCP Options and BOOTP Vendor Extensions*, March 1997.
- [5] Amazon.com, Inc. Amazon elastic compute cloud. <http://aws.amazon.com/ec2>, 2009.
- [6] D. P. Anderson. Boinc: A system for public-resource computing and storage. In *the International Workshop on Grid Computing*, 2004.
- [7] N. Andrade, L. Costa, G. Germoglio, and W. Cirne. Peer-to-peer grid computing with the ourgrid community. In *Brazilian Symposium on Computer Networks*, May 2005.
- [8] N. Andrade, L. Costa, G. Germoglio, and W. Cirne. Peer-to-peer grid computing with the ourgrid community. In *Brazilian Symposium on Computer Networks (SBRC) - 4th Special Tools Session*, 2005.
- [9] P. Andreetto, S. Andreozzi, G. Avellino, S. Beco, A. Cavallini, M. Cecchi, V. Ciaschini, A. Dorise, F. Giacomini, A. Gianelle, U. Grandinetti, A. Guarise, A. Krop, R. Lops, A. Maraschini, V. Martelli, M. Marzolla, M. Mezzadri, E. Molinari, S. Monforte, F. Pacini, M. Pappalardo, A. Parrini, G. Patania, L. Petronzio, R. Piro, M. Porciani, F. Prelz, D. Rebatto, E. Ronchieri, M. Sgaravatto, V. Venturi, and L. Zangrando. The glite workload management system. *Journal of Physics: Conference Series*, 119(6):062007, 2008.
- [10] Azurues. Message stream encryption. http://www.azureuswiki.com/index.php/Message_Stream_Encryption, December 2007.
- [11] A. Biggadike, D. Ferullo, G. Wilson, and A. Perrig. NATBLASTER: Establishing TCP connections between hosts behind NATs. In *ACM SIGCOMM Asia Workshop*, April 2005.
- [12] F. Bondoux. Campagnol : distributed vpn over udp/dtls. <http://campagnol.sourceforge.net>, 2010.

- [13] P. O. Boykin, J. S. A. Bridgewater, J. S. Kong, K. M. Lozev, B. A. Rezaei, and V. P. Roychowdhury. A symphony conducted by brunet. <http://arxiv.org/abs/0709.4048>, 2007.
- [14] D. Bryan, B. Lowekamp, and C. Jennings. Sosimple: A serverless, standards-based, p2p sip communication system. In *Advanced Architectures and Algorithms for Internet Delivery and Applications*, June 2005.
- [15] S. Buchegger and A. Datta. A case for P2P infrastructure for social networks - opportunities & challenges. In *WONS '09: The Sixth International Conference on Wireless On-demand Network Systems and Services*, 2009.
- [16] S. Buchegger, D. Schiöberg, L. H. Vu, and A. Datta. Peerson: P2p social networking: early experiences and insights. In *Workshop on Social Network Systems*, 2009.
- [17] A. R. Butt, T. A. Johnson, Y. Zheng, and Y. C. Hu. Kosha: A peer-to-peer enhancement for the network file system. In *IEEE/ACM Supercomputing*, 2004.
- [18] S. Carl-Mitchell and J. S. Quarterman. RFC 1027 - using arp to implement transparent subnet gateways, October 1987.
- [19] M. Castro, M. Costa, and A. Rowstron. Debunking some myths about structured and unstructured overlays. In *Symposium on Networked Systems Design & Implementation*, 2005.
- [20] M. Castro, P. Druschel, A. Ganesh, A. Rowstron, and D. S. Wallach. Security for structured peer-to-peer overlay networks. In *Symposium on Operating Systems Design and Implementation (OSDI)*, December 2002.
- [21] M. Castro, P. Druschel, A.-M. Kermarrec, and A. Rowstron. One ring to rule them all: Service discover and binding in structured peer-to-peer overlay networks. In *Symposium on Operating Systems Principles (SOSP) European Workshop*, Sept. 2002.
- [22] F. Chang, J. Dean, S. Ghemawat, W. C. Hsieh, D. A. Wallach, M. Burrows, T. Chandra, A. Fikes, and R. E. Gruber. Bigtable: a distributed storage system for structured data. In *Symposium on Operating Systems Design and Implementation (OSDI)*, 2006.
- [23] B. Chun, D. Culler, T. Roscoe, A. Bavier, L. Peterson, M. Wawrzoniak, and M. Bowman. Planetlab: an overlay testbed for broad-coverage services. *SIG-COMM Comput. Commun. Rev.*, 2003.
- [24] M. Conrad and H.-J. Hof. A generic, self-organizing, and distributed bootstrap service for peer-to-peer networks. In *International Workshop on Self-Organizing Systems (IWSOS)*, 2007.

- [25] C. Cramer, K. Kutzner, and T. Fuhrmann. Bootstrapping locality-aware p2p networks. In *International Conference on Networks (ICON)*, 2004.
- [26] L. A. Cutillo, R. Molva, and T. Strufe. Privacy preserving social networking through decentralization. In *Wireless On-Demand Network Systems and Services (WONS)*, 2009.
- [27] H. Damfpling. Gnutella web caching system. <http://www.gnucleus.com/gwebcache/specs.html>, 2003.
- [28] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Voshall, and W. Vogels. Dynamo: amazon's highly available key-value store. In *Symposium on Operating Systems Principles (SOSP)*, New York, NY, USA, 2007. ACM.
- [29] L. Deri and R. Andrews. N2N: A layer two peer-to-peer vpn. In *International conference on Autonomous Infrastructure, Management and Security*, 2008.
- [30] J. R. Douceur. The sybil attack. In *International Workshop on Peer-to-Peer Systems*, pages 251–260. Springer-Verlag, 2002.
- [31] R. Droms. *RFC 2131 Dynamic Host Configuration Protocol*, March 1997.
- [32] S. El-Ansary, L. Alima, P. Brand, and S. Haridi. Efficient broadcast in structured p2p networks. In *International Workshop on Peer-to-Peer Systems (IPTPS)*, 2003.
- [33] D. H. J. Epema, M. Livny, R. van Dantzig, X. Evers, and J. Pruyne. A worldwide flock of condors: Load sharing among workstation clusters. *Future Generation Computer Systems*, 12(1):53 – 65, 1996.
- [34] E. Exa. CloudVPN. <http://e-x-a.org/?view=cloudvpn>, September 2009.
- [35] D. Fabrice. Skype uncovered. http://www.ossir.org/windows/supports/2005/2005-11-07/EADS-CCR_Fabrice_Skype.pdf, November 2005.
- [36] Facebook. Facebook. <http://www.facebook.com>, January 2010.
- [37] R. J. Figueiredo, P. O. Boykin, J. A. B. Fortes, T. Li, J. Peir, D. Wolinsky, L. K. John, D. R. Kaeli, D. J. Lilja, S. A. McKee, G. Memik, A. Roy, and G. S. Tyson. Archer: A community distributed computing infrastructure for computer architecture research and education. In *CollaborateCom*, November 2008.
- [38] R. J. Figueiredo, P. O. Boykin, P. S. Juste, and D. Wolinsky. Integrating overlay and social networks for seamless p2p networking. In *Workshop on Enabling Technologies: Infrastructure for Collaborative Enterprises*, 2008.
- [39] R. J. Figueiredo, P. A. Dinda, and J. A. B. Fortes. A case for grid computing on virtual machines. In *International Conference on Distributed Computing Systems*. IEEE Computer Society, 2003.

- [40] I. Foster. Globus toolkit version 4: Software for service-oriented systems. *Journal of Computer Science and Technology*, 21:513–520, 2006. 10.1007/s11390-006-0513-y.
- [41] A. Ganguly, A. Agrawal, O. P. Boykin, and R. Figueiredo. IP over P2P: Enabling self-configuring virtual IP networks for grid computing. In *International Parallel and Distributed Processing Symposium*, 2006.
- [42] A. Ganguly, A. Agrawal, P. O. Boykin, and R. Figueiredo. Wow: Self-organizing wide area overlay networks of virtual workstations. In *IEEE High Performance Distributed Computing (HPDC)*, June 2006.
- [43] A. Ganguly, P. O. Boykin, D. Wolinsky, and R. J. Figueiredo. Improving peer connectivity in wide-area overlays of virtual workstations. *Cluster Computing Journal*, 7 2009.
- [44] A. Ganguly, D. Wolinsky, P. Boykin, and R. Figueiredo. Decentralized dynamic host configuration in wide-area overlays of virtual workstations. In *International Parallel and Distributed Processing Symposium*, March 2007.
- [45] C. GauthierDickey and C. Grothoff. Bootstrapping of peer-to-peer networks. In *International Symposium on Applications and the Internet*, 2008.
- [46] W. Ginolas. P2PVPN. <http://p2pvpn.org>, 2009.
- [47] B. Gleeson, A. Lin, J. Heinanen, T. Finland, G. Armitage, and A. Malis. RFC 2764 a framework for IP based virtual private networks, February 2000.
- [48] S. Guha, N. Daswani, and R. Jain. An experimental study of the skype peer-to-peer voip system. In *International Workshop on Peer-to-Peer Systems*, 2006.
- [49] K. P. Gummadi, S. Saroiu, and S. D. Gribble. King: Estimating latency between arbitrary internet end hosts. In *SIGCOMM Internet Measurement Workshop*, 2002.
- [50] M. Harren, J. M. Hellerstein, R. Huebsch, B. T. Loo, S. Shenker, and I. Stoica. Complex queries in dht-based peer-to-peer networks. In *International Workshop on Peer-to-Peer Systems*, 2002.
- [51] A. Harutyunyan, P. Buncic, T. Freeman, and K. Keahey. Dynamic virtual AliEn grid sites on nimbus with CernVM. *Journal of Physics: Conference Series*, 2010.
- [52] B. Hubert. tc - linux advanced routing & traffic control. <http://lartc.org/>, June 2009.
- [53] X. Jiang and D. Xu. Violin: Virtual internetworking on overlay. In *International Symposium on Parallel and Distributed Processing and Applications*, pages 937–946, 2003.

- [54] R. Jones. Netperf: A network performance monitoring tool. <http://www.netperf.org>, 2009.
- [55] D. Joseph, J. Kannan, A. Kubota, K. Lakshminarayanan, I. Stoica, and K. Wehrle. Ocala: an architecture for supporting legacy applications over overlays. In *Symposium on Networked Systems Design & Implementation*, pages 20–20, 2006.
- [56] P. S. Juste, D. Wolinsky, P. Oscar Boykin, M. J. Covington, and R. J. Figueiredo. SocialVPN: Enabling wide-area collaboration with integrated social and overlay networks. *Computer Networking*, 54:1926–1938, August 2010.
- [57] M. Kallahalla, M. Uysal, R. Swaminathan, D. E. Lowell, M. Wray, T. Christian, N. Edwards, C. I. Dalton, and F. Gittler. SoftUDC: A software-based data center for utility computing. *Computer*, 37(11):38–46, 2004.
- [58] K. Keahey, K. Doering, and I. Foster. From sandbox to playground: Dynamic virtual environments in the grid. In *International Workshop in Grid Computing*, November 2004.
- [59] K. Keahey and T. Freeman. Contextualization: Providing one-click virtual clusters. In *eScience*, 2008.
- [60] K. Keahey and T. Freeman. Science clouds: Early experiences in cloud computing for scientific applications. In *Cloud Computing and Its Applications*, 2008.
- [61] C. C. Keir, C. Clark, K. Fraser, S. H, J. G. Hansen, E. Jul, C. Limpach, I. Pratt, and A. Warfield. Live migration of virtual machines. In *Symposium on Networked Systems Design and Implementation*, pages 273–286, 2005.
- [62] M. Knoll, A. Wacker, G. Schiele, and T. Weis. Bootstrapping in peer-to-peer systems. In *International Conference on Parallel and Distributed Systems (IPDPS)*, 2008.
- [63] M. Krasnyansky. Universal tun/tap device driver. <http://vtun.sourceforge.net/tun>, 2005.
- [64] A. Lakshman. Cassandra a structured storage system on a P2P network. http://www.facebook.com/note.php?note_id=24413138919, August 2008.
- [65] M. Livny, J. Basney, R. Raman, and T. Tannenbaum. Mechanisms for high throughput computing. *SPEEDUP Journal*, 11(1), June 1997.
- [66] G. LLC. Gbridge. <http://www.gbridge.com>, September 2009.
- [67] LogMeIn. Hamachi. <https://secure.logmein.com/products/hamachi2/>, 2009.
- [68] LogMeIn, Inc. LogMeIn hamachi2 security, 2009.

- [69] S. Ludwig, J. Beda, P. Saint-Andre, R. McQueen, S. Egan, and J. Hildebran. XEP-0166: Jingle, December 2009.
- [70] H. Madhyastha, T. Isdal, M. Piatek, C. Dixon, T. Anderson, A. Krishnamurthy, and A. Venkataramani. iplane: an information plane for distributed services. In *USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2006.
- [71] G. S. Manku, M. Bawa, and P. Raghavan. Symphony: distributed hashing in a small world. In *USITS*, 2003.
- [72] P. Maymounkov and D. Mazières. Kademlia: A peer-to-peer information system based on the XOR metric. In *International Workshop on Peer-to-Peer Systems*, 2002.
- [73] A. Mislove, A. Post, A. Haeberlen, and P. Druschel. Experiences in building and operating epost, a reliable peer-to-peer application. In *Symposium on Operating Systems Principles (SOSP)/EuroSys European Conference on Computer Systems*, 2006.
- [74] MySpace, Inc. Myspace. <http://www.myspace.com>, January 2010.
- [75] M. Nelson, B.-H. Lim, and G. Hutchins. Fast transparent migration for virtual machines. In *USENIX Annual Technical Conference*, pages 25–25, Berkeley, CA, USA, 2005.
- [76] D. Nurmi, R. Wolski, C. Grzegorzcyk, G. Obertelli, S. Soman, L. Youseff, and D. Zagorodnov. The eucalyptus open-source cloud-computing system. In *IEEE/ACM International Symposium on Cluster Computing and the Grid (CCGrid)*, 2009.
- [77] J. C. Perez. Facebook’s beacon more intrusive than previously thought. http://www.pcworld.com/article/140182/facebook_beacon_more_intrusive_than_previously_thought.html, 2007.
- [78] S. Perreault and J. Rosenberg. TCP candidates with interactive connectivity establishment (ICE). <http://tools.ietf.org/html/draft-ietf-mmusic-ice-tcp-08>, October 2009.
- [79] K. Petric. Wippien. <http://wippien.com/>, August 2009.
- [80] J. Postel. RFC 0925 - multi-lan address resolution, 1984.
- [81] Qumranet. Kernel-based virtual machine for linux. <http://kvm.qumranet.com/kvmwiki>, March 2007.
- [82] S. Ratnasamy, P. Francis, S. Shenker, and M. Handley. A scalable content-addressable network. In *ACM SIGCOMM*, 2001.

- [83] E. Rescorla and N. Modadugu. RFC 4347 datagram transport layer security, April 2006.
- [84] C. Resources. Torque resource manager. <http://www.clusterresources.com/pages/products/torque-resource-manager.php>, March 2007.
- [85] A. Rezmerita, T. Morlier, V. Neri, and F. Cappello. Private virtual cluster: Infrastructure and protocol for instant grids. In *Euro-Par*, November 2006.
- [86] S. Rhea, B. Godfrey, B. Karp, J. Kubiawicz, S. Ratnasamy, S. Shenker, I. Stoica, and H. Yu. Opendht: a public dht service and its uses. In *Conference on Applications, technologies, architectures, and protocols for computer communications*, pages 73–84, New York, NY, USA, 2005. ACM.
- [87] M. Ripeanu. Peer-to-peer architecture case study: Gnutella network. <http://www.cs.uchicago.edu/%7Ematei/PAPERS/gnutella-rc.pdf>, 2001.
- [88] E. Rosen and Y. Rekhter. *RFC 2547 BGP/MPLS VPNs*, March 1999.
- [89] J. Rosenberg. Interactive connectivity establishment (ICE): A protocol for network address translator (NAT) traversal for offer/answer protocols. <http://tools.ietf.org/html/draft-ietf-mmusic-ice-19>, October 2008.
- [90] J. Rosenberg, R. Mahy, and P. Matthews. "traversal using relays around nat (turn)". <http://tools.ietf.org/html/draft-ietf-behave-turn-16>, 2009.
- [91] J. Rosenberg, R. Mahy, P. Matthews, and D. Wing. RFC 3489 session traversal utilities for nat (STUN), October 2008.
- [92] J. Rosenberg, R. Mahy, P. Matthews, and D. Wing. RFC 3489 session traversal utilities for nat (STUN), October 2008.
- [93] J. Rosenberg, J. Weinberger, C. Huitema, and R. Mahy. RFC 3489 stun - simple traversal of user datagram protocol (udp) through network address translators (nats), 2003.
- [94] A. Rowstron and P. Druschel. Pastry: Scalable, decentralized object location and routing for large-scale peer-to-peer systems. In *IFIP/ACM International Conference on Distributed Systems Platforms (Middleware)*, November 2001.
- [95] A. Rowstron and P. Druschel. Storage management and caching in PAST, a large-scale, persistent peer-to-peer storage utility. In *Symposium on Operating Systems Principles (SOSP)*, 2001.
- [96] P. Saint-Andre. RFC 3920 extensible messaging and presence protocol (XMPP): Core, October 2004.
- [97] S. Santhanam, P. Elango, A. A. Dusseau, and M. Livny. Deploying virtual machines as sandboxes for the grid. In *WORLDS*, 2005.

- [98] A. Shakimov, H. Lim, L. P. Cox, and R. Caceres. Vis-à-vis:online social networking via virtual individual servers, May 2008.
- [99] Skype Limited. Skype. <http://www.skype.com>.
- [100] G. Sliepen. tinc. <http://www.tinc-vpn.org/>, September 2009.
- [101] P. Srisuresh, B. Ford, and D. Kegel. *RFC 5128 State of Peer-to-Peer (P2P) Communication across Network Address Translators (NATs)*, March 2008.
- [102] Standard Performance Evaluation Corporation. Specjbb2005. <http://www.spec.org/jbb2005/>, 2005.
- [103] I. Stoica, D. Adkins, S. Zhuang, S. Shenker, and S. Surana. Internet indirection infrastructure. *IEEE/ACM Transactions on Networking*, 2004.
- [104] I. Stoica, R. Morris, D. Liben-Nowell, D. R. Karger, M. F. Kaashoek, F. Dabek, and H. Balakrishnan. Chord: a scalable peer-to-peer lookup protocol for internet applications. *IEEE/ACM Transactions on Networking*, 11(1), 2003.
- [105] Sun. gridengine. <http://gridengine.sunsource.net/>, March 2007.
- [106] A. I. Sundararaj and P. A. Dinda. Towards virtual networks for virtual machine grid computing. In *Conference on Virtual Machine Research And Technology Symposium*, pages 14–14, 2004.
- [107] W. Townsley, A. Valencia, A. Rubens, G. Pall, G. Zorn, and B. Palter. *RFC 2661 Layer Two Tunneling Protocol*, August 1999.
- [108] M. Tsugawa and J. Fortes. A virtual network (vine) architecture for grid computing. *International Parallel and Distributed Processing Symposium*, 2006.
- [109] M. Tsugawa and J. Fortes. Characterizing user-level network virtualization: Performance, overheads and limits. In *IEEE International Conference on eScience*, pages 206–213, Dec. 2008.
- [110] UPnP Forum. UPnP device architecture 1.1. <http://www.upnp.org/specs/arch/UPnP-arch-DeviceArchitecture-v1.1.pdf>, October 2008.
- [111] V. Vishnevsky, A. Safonov, M. Yakimov, E. Shim, and A. D. Gelman. Scalable blind search and broadcasting over distributed hash tables. *Computer Communications*, 31(2), 2008.
- [112] VMware, Inc. Timekeeping in vmware virtual machines. http://www.vmware.com/pdf/vmware_timekeeping.pdf, 2008.
- [113] D. I. Wolinsky, A. Agrawal, P. O. Boykin, J. Davis, A. Ganguly, V. Paramygin, P. Sheng, and R. J. Figueiredo. On the design of virtual machine sandboxes for distributed computing in wide area overlays of virtual workstations. In *International Workshop on Virtualization Technologies in Distributed Computing*, 2006.

- [114] D. I. Wolinsky and R. Figueiredo. Grid appliance user interface. <http://www.grid-appliance.org>, September 2009.
- [115] D. I. Wolinsky, K. Lee, P. O. Boykin, and R. Figueiredo. On the design of autonomic, decentralized vpns. In *International Conference on Collaborative Computing: Networking, Applications and Worksharing*, 2010.
- [116] D. I. Wolinsky, Y. Liu, P. S. Juste, G. Venkatasubramanian, and R. Figueiredo. On the design of scalable, self-configuring virtual networks. In *IEEE/ACM Supercomputing 2009*, November 2009.
- [117] D. I. Wolinsky, P. St. Juste, P. O. Boykin, and R. Figueiredo. Addressing the P2P bootstrap problem for small overlay networks. In *10th IEEE International Conference on Peer-to-Peer Computing (P2P)*, 2010.
- [118] C. P. Wright and E. Zadok. Unionfs: Bringing file systems together. In *Linux Journal*, December 2004.
- [119] XMPP Standards Foundation. Public XMPP services. <http://xmpp.org/services/>, December 2009.
- [120] J. Yonan. OpenVPN. <http://openvpn.net/>, 2009.

BIOGRAPHICAL SKETCH

David Isaac Wolinsky was born on October 31, 1982. He was blessed with an awesome, Isaac Emmanuel, born November 30, 2009. Beginning his studies in August 2001 at the University of Florida, David obtained the following degrees in electrical and computer engineering: Bachelor of Science in Spring 2005, Master of Science in Spring 2007, and Doctorate of Philosophy in Spring 2011. His advisor at the University of Florida was Professor Renato Figueiredo, whom he began working with since the during the Spring of 2006 at the Advanced Computing and Information Systems Lab.

His primary research focuses are network virtualization using structured P2P overlays and grid computing. The networking research has been realized in IPOP, a free (BSD) network virtualization software. Additionally, he has worked on enabling DHTs, decentralized NAT traversal through relays, software models for improved network virtualization, and autonomic virtual networking stacks. This work is a major contribution to his grid computing research focus, Grid Appliance, which enables the creation of decentralized, distributed grids using virtualized, physical, and cloud resources. Going forward, he expressed great interested in using these concepts in other distributed systems such as sensor networks, social networks, cloud services, or even web services.

During his free time, he enjoys time with my boy, running, playing basketball, and occasionally playing video games. At one point, he was ranked in the top 20 on the US East Warcraft III Free For All Ladder. Most of his time up to this point has been split between the Archer project and attempting to finish his Ph.D, which he proudly did before turning 30. Moving forward, he plans to research the concept known as hobbies and hopes to have a good time as a result.